R, Databases and Docker

John David Smith, Sophie Yang, M. Edward (Ed) Borasky, Scott Came, Mary Anne Thygesen, Ian Frantz, and Dipti Muni

2018-12-18

Contents

Chapter 1

Introduction

At the end of this chapter, you will be able to

- Understand the importance of using R and Docker to query a DBMS and access a service like Postgres outside of R.
- Setup your environment to explore the use-case for useRs.

1.1 Using R to query a DBMS in your organization

1.1.1 Why write a book about DBMS access from R using Docker?

- Large data stores in organizations are stored in databases that have specific access constraints and structural characteristics.
  * Data documentation may be incomplete, often emphasizes operational issues rather than analytic ones, and often needs to be confirmed on the fly.
  * Data volumes and query performance are important design constraints.

- R users frequently need to make sense of complex data structures and coding schemes to address incompletely formed questions so that exploratory data analysis has to be fast. * Exploratory and diagnostic techniques for the purpose should not be reinvented and would benefit from more public instruction or discussion.

- Learning to navigate the interfaces (passwords, packages, etc.) or gap between R and a database is difficult to simulate outside corporate walls.
  * Resources for interface problem diagnosis behind corporate walls may or may not address all the issues that R users face, so a simulated environment is needed.

- Docker is a relatively easy way to simulate the relationship between an R/Rstudio session and database – all on a single machine.

1.2 Docker as a tool for UseRs

Noam Ross's "Docker for the UseR" suggests that there are four distinct Docker use-cases for useRs.

1. Make a fixed working environment for reproducible analysis
2. Access a service outside of R **(e.g., Postgres)**
3. Create an R based service (e.g., with `plumber`)
4. Send our compute jobs to the cloud with minimal reconfiguration or revision

This book explores #2 because it allows us to work on the database access issues described above and to practice on an industrial-scale DBMS.

- Docker is a relatively easy way to simulate the relationship between an R/RStudio session and a database – all on on a single machine, provided you have Docker installed and running.
- You may want to run PostgreSQL on a Docker container, avoiding any OS or system dependencies that might come up.

1.3   Docker and R on your machine

Here is how R and Docker fit on your operating system in this tutorial:

(This diagram needs to be updated as our directory structure evolves.)

1.4   Who are we?

We have been collaborating on this book since the Summer of 2018, each of us chipping into the project as time permits:

- Dipti Muni - @deemuni
- Ian Franz - @ianfrantz
- Jim Tyhurst - @jimtyhurst
- John David Smith - @smithjd
- M. Edward (Ed) Borasky - @znmeb
- Maryanne Thygesen @maryannet
- Scott Came - @scottcame
- Sophie Yang - @SophieMYang

1.5   How did this project come about?

We trace this book back to the June 2, 2018 Cascadia R Conf where Aaron Makubuya gave a presentation using Vagrant hosting. After that John Smith, Ian Franz, and Sophie Yang had discussions after the monthly Data Discussion Meetups about the difficulties around setting up Vagrant, (a virtual environment), connecting to a corporate database and having realistic **public** environment to demo or practice the issues

that come up behind corporate firewalls. Scott Came's tutorial on R and Docker (an alternative to Vagrant) at the 2018 UseR Conference in Melbourne was provocative and it turned out he lived nearby. We reconnected with M. Edward (Ed) Borasky who had done extensive development for a Hack Oregon data science containerization project.

Chapter 2

Setup instructions (00)

This chapter explains:

- What you need to run the code in this book
- Where to get documentation for Docker
- How you can contribute to the book project

2.1 R, RStudio and Git

Most of you will probably have these already, but if you don't:

1. If you do not have R:
   - Go to https://cran.rstudio.com/.
   - Select the download link for your system. For Linux, choose your distro. We recommend Ubuntu 18.04 LTS "Bionic Beaver". It's much easier to find support answers on the web for Ubuntu than other distros.
   - Follow the instructions.
   - Note: if you already have R, make sure it's upgraded to R 3.5.1. We don't test on older versions!
2. If you do not have RStudio: go to https://www.rstudio.com/products/rstudio/download/#download. Make sure you have version 1.1.463 or later.
3. If you do not have Git:
   - On Windows, go to https://git-scm.com/download/win and follow instructions. There are a lot of options. Just pick the defaults!!!
   - On MacOS, go to https://sourceforge.net/projects/git-osx-installer/files/ and follow instructions.
   - On Linux, install Git from your distribution.

2.2 Docker

You will need Docker Community Edition (Docker CE).

- Windows: Go to https://store.docker.com/editions/community/docker-ce-desktop-windows. If you don't have a Docker Store login, you'll need to create one. Then:
  - If you have Windows 10 Pro, download and install Docker for Windows.
  - If you have an older version of Windows, download and install Docker Toolbox (https://docs.docker.com/toolbox/overview/).
  - Note that both versions require 64-bit hardware and the virtualization needs to be enabled in the firmware.

- MacOS: Go to https://store.docker.com/editions/community/docker-ce-desktop-mac. If you don't have a Docker Store login, you'll need to create one. Then download and install Docker for Mac. Your MacOS must be at least release Yosemite (10.10.3).
- Linux: note that, as with Windows and MacOS, you'll need a Docker Store login. Although most Linux distros ship with some version of Docker, chances are it's not the same as the official Docker CE version.
  - Ubuntu: https://store.docker.com/editions/community/docker-ce-server-ubuntu,
  - Fedora: https://store.docker.com/editions/community/docker-ce-server-fedora,
  - CentOS: https://store.docker.com/editions/community/docker-ce-server-centos,
  - Debian: https://store.docker.com/editions/community/docker-ce-server-debian.

Note that on Linux, you will need to be a member of the `docker` group to use Docker. To do that, execute `sudo usermod -aG docker ${USER}`. Then, log out and back in again.

2.3   Defining the PostgreSQL connection parameters

We use a PostgreSQL database server running in a Docker container for the database functions. To connect to it, you have to define some parameters. These parameters are used in two places:

1. When the Docker container is created, they're used to initialize the database, and
2. Whenever we connect to the database, we need to specify them to authenticate.

We define the parameters in an environment file that R reads when starting up. The file is called `.Renviron`, and is located in your home directory.

The easiest way to make this file is to copy the following R code and paste it into the R console:

```
cat(
  "\nDEFAULT_POSTGRES_USER_NAME=postgres",
  file = "~/.Renviron",
  sep = "",
  append = TRUE
)
cat(
  "\nDEFAULT_POSTGRES_PASSWORD=postgres\n",
  file = "~/.Renviron",
  sep = "",
  append = TRUE
)
```

2.4   Next steps

2.4.1   Browsing the book

If you just want to read the book and copy / paste code into your working environment, simply browse to https://smithjd.github.io/sql-pet. If you get stuck, or find things aren't working, open an issue at https://github.com/smithjd/sql-pet/issues/new/.

2.4.2   Diving in

If you want to experiment with the code in the book, run it in RStudio and interact with it, you'll need to do two more things:

1. Install the `sqlpetr` R package. See https://smithjd.github.io/sqlpetr for the package documentation. This will take some time; it is installing a number of packages.
2. Clone the Git repository https://github.com/smithjd/sql-pet.git and open the project file `sql-pet.Rproj` in RStudio.

Onward!

Chapter 3

How to use this book (01)

This chapter explains:

- The prerequisites for running the code in this book
- What R packages are used in the book

This book is full of examples that you can replicate on your computer.

3.1 Prerequisites

You will need:

- A computer running
  - Windows (Windows 7 64-bit or late - Windows 10-Pro is recommended)
  - MacOS
  - Linux (any Linux distro that will run Docker Community Edition, R and RStudio will work)
- Current versions of R and RStudio required.
- Docker (instructions below)
- Our companion package `sqlpetr` installs with: `devtools::install_github("smithjd/sqlpetr")`. Note that when you install the package it will ask you to update the packages it uses and that can take some time.

The database we use is PostgreSQL 10, but you do not need to install it - it's installed via a Docker image.

In addition to the current version of R and RStudio, you will need current versions of the following packages:

- `DBI`
- `DiagrammeR`
- `RPostgres`
- `dbplyr`
- `devtools`
- `downloader`
- `glue`
- `here`
- `knitr`
- `skimr`
- `tidyverse`

- `bookdown` (for compiling the book, if you want to)

3.2   Installing Docker

Install Docker. Installation depends on your operating system:

- On a Mac
- On UNIX flavors
- For Windows, consider these issues and follow these instructions.

3.3   Download the repo

The code to generate the book and the exercises it contains can be downloaded from this repo.

3.4   Read along, experiment as you go

We have never been sure whether we're writing an expository book or a massive tutorial. You may use it either way.

After the introductory chapters and the chapter that creates the persistent database ("The dvdrental database in Postgres in Docker (05)), you can jump around and each chapter stands on its own.

Chapter 4

Docker Hosting for Windows (02)

This chapter explains:

- How to setup your environment for Windows
- How to use Git and GitHub effectively on Windows

Skip these instructions if your computer has either OSX or a Unix variant.

4.1 Hardware requirements

You will need an Intel or AMD processor with 64-bit hardware and the hardware virtualization feature. Most machines you buy today will have that, but older ones may not. You will need to go into the BIOS / firmware and enable the virtualization feature. You will need at least 4 gigabytes of RAM!

4.2 Software requirements

You will need Windows 7 64-bit or later. If you can afford it, I highly recommend upgrading to Windows 10 Pro.

4.2.1 Windows 7, 8, 8.1 and Windows 10 Home (64 bit)

Install Docker Toolbox. The instructions are here: https://docs.docker.com/toolbox/toolbox_install_windows/. Make sure you try the test cases and they work!

4.2.2 Windows 10 Pro

Install Docker for Windows *stable*. The instructions are here: https://docs.docker.com/docker-for-windows/install/#start-docker-for-windows. Again, make sure you try the test cases and they work.

4.3 Additional technical details

See the Chapter on Additional technical details for Windows users (95) for more information.

Chapter 5

Learning Goals and Use Cases (03)

This chapter sets the context for the book by:

- Challenging you to think about your goals and expectations
- Imagining the setting where our sample database would be used
- Posing some imaginary use cases that a data analyst might face
- Discussing the different elements involved in answering questions from an organization's database

5.1 Ask yourself, what are you aiming for?

- Differences between production and data warehouse environments.
- Learning to keep your DBAs happy:
  - You are your own DBA in this simulation, so you can wreak havoc and learn from it, but you can learn to be DBA-friendly here.
  - In the end it's the subject-matter experts that understand your data, but you have to work with your DBAs first.

5.2 Learning Goals

After working through this tutorial, you can expect to be able to:

- Set up a PostgreSQL database in a Docker environment.
- Run queries against PostgreSQL in an environment that simulates what you will find in a corporate setting.
- Understand techniques and some of the trade-offs between:
  1. queries aimed at exploration or informal investigation using dplyr; and
  2. those where performance is important because of the size of the database or the frequency with which a query is run.
- Understand the equivalence between `dplyr` and SQL queries, and how R translates one into the other
- Understand some advanced SQL techniques.
- Gain familiarity with the standard metadata that a SQL database contains to describe its own contents.
- Gain some understanding of techniques for assessing query structure and performance.
- Understand enough about Docker to swap databases, e.g. Sports DB for the DVD rental database used in this tutorial. Or swap the database management system (DBMS), e.g. MySQL for PostgreSQL.

5.3   Imagining a DVD rental business

- Years ago people rented videos on DVD disks, and video stores were a big business.
- Imagine managing a video rental store like Movie Madness in Portland, Oregon.

- What data would be needed and what questions would you have to answer about the business?

This tutorial uses the Postgres version of "dvd rental" database which represents the transaction database for running a movie (e.g., dvd) rental business. The database can be downloaded here. Here's a glimpse of it's structure, which will be discussed in some detail:

A data analyst uses the database abstraction and the practical business questions to make better decision and solve problems.

5.4   Use cases

Imagine that you have one of following several roles at our fictional company **DVDs R Us** and you have a following need to be met:

- As a data scientist, I want to know the distribution of number of rentals per month per customer, so that the Marketing department can create incentives for customers in 3 segments: Frequent Renters, Average Renters, Infrequent Renters.
- As the Director of Sales, I want to see the total number of rentals per month for the past 6 months and I want to know how fast our customer base is growing/shrinking per month for the past 6 months.

Figure 5.1: Entity Relationship diagram for the dvdrental database

- As the Director of Marketing, I want to know which categories of DVDs are the least popular, so that I can create a campaign to draw attention to rarely used inventory.
- As a shipping clerk, I want to add rental information when I fulfill a shipment order.
- As the Director of Analytics, I want to test as much of the production R code in my shop as possible against a new release of the DBMS that the IT department is implementing next month.
- etc.

5.5   Investigating a question using with an organization's database

- Need both familiarity with the data and a focus question
    - An iterative process where
        * the data resource can shape your understanding of the question
        * the question you need to answer will frame how your see the data resource
    - You need to go back and forth between the two, asking
        * do I understand the question?
        * do I understand the data?
- How well do you understand the data resource (in the DBMS)?
    - Use all available documentation and understand its limits
    - Use your own tools and skills to examine the data resource
    - what's *missing* from the database: (columns, records, cells)
    - why is the missing data?
- How well do you understand the question you seek to answer?
    - How general or specific is your question?
    - How aligned is it with the purpose for which the database was designed and is being operated?
    - How different are your assumptions and concerns from those of the people who enter and use the data on a day to day basis?

Chapter 6

Connecting Docker, Postgres, and R (04)

This chapter demonstrates how to:

- Run, clean-up and close postgreSQL in docker containers.
- Keep necessary credentials secret while being available to R when it executes.
- Interact with PostgreSQL when it's running inside a Docker container.
- Read and write to PostgreSQL from R.

The following packages are used in this chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
require(knitr)
library(sqlpetr)
```

Please install the `sqlpetr` package if not already installed:

```
library(devtools)
if (!require(sqlpetr)) devtools::install_github("smithjd/sqlpetr")
```

Note that when you install the package the first time, it will ask you to update the packages it uses and that can take some time.

6.1 Verify that Docker is running

Docker commands can be run from a terminal (e.g., the Rstudio Terminal pane) or with a `system()` command. We provide the necessary functions to start, stop Docker containers and do other busy work in the `sqlpetr` package. As time permits and curiosity dictates, feel free to look at those functions to see how they work.

Check that docker is up and running:

```
sp_check_that_docker_is_up()
```

```
## [1] "Docker is up, running these containers:"
## [2] "CONTAINER ID      IMAGE           COMMAND                  CREATED         STATUS          PORTS
## [3] "933f5ff29218      postgres:10     \"docker-entrypoint.s…\"  3 days ago      Up 2 days        0.0.
```

6.2   Clean up if appropriate

Remove the `cattle` and `sql-pet` containers if they exists (e.g., from a prior experiments).

```
sp_docker_remove_container("cattle")
```

```
## [1] 0
```

```
sp_docker_remove_container("sql-pet")
```

```
## [1] 0
```

The convention we use in this book is to put docker commands in the `sqlpetr` package so that you can ignore them if you want. However, the functions are set up so that you can easily see how to do things with Docker and modify if you want.

We name containers `cattle` for "throw-aways" and `pet` for ones we treasure and keep around. :-)

```
sp_make_simple_pg("cattle")
```

```
## [1] 0
```

Docker returns a long string of numbers. If you are running this command for the first time, Docker downloads the PostgreSQL image, which takes a bit of time.

The following command shows that a container named `cattle` is running `postgres:10`. `postgres` is waiting for a connection:

```
sp_check_that_docker_is_up()
```

```
## [1] "Docker is up, running these containers:"
## [2] "CONTAINER ID     IMAGE           COMMAND               CREATED         STATUS             PORTS
## [3] "503fb3dc6e03     postgres:10       \"docker-entrypoint.s…\"  1 second ago     Up Less than a second
```

6.3   Connect, read and write to Postgres from R

6.3.1   Connect with Postgres

Connect to the postgrSQL using the `sp_get_postgres_connection` function:

```
con <- sp_get_postgres_connection(user = "postgres",
                       password = "postgres",
                       dbname = "postgres",
                       seconds_to_test = 10)
```

Notice that we are using the postgreSQL default username and password at this point and that it's in plain text. That is bad practice because user credentials should not be shared in this way. In a subsequent chapter we'll demonstrate how to store and use credentials to access the dbms.

Make sure that you can connect to the PostgreSQL database that you started earlier. If you have been executing the code from this tutorial, the database will not contain any tables yet:

```
dbListTables(con)
```

```
## character(0)
```

6.3.2   Interact with Postgres

Write `mtcars` to PostgreSQL

```r
dbWriteTable(con, "mtcars", mtcars, overwrite = TRUE)
```

List the tables in the PostgreSQL database to show that `mtcars` is now there:

```r
dbListTables(con)
```

```
## [1] "mtcars"
```

```r
# list the fields in mtcars:
dbListFields(con, "mtcars")
```

```
##  [1] "mpg"  "cyl"  "disp" "hp"   "drat" "wt"   "qsec" "vs"   "am"   "gear"
## [11] "carb"
```

Download the table from the DBMS to a local data frame:

```r
mtcars_df <- tbl(con, "mtcars")

# Show a few rows:
knitr::kable(head(mtcars_df))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

6.4 Clean up

Afterwards, always disconnect from the dbms:

```r
dbDisconnect(con)
```

Tell Docker to stop the `cattle` container:

```r
sp_docker_stop("cattle")
```

Tell Docker to remove the `cattle` container from it's library of active containers:

```r
sp_docker_remove_container("cattle")
```

```
## [1] 0
```

If we just **stop** the docker container but don't remove it (as we did with the `sp_docker_remove_container("cattle")` command), the `cattle` container will persist and we can start it up again later with `sp_docker_start("cattle")`. In that case, `mtcars` would still be there and we could retrieve it from postgreSQL again. Since `sp_docker_remove_container("cattle")` has removed it, the updated database has been deleted. (There are enough copies of `mtcars` in the world, so no great loss.)

Chapter 7

The dvdrental database in Postgres in Docker (05a)

This chapter demonstrates how to:

- Setup the `dvdrental` database in Docker
- Stop and start Docker container to demonstrate persistence
- Connect to and disconnect R from the `dvdrental` database
- Set up the environment for subsequent chapters

7.1 Overview

In the last chapter we connected to PostgreSQL from R. Now we set up a "realistic" database named `dvdrental`. There are different approaches to doing this: this chapter sets it up in a way that doesn't delve into the Docker details. If you are interested, you can look at an alternative approach in Creating the sql-pet Docker container a step at a time that breaks the process down into smaller chunks.

These packages are called in this Chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
library(glue)
require(knitr)
library(dbplyr)
library(sqlpetr)
library(bookdown)
```

7.2 Verify that Docker is up and running

```
sp_check_that_docker_is_up()
```

```
## [1] "Docker is up but running no containers"
```

7.3   Clean up if appropriate

Remove the `cattle` and `sql-pet` containers if they exist (e.g., from a prior runs):

```
sp_docker_remove_container("cattle")
```

```
## [1] 0
```

```
sp_docker_remove_container("sql-pet")
```

```
## [1] 0
```

7.4   Build the pet-sql Docker Image

Build an image that derives from postgres:10. The commands in `dvdrental.Dockerfile` creates a Docker container running PostgreSQL, and loads the `dvdrental` database. The dvdrental.Dockerfile is discussed below.

```r
docker_messages <- system2("docker",
        glue("build ", # tells Docker to build an image that can be loaded as a container
            "--tag postgres-dvdrental ", # (or -t) tells Docker to name the image
            "--file dvdrental.Dockerfile ", #(or -f) tells Docker to read `build` instructions from the d
            " . "),   # tells Docker to look for dvdrental.Dockerfile, and files it references, in the cur
            stdout = TRUE, stderr = TRUE)

cat(docker_messages, sep = "\n")
```

```
## Sending build context to Docker daemon   74.39MB


## Step 1/4 : FROM postgres:10
##  ---> ac25c2bac3c4
## Step 2/4 : WORKDIR /tmp
##  ---> Using cache
##  ---> 3f00a18e0bdf
## Step 3/4 : COPY init-dvdrental.sh /docker-entrypoint-initdb.d/
##  ---> Using cache
##  ---> 3453d61d8e3e
## Step 4/4 : RUN apt-get -qq update &&  apt-get install -y -qq curl zip  > /dev/null 2>&1 &&  curl -Os http:/
##  ---> Using cache
##  ---> f5e93aa64875
## Successfully built f5e93aa64875
## Successfully tagged postgres-dvdrental:latest
```

7.5   Run the pet-sql Docker Image

Run docker to bring up postgres. The first time it runs it will take a minute to create the PostgreSQL environment. There are two important parts to this that may not be obvious:

- The `source=` parameter points to dvdrental.Dockerfile, which does most of the heavy lifting. It has detailed, line-by-line comments to explain what it is doing.

- *Inside* dvdrental.Dockerfile the command `COPY init-dvdrental.sh /docker-entrypoint-initdb.d/` copies init-dvdrental.sh from the local file system into the specified location in the Docker container. When the PostgreSQL Docker container initializes, it looks for that file and executes it.

Doing all of that work behind the scenes involves two layers. Depending on how you look at it, that may be more or less difficult to understand than an alternative method.

```r
wd <- getwd()

docker_cmd <- glue(
  "run ",        # Run is the Docker command.  Everything that follows are `run` parameters.
  "--detach ", # (or `-d`) tells Docker to disconnect from the terminal / program issuing the command
  " --name sql-pet ",    # tells Docker to give the container a name: `sql-pet`
  "--publish 5432:5432 ", # tells Docker to expose the Postgres port 5432 to the local network with 543
  "--mount ", # tells Docker to mount a volume -- mapping Docker's internal file structure to the host
  "type=bind,", # tells Docker that the mount command points to an actual file on the host system
  'source="', # specifies the directory on the host to mount into the container at the mount point spec
  wd, '","', # the current working directory, as retrieved above
  "target=/petdir", # tells Docker to refer to the current directory as "/petdir" in its file system
  " postgres-dvdrental" # tells Docker to run the image was built in the previous step
)
```

If you are curious you can paste `docker_cmd` into a terminal window after the command 'docker':

```r
system2("docker", docker_cmd, stdout = TRUE, stderr = TRUE)
```

```
## [1] "5802b471a9cfebd31910a7913fc9476f3c5789e08e8439cf85564f846ace9e18"
```

7.6 Connect to Postgres with R

Use the DBI package to connect to the `dvdrental` database in PostgreSQL. Remember the settings discussion about [keeping passwords hidden][Pause for some security considerations]

```r
con <- sp_get_postgres_connection(password = "postgres",
                        user = "postgres",
                        dbname = "dvdrental",
                        seconds_to_test = 30)
```

List the tables in the database and the fields in one of those tables.

```r
dbListTables(con)
```

```
##  [1] "actor_info"            "customer_list"
##  [3] "film_list"             "nicer_but_slower_film_list"
##  [5] "sales_by_film_category"   "staff"
##  [7] "sales_by_store"        "staff_list"
##  [9] "category"              "film_category"
## [11] "country"               "actor"
## [13] "language"              "inventory"
## [15] "payment"               "rental"
## [17] "city"                  "store"
## [19] "film"                  "address"
## [21] "film_actor"            "customer"
```

```r
dbListFields(con, "rental")
```

```
## [1] "rental_id"    "rental_date"  "inventory_id" "customer_id"
## [5] "return_date"  "staff_id"     "last_update"
```

Disconnect from the database:

```
dbDisconnect(con)
```

7.7   Stop and start to demonstrate persistence

Stop the container:

```
sp_docker_stop("sql-pet")
```

Restart the container and verify that the dvdrental tables are still there:

```
sp_docker_start("sql-pet")
```

Connect to the `dvdrental` database in postgreSQL:

```
con <- sp_get_postgres_connection(user = "postgres",
                      password = "postgres",
                      dbname = "dvdrental",
                      seconds_to_test = 30)
```

Check that you can still see the fields in the `rental` table:

```
dbListFields(con, "rental")
```

```
## [1] "rental_id"    "rental_date"  "inventory_id" "customer_id"
## [5] "return_date"  "staff_id"     "last_update"
```

7.8   Cleaning up

Always have R disconnect from the database when you're done.

```
dbDisconnect(con)
```

Stop the `sql-pet` container:

```
sp_docker_stop("sql-pet")
```

Show that the container still exists even though it's not running

```
sp_show_all_docker_containers()
```

```
## CONTAINER ID    IMAGE            COMMAND            CREATED        STATUS             PORTS
## 5802b471a9cf    postgres-dvdrental  "docker-entrypoint.s…"  8 seconds ago    Exited (0) Less than a s
```

Next time, you can just use this command to start the container:

```
    sp_docker_start("sql-pet")
```

And once stopped, the container can be removed with:

```
    sp_check_that_docker_is_up("sql-pet")
```

7.9   Using the `sql-pet` container in the rest of the book

After this point in the book, we assume that Docker is up and that we can always start up our *sql-pet database* with:

```
    sp_docker_start("sql-pet")
```

Chapter 8

Securing and using your dbms credentials (05b)

This chapter demonstrates how to:

- Keep necessary credentials secret while being available to R when it executes.
- Interact with PostgreSQL using your secret dbms credentials

Connecting to a dbms can be very frustrating at first. In many organizations, simply **getting** access credentials takes time and may involve jumping through multiple hoops.

In addition, a dbms is terse or deliberately inscrutable when your credetials are incorrect. That's a security strategy, not a limitation of your understanding or your software. When R can't log you on to a dbms, you will have no information as to what went wrong.

The following packages are used in this chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
require(knitr)
library(sqlpetr)
```

8.1 Set up the sql-pet docker container

8.1.1 Verify that Docker is running

Check that docker is up and running:

```
sp_check_that_docker_is_up()
```

```
## [1] "Docker is up but running no containers"
```

8.1.2 Start the docker container:

Start the sql-pet docker container:

```
sp_docker_start("sql-pet")
```

8.2   Storing your dbms credentials

In previous chapters the connection string for connecting to the dbms has used default credentials specified in play text as follows:

```
user= 'postgres', password = 'postgres'
```

When we call `sp_get_postgres_connection` below we'll use environment variables that R obtains from reading the *.Renviron* file when R starts up. This approach has two benefits: that file is not uploaded to GitHub. R looks for it in your default directory every time it loads. To see whether you have already created that file, use the R Studio Files tab to look at your **home directory**:

That file should contain lines that **look like** the example below. Although in this example it contains the PostgreSQL default values for the username and password, they are obviously not secret. But this approach demonstrates where you should put secrets that R needs while not risking accidental uploaded to GitHub or some other public location..

Open your `.Renviron` file with this command:

```
file.edit("~/.Renviron")
```

Or you can execute define_postgresql_params.R to create the file or you could copy / paste the following into your **.Renviron** file:

```
DEFAULT_POSTGRES_PASSWORD=postgres
DEFAULT_POSTGRES_USER_NAME=postgres
```

Once that file is created, restart R, and after that R reads it every time it comes up.

8.2.1   Connect with Postgres using the Sys.getenv function

Connect to the postgrSQL using the `sp_get_postgres_connection` function:

```
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                        password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                        dbname = "dvdrental",
                        seconds_to_test = 10)
```

Once the connection object has been created, you can list all of the tables in the database:

```
dbListTables(con)
```

```
##  [1] "actor_info"                "customer_list"
##  [3] "film_list"                 "nicer_but_slower_film_list"
##  [5] "sales_by_film_category"    "staff"
##  [7] "sales_by_store"            "staff_list"
##  [9] "category"                  "film_category"
## [11] "country"                   "actor"
## [13] "language"                  "inventory"
## [15] "payment"                   "rental"
## [17] "city"                      "store"
## [19] "film"                      "address"
## [21] "film_actor"                "customer"
```

8.3 Clean up

Afterwards, always disconnect from the dbms:

```r
dbDisconnect(con)
```

Tell Docker to stop the `cattle` container:

```r
sp_docker_stop("sql-pet")
```

Chapter 9

Mapping your local environment (10)

This chapter explores:

- The different entities involved in running the examples in this book's sandbox
- The different roles that each entity plays in the sandbox
- How those entities are connected and how communication between those entities happens
- Pointers to the commands that go with each entity

These packages are used in this chapter:

```r
library(tidyverse)
library(DBI)
library(RPostgres)
require(knitr)
library(dbplyr)
library(sqlpetr)
library(DiagrammeR)
display_rows <- 5
```

9.1 Set up our standard pet-sql environment

Assume that the Docker container with PostgreSQL and the dvdrental database are ready to go. Start up the `docker-pet` container:

```r
sp_docker_start("sql-pet")
```

Connect to the `dvdrental` database with R.

```r
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

9.2 Sandbox Environment

Here is an overview of our sandbox environment. In this chapter we explore each of the entities in the sandbox, how they are connected and how they communicate with each other. You can skip this chapter

and come back later when you are curious about the setup that we're using in this book.
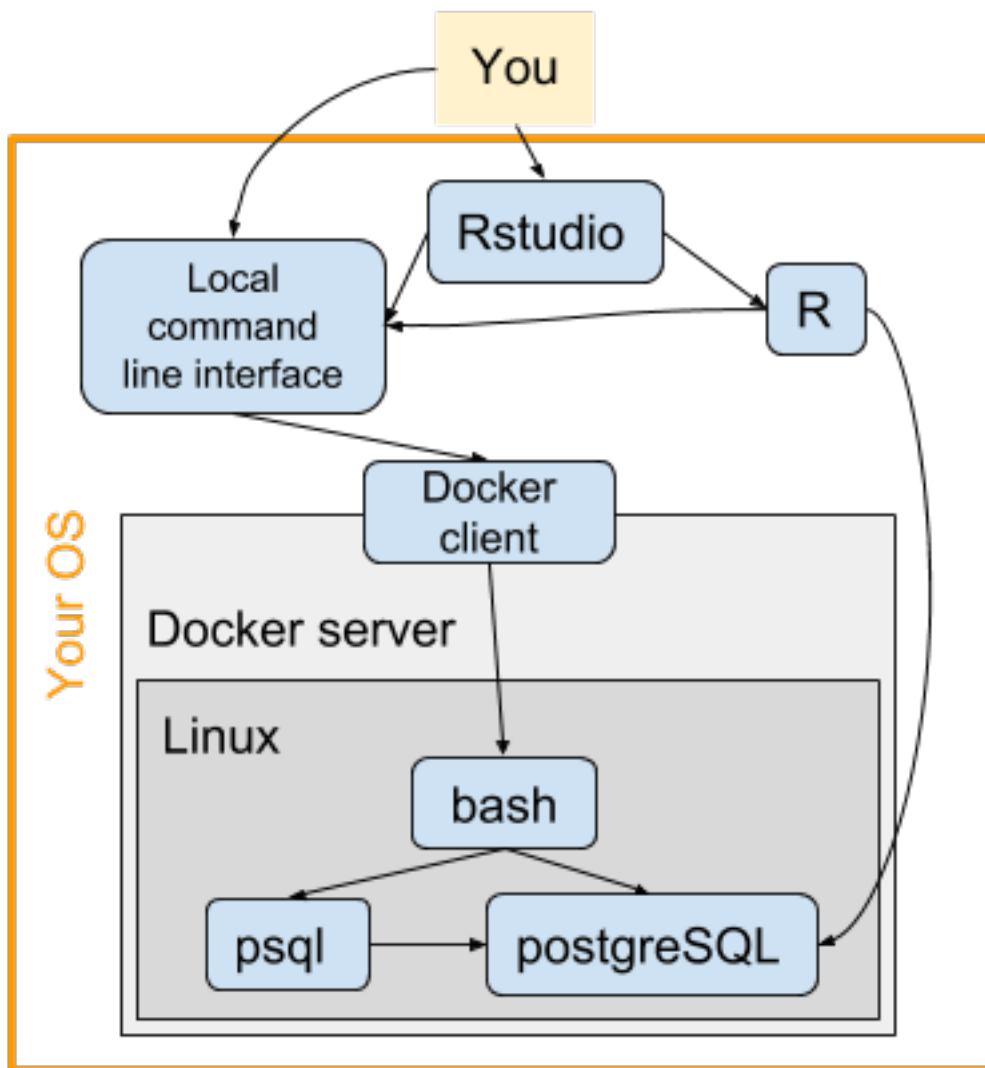
9.2.1   Sandbox entities and their roles

9.2.2   RStudio

You communicate with Rstudio, which can send commands to both R and to Unix. Commands to your OS can be entered directly in the terminal pane or via an R function like `exec2()`. On a Unix or Mac computer, you typically communicate with `bash`, while you have several choices on a Windows computer.

To check on the RStudio version you are using, enter this R command:

```
require(rstudioapi) versionInfo()
```

The RStudio cheat sheet is handy for learning your way around the IDE.

9.2.3 OS / local command line interface

You can type commands directly into a terminal window on your computer to communicate with your operating system (OS). It will be a `bash` prompt on a Unix or Mac, but could be one of several flavors on Windows. Our diagram conflates the operating system with the command line interface (CLI) which is a bit of a simplification as discussed below.

In addition to operating system commands, you can communicate with the Docker client through the CLI to start and stop the Docker server, load containers with programs such as Unix, postgreSQL, communicae with those programs, etc.

To check on the OS version you are using, enter this on your RStudio terminal or local CLI:

```
version -a
```

An OS can contain different comand line interfaces. Check on it with this on your RStudio terminal or local CLI:

```
echo $0
```

A Unix / Linux command line cheet sheet is a handy reference.

9.2.4 R

R processes instructions from Rstudio. It can send instructions to your OS via the `system2` function. R can also talks directly to postgreSQL through the DBI package.

R functions like `file.info("file.typ")` communicate with your operating system but do not visibly issue a command to your CLI. That's an example of an equivalence that can be useful or confusing (as in our environment diagram): you can get the same information from `ls -ql README.md` on a Unix command line as `file.info("README.md")` on the R console.

Although this sandbox seeks to make it easy, connecting to the database often involves technical and organizational hurdles like getting authorization. The main purpose of this book is to provide a sandbox for database queries to experiment with sending commands with the one of the *DBI* functions to the dbms directly from R. However, Docker and postreSQL commands are useful to know and may be necessary in extending the book's examples.

To check on the version of R that you are using, enter this on your R command line:

```
R.version
```

The growing collection of RStudio cheat sheets is indispensable.

9.2.5   Docker client

The docker client sets up the Docker server, loads containers, and passes instructions from your OS to the programs running in the Docker server. A Docker container will always contain a subset of the Linux operating system, so that it contains a second CLI in your sandbox. See more about the Docker environment.

In addition to interaction with docker through your computer's CLI or the RStudio terminal pane, the `docker` and `stevedore` packages can communicate with Docker from R. Both packages rely on the `reticulate` packaage and python.

For this book, we chose to send instructions to Docker through R's `system2()` function calls which do pass commands along to Docker through your computer's CLI. We chose that route in order to be as transparent as possible and because the book's sandbox environment is fairly simple. Although docker has different 44 commands, in this book we only use a subset: `ps`, `build`, `run`, `exec`, `start`, `stop`, and `rm`. We wrap all of these commands in `sqlpetr` package functions to encourage you to focus on R and postgreSQL.

To check on the Docker version you are using, enter this on your RStudio terminal or local CLI:

```
docker version
```

There are many Docker command line cheat sheets; this one is recommended.

9.2.6   In Docker: Linux

Docker runs a subset of the Linux operating system that in turn runs other programs like psql or postgreSQL. You may want to poke around the Linux environment inside Docker. To find what version of Linux Docker is running, enter the following command on your local CLI or in the RStudio terminal pane:

```
docker exec -ti sql-pet uname -a
```

As Linux can itself have different CLIs, enter the following command on your local CLI or in the RStudio terminal pane to find out which CLI is running inside Docker:

```
docker exec -ti sql-pet echo $0
```

To enter an interactive session inside Docker's Linux environment, enter the following command on your local CLI or in the RStudio terminal pane:

```
docker exec -ti sql-pet bash
```

To exit, enter:

```
exit
```

A Unix / Linux command line cheet sheet is a handy reference.

9.2.7   In Docker: `psql`

If you are comfortable executing SQL from a command line directly against the database, you can run the `psql` application in our Docker environment. To start up a `psql` session to investigate postgreSQL from a command line enter the following command on your computer's CLI or the RStudio terminal pane:

```
$ docker exec -ti sql-pet psql -a -p 5432 -d dvdrental -U postgres
```

Exit that environment with:

```
\q
```

Us this handy psql cheat sheet to get around.

9.2.8 In Docker: `postgreSQL`

The postgreSQL database is a whole environment unto itself. It can receive instructions through bash from `psql`, and it will respond to `DBI` queries from R on port 5282.

To check on the version of postgreSQL *client* (e.g., `psql`) you are using, enter this on your RStudio terminal or local command line interface:

```
docker exec -ti sql-pet psql --version
```

To check on the version of postgreSQL *server* you are running in Docker, enter this on your RStudio terminal or local command line interface:

```
docker exec -ti sql-pet psql -U postgres -c 'select version();'
```

Here's a recommended PostgreSQL cheat sheet.

9.3 Getting there from here: entity connections, equivalence, and commands

pathways, equivalences, command structures.

We use two trivial commands to explore the various *interfaces*. `ls -l` is the unix command for listing information about a file and `\du` is the psql command to list the users that exist in postgreSQL.

Your OS and the OS inside docker may be looking at the same file but they are in different time zones.

9.3.1 Get info on a local file from R code

```r
file.info("README.md")
```

```
##           size isdir mode                 mtime               ctime
## README.md 4973 FALSE  644 2018-11-01 14:23:56 2018-11-01 14:23:56
##                            atime uid gid uname grname
## README.md 2018-12-18 12:39:03 502  80   jds  admin
```

The equivalent information from executing a command on the CLI or terminal would be

```r
system2("ls",  "-l README.md", stdout = TRUE, stderr = FALSE)
```

```
## [1] "-rw-r--r--  1 jds  admin  4973 Nov  1 14:23 README.md"
```

9.3.2 Get info on the same OS file inside Docker from R Code

```r
system2("docker", "exec sql-pet ls -l petdir/README.md", stdout = TRUE, stderr = FALSE)
```

```
## Warning in system2("docker", "exec sql-pet ls -l petdir/README.md", stdout
## = TRUE, : running command ''docker' exec sql-pet ls -l petdir/README.md 2>/
## dev/null' had status 2
```

```
## character(0)
## attr(,"status")
## [1] 2
```

9.3.3   Docker and psql together from R or your CLI

As you become familiar with using docker, you'll see that there are various ways to do any given task. Here's an illustration of how to get a list of users who have access to the postegreSQL database.

```r
system2("docker", "exec sql-pet psql -U postgres -c '\\du' ",
        stdout = TRUE, stderr = FALSE)
```

```
## [1] "                                   List of roles"
## [2] " Role name |                        Attributes                         | Member of "
## [3] "-----------+------------------------------------------------------------+-----------"
## [4] " postgres  | Superuser, Create role, Create DB, Replication, Bypass RLS | {}"
## [5] ""
```

From the RStudio terminal window, the equivalent would be a matter of dropping off some of the R code:

```
docker exec -it sql-pet psql -U postgres -c '\du'
```

9.3.4   Nesting commands illustrates how entities are connected

The following tables illustrates how the different entities communicate with each other by decomposing a command from the chapter on creating a docker container one step at a time:

```
system2("docker", "exec sql-pet pg_restore -U postgres -d dvdrental petdir/dvdrental.tar",
stdout = TRUE, stderr = TRUE)
```

Code element	Comment
`system2(`	R command to send instructions to your computer's CLI.
`"docker",`	The program (docker) on your computer that will interpret the commands passed from the `system2` function.
`"`	The entire string within the quotes is passed to docker
`exec sql-pet`	`exec` will pass a command to any program running in the `sql-pet` container.
`pg_restore`	`pg_restore` is the program inside the `sql-pet` container that processes instructions to restore a previously downloaded backup file.
`-U postgres -d dvdrental petdir/dvdrental.tar`	The `pg_restore` program requires a username, a database and a backup file to be restored.
`",`	End of the docker commands passed to the `system2` function in R.
`stdout = TRUE, stderr = TRUE)`	The `system2` function needs to know what to do with its output, which in this case is to print all of it.

9.4   Exercises

Docker containers have a small foot print. In our container, we are running a limited Linux kernel and a Postgres database. To show how tiny the docker environment is, we will look at all the processes running

inside Docker and the top level file structure.

In the following exercies, use the `-i` option and the CONTAINER = `sql-pet`.

Start up R/RStudio and convert the CLI command to an R/RStudio command

#	Question	Docker CLI Command	R RStudio command	Local Command LINE
1	How many processes are running inside the Docker container?	docker exec -i sql-pet ps -eF		
1a	How many process are running on your local machine?			widows: tasklist Mac/Linux: ps -ef
2	What is the total number of files and directories in Docker?	docker exec -i sql-pet ls -al		
2a	What is the total number of files and directories on your local machine?			
3	Is Docker Running?	docker version		
3a	What are your Client and Server Versions?			
4	Does Postgres exist in the container?	docker ps -a		
4a	What is the status of Postgres?	docker ps -a		
4b	What is the size of Postgres?	docker images		
4c	What is the size of your laptop OS			https://www. quora.com/ What-is-the-actual-size-of-Windows
5	If sql-pet status is Up, How do I stop it?	docker stop sql-pet		
5a	If sql-pet status is Exited, How do I start it?	docker start sql-pet		

Chapter 10

Introduction to DBMS queries (11a)

This chapter demonstrates how to:

- Get a glimpse of what tables are in the database and what fields a table contains
- Download all or part of a table from the dbms
- See how `dplyr` code is translated into `SQL` commands
- Get acquainted with some useful tools for investigating a single table
- Begin thinking about how to divide the work between your local R session and the dbms

The following packages are used in this chapter:

```r
library(tidyverse)
library(DBI)
library(RPostgres)
library(dbplyr)
require(knitr)
library(bookdown)
library(sqlpetr)
```

Assume that the Docker container with PostgreSQL and the dvdrental database are ready to go. If not go back to Chapter 7

```r
sp_docker_start("sql-pet")
```

Connect to the database:

```r
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

10.1 Getting data from the database

As we show later on, the database serves as a store of data and as an engine for sub-setting, joining, and computation on the data. We begin with getting data from the dbms, or "downloading" data.

10.1.1   Finding out what's there

We've already seen the simplest way of getting a list of tables in a database with DBI functions that list tables and fields. Generate a vector listing the (public) tables in the database:

```
tables <- DBI::dbListTables(con)
tables
```

```
##  [1] "actor_info"               "customer_list"
##  [3] "film_list"                "nicer_but_slower_film_list"
##  [5] "sales_by_film_category"   "staff"
##  [7] "sales_by_store"           "staff_list"
##  [9] "category"                 "film_category"
## [11] "country"                  "actor"
## [13] "language"                 "inventory"
## [15] "payment"                  "rental"
## [17] "city"                     "store"
## [19] "film"                     "address"
## [21] "film_actor"               "customer"
```

Print a vector with all the fields (or columns or variables) in one specific table:

```
DBI::dbListFields(con, "rental")
```

```
## [1] "rental_id"    "rental_date"  "inventory_id" "customer_id"
## [5] "return_date"  "staff_id"     "last_update"
```

10.1.2   Listing all the fields for all the tables

The first example, DBI::dbListTables(con) returned 22 tables and the second example, DBI::dbListFields(con, "rental") returns 7 fields. Here we combine the two calls to return a list of tables which has a list of all the fields in the table. The code block just shows the first two tables.

```
table_columns <- lapply(tables, dbListFields, conn = con)

# or using purr:

table_columns <- map(tables, ~ dbListFields(.,conn = con) )

# rename each list [[1]] ... [[22]] to meaningful table name
names(table_columns) <- tables

head(table_columns)
```

```
## $actor_info
## [1] "actor_id"   "first_name" "last_name"  "film_info"
##
## $customer_list
## [1] "id"       "name"     "address"  "zip code" "phone"    "city"
## [7] "country"  "notes"    "sid"
##
## $film_list
## [1] "fid"         "title"       "description" "category"    "price"
## [6] "length"      "rating"      "actors"
##
## $nicer_but_slower_film_list
## [1] "fid"         "title"       "description" "category"    "price"
```

```
## [6] "length"      "rating"      "actors"
##
## $sales_by_film_category
## [1] "category"    "total_sales"
##
## $staff
##  [1] "staff_id"    "first_name" "last_name"  "address_id"  "email"
##  [6] "store_id"    "active"     "username"   "password"    "last_update"
## [11] "picture"
```

Later on we'll discuss how to get more extensive data about each table and column from the database's own store of metadata using a similar technique. As we go further the issue of scale will come up again and again: you need to be careful about how much data a call to the dbms will return, whether it's a list of tables or a table that could have millions of rows.

It's improtant to connect with people who own, generate, or are the subjects of the data. A good chat with people who own the data, generate it, or are the subjects can generate insights and set the context for your investigation of the database. The purpose for collecting the data or circumsances where it was collected may be burried far afield in an organization, but *usually someone knows*. The metadata discussed in a later chapter is essential but will only take you so far.

There are different ways of just **looking at the data**, which we explore below.

10.1.3 Downloading an entire table

There are many different methods of getting data from a DBMS, and we'll explore the different ways of controlling each one of them.

`DBI::dbReadTable` will download an entire table into an R tibble.

```
rental_tibble <- DBI::dbReadTable(con, "rental")
str(rental_tibble)
```

```
## 'data.frame':    16044 obs. of  7 variables:
##  $ rental_id   : int  2 3 4 5 6 7 8 9 10 11 ...
##  $ rental_date : POSIXct, format: "2005-05-24 22:54:33" "2005-05-24 23:03:39" ...
##  $ inventory_id: int  1525 1711 2452 2079 2792 3995 2346 2580 1824 4443 ...
##  $ customer_id : int  459 408 333 222 549 269 239 126 399 142 ...
##  $ return_date : POSIXct, format: "2005-05-28 19:40:33" "2005-06-01 22:12:39" ...
##  $ staff_id    : int  1 1 2 1 1 2 2 1 2 2 ...
##  $ last_update : POSIXct, format: "2006-02-16 02:30:53" "2006-02-16 02:30:53" ...
```

That's very simple, but if the table is large it may not be a good idea, since R is designed to keep the entire table in memory. Note that the first line of the str() output reports the total number of observations.

10.1.4 A table object that can be reused

The `dplyr::tbl` function gives us more control over access to a table by enabling control over which columns and rows to download. It creates an object that might **look** like a data frame, but it's actually a list object that `dplyr` uses for constructing queries and retrieving data from the DBMS.

```
rental_table <- dplyr::tbl(con, "rental")
```

10.1.5   Controlling the number of rows returned

The `collect` function triggers the creation of a tibble and controls the number of rows that the DBMS sends to R.

```
rental_table %>% collect(n = 3) %>% dim
```

```
## [1] 3 7
```

```
rental_table %>% collect(n = 500) %>% dim
```

```
## [1] 500   7
```

10.1.6   Random rows from the dbms

When the dbms contains many rows, a sample of the data may be plenty for your purposes. Although `dplyr` has nice functions to sample a data frame that's already in R (e.g., the `sample_n` and `sample_frac` functions), to get a sample from the dbms we have to use `dbGetQuery` to send native SQL to the database. To peak ahead, here is one example of a query that retrieves 20 rows from a 1% sample:

```
one_percent_sample <- DBI::dbGetQuery(
  con,
  "SELECT rental_id, rental_date, inventory_id, customer_id FROM rental TABLESAMPLE SYSTEM(1) LIMIT 20;
  "
)

one_percent_sample
```

```
## [1] rental_id    rental_date  inventory_id customer_id
## <0 rows> (or 0-length row.names)
```

10.1.7   Sub-setting variables

A table in the dbms may not only have many more rows than you want and also many more columns. The `select` command controls which columns are retrieved.

```
rental_table %>% select(rental_date, return_date) %>% head()
```

```
## # Source:    lazy query [?? x 2]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##   rental_date         return_date
##   <dttm>              <dttm>
## 1 2005-05-24 22:54:33 2005-05-28 19:40:33
## 2 2005-05-24 23:03:39 2005-06-01 22:12:39
## 3 2005-05-24 23:04:41 2005-06-03 01:43:41
## 4 2005-05-24 23:05:21 2005-06-02 04:33:21
## 5 2005-05-24 23:08:07 2005-05-27 01:32:07
## 6 2005-05-24 23:11:53 2005-05-29 20:34:53
```

That's exactly equivalent to submitting the following SQL commands dirctly:

```
DBI::dbGetQuery(
  con,
  'SELECT "rental_date", "return_date"
FROM "rental"
LIMIT 6')
```

```
##           rental_date          return_date
## 1 2005-05-24 22:54:33 2005-05-28 19:40:33
## 2 2005-05-24 23:03:39 2005-06-01 22:12:39
## 3 2005-05-24 23:04:41 2005-06-03 01:43:41
## 4 2005-05-24 23:05:21 2005-06-02 04:33:21
## 5 2005-05-24 23:08:07 2005-05-27 01:32:07
## 6 2005-05-24 23:11:53 2005-05-29 20:34:53
```

We won't discuss `dplyr` methods for sub-setting variables, deriving new ones, or sub-setting rows based on the values found in the table because they are covered well in other places, including:

- Comprehensive reference: https://dplyr.tidyverse.org/
- Good tutorial: https://suzan.rbind.io/tags/dplyr/

In practice we find that, **renaming variables** is often quite important because the names in an SQL database might not meet your needs as an analyst. In "the wild" you will find names that are ambiguous or overly specified, with spaces in them, and other problems that will make them difficult to use in R. It is good practice to do whatever renaming you are going to do in a predictable place like at the top of your code. The names in the `dvdrental` database are simple and clear, but if they were not, you might rename them for subsequent use in this way:

```
tbl(con, "rental") %>%
  rename(rental_id_number = rental_id, inventory_id_number = inventory_id) %>%
  select(rental_id_number, rental_date, inventory_id_number) %>%
  head()
```

```
## # Source:    lazy query [?? x 3]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##   rental_id_number rental_date         inventory_id_number
##              <int> <dttm>                            <int>
## 1                2 2005-05-24 22:54:33                1525
## 2                3 2005-05-24 23:03:39                1711
## 3                4 2005-05-24 23:04:41                2452
## 4                5 2005-05-24 23:05:21                2079
## 5                6 2005-05-24 23:08:07                2792
## 6                7 2005-05-24 23:11:53                3995
```

That's equivalent to the following SQL code:

```
DBI::dbGetQuery(
  con,
  'SELECT "rental_id_number", "rental_date", "inventory_id_number"
FROM (SELECT "rental_id" AS "rental_id_number", "rental_date", "inventory_id" AS "inventory_id_number",
FROM "rental") "ihebfvnxvb"
LIMIT 6' )
```

```
##   rental_id_number         rental_date inventory_id_number
## 1                2 2005-05-24 22:54:33                1525
## 2                3 2005-05-24 23:03:39                1711
## 3                4 2005-05-24 23:04:41                2452
## 4                5 2005-05-24 23:05:21                2079
## 5                6 2005-05-24 23:08:07                2792
## 6                7 2005-05-24 23:11:53                3995
```

The one difference is that the `SQL` code returns a regular data frame and the `dplyr` code returns a `tibble`. Notice that the seconds are greyed out in the `tibble` display.

10.1.8   Translating `dplyr` code to `SQL` queries

Where did the translations we've showon above come from? The `show_query` function shows how `dplyr` is translating your query to the dialect of the target dbms:

```
rental_table %>%
  count(staff_id) %>%
  show_query()
```

```
## <SQL>
## SELECT "staff_id", COUNT(*) AS "n"
## FROM "rental"
## GROUP BY "staff_id"
```

Here is an extensive discussion of how `dplyr` code is translated into SQL:

- https://dbplyr.tidyverse.org/articles/sql-translation.html

The SQL code can submit the same query directly to the DBMS with the `DBI::dbGetQuery` function:

```
DBI::dbGetQuery(
  con,
  'SELECT "staff_id", COUNT(*) AS "n"
   FROM "rental"
   GROUP BY "staff_id";
   '
)
```

```
##   staff_id    n
## 1        2 8004
## 2        1 8040
```

<<smy We haven't investigated this, but it looks like `dplyr` collect() function triggers a call simmilar to the dbGetQuery call above. The default `dplyr` behavior looks like dbSendQuery() and dbFetch() model is used.>>

When you create a report to run repeatedly, you might want to put that query into R markdown. That way you can also execute that SQL code in a chunk with the following header:

```
{sql, connection=con, output.var = "query_results"}
```

```
SELECT "staff_id", COUNT(*) AS "n"
FROM "rental"
GROUP BY "staff_id";
```

Rmarkdown stores that query result in a tibble which can be printed by referring to it:

```
query_results
```

```
##   staff_id    n
## 1        2 8004
## 2        1 8040
```

10.2   Examining a single table with R

Dealing with a large, complex database highlights the utility of specific tools in R. We include brief examples that we find to be handy:

- Base R structure: `str`
- printing out some of the data: `datatable`, `kable`, and `View`

- summary statistics: `summary`
- `glimpse` oin the `tibble` package, which is included in the `tidyverse`
- `skim` in the `skimr` package

10.2.1 `str` - a base package workhorse

`str` is a workhorse function that lists variables, their type and a sample of the first few variable values.

```
str(rental_tibble)
```

```
## 'data.frame':    16044 obs. of  7 variables:
##  $ rental_id   : int  2 3 4 5 6 7 8 9 10 11 ...
##  $ rental_date : POSIXct, format: "2005-05-24 22:54:33" "2005-05-24 23:03:39" ...
##  $ inventory_id: int  1525 1711 2452 2079 2792 3995 2346 2580 1824 4443 ...
##  $ customer_id : int  459 408 333 222 549 269 239 126 399 142 ...
##  $ return_date : POSIXct, format: "2005-05-28 19:40:33" "2005-06-01 22:12:39" ...
##  $ staff_id    : int  1 1 2 1 1 2 2 1 2 2 ...
##  $ last_update : POSIXct, format: "2006-02-16 02:30:53" "2006-02-16 02:30:53" ...
```

10.2.2 Always look at your data with `head`, `View`, or `kable`

There is no substitute for looking at your data and R provides several ways to just browse it. The `head` function controls the number of rows that are displayed. Note that tail does not work against a database object. In every-day practice you would look at more than the default 6 rows, but here we wrap `head` around the data frame:

```
sp_print_df(head(rental_tibble))
```

rental_id	rental_date	inventory_id	customer_id	return_date	staff_id	last_update
2	2005-05-24 22:54:33	1525	459	2005-05-28 19:40:33	1	2006-02-16 02:30:53
3	2005-05-24 23:03:39	1711	408	2005-06-01 22:12:39	1	2006-02-16 02:30:53
4	2005-05-24 23:04:41	2452	333	2005-06-03 01:43:41	2	2006-02-16 02:30:53
5	2005-05-24 23:05:21	2079	222	2005-06-02 04:33:21	1	2006-02-16 02:30:53
6	2005-05-24 23:08:07	2792	549	2005-05-27 01:32:07	1	2006-02-16 02:30:53
7	2005-05-24 23:11:53	3995	269	2005-05-29 20:34:53	2	2006-02-16 02:30:53

10.2.3 The `summary` function in base

The basic statistics that the base package `summary` provides can serve a unique diagnostic purpose in this context. For example, the following output shows that `rental_id` is a sequential number from 1 to 16,049 with no gaps. The same is true of `inventory_id`. The number of NA's is a good first guess as to the number of dvd's rented out or lost on 2005-09-02 02:35:22.

```
summary(rental_tibble)
```

```
##    rental_id     rental_date                    inventory_id
##  Min.   :    1  Min.   :2005-05-24 22:53:30  Min.   :   1
##  1st Qu.: 4014  1st Qu.:2005-07-07 00:58:40  1st Qu.:1154
##  Median : 8026  Median :2005-07-28 16:04:32  Median :2291
##  Mean   : 8025  Mean   :2005-07-23 08:13:34  Mean   :2292
##  3rd Qu.:12037  3rd Qu.:2005-08-17 21:16:23  3rd Qu.:3433
##  Max.   :16049  Max.   :2006-02-14 15:16:03  Max.   :4581
##
##   customer_id     return_date                    staff_id
```

```
## Min.   :  1.0   Min.    :2005-05-25 23:55:21   Min.    :1.000
## 1st Qu.:148.0   1st Qu.:2005-07-10 15:49:36   1st Qu.:1.000
## Median :296.0   Median :2005-08-01 19:45:29   Median :1.000
## Mean   :297.1   Mean    :2005-07-25 23:58:03   Mean    :1.499
## 3rd Qu.:446.0   3rd Qu.:2005-08-20 23:35:55   3rd Qu.:2.000
## Max.   :599.0   Max.    :2005-09-02 02:35:22   Max.    :2.000
##                 NA's    :183
##   last_update
## Min.   :2006-02-15 21:30:53
## 1st Qu.:2006-02-16 02:30:53
## Median :2006-02-16 02:30:53
## Mean   :2006-02-16 02:31:31
## 3rd Qu.:2006-02-16 02:30:53
## Max.   :2006-02-23 09:12:08
##
```

10.2.4 The `glimpse` function in the `tibble` package

The `tibble` package's `glimpse` function is a more compact version of `str`:

```
tibble::glimpse(rental_tibble)
```

```
## Observations: 16,044
## Variables: 7
## $ rental_id    <int> 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 1...
## $ rental_date  <dttm> 2005-05-24 22:54:33, 2005-05-24 23:03:39, 2005-0...
## $ inventory_id <int> 1525, 1711, 2452, 2079, 2792, 3995, 2346, 2580, 1...
## $ customer_id  <int> 459, 408, 333, 222, 549, 269, 239, 126, 399, 142,...
## $ return_date  <dttm> 2005-05-28 19:40:33, 2005-06-01 22:12:39, 2005-0...
## $ staff_id     <int> 1, 1, 2, 1, 1, 2, 2, 1, 2, 2, 2, 1, 1, 1, 2, 1, 2...
## $ last_update  <dttm> 2006-02-16 02:30:53, 2006-02-16 02:30:53, 2006-0...
```

10.2.5 The `skim` function in the `skimr` package

The `skimr` package has several functions that make it easy to examine an unknown data frame and assess what it contains. It is also extensible.

```
library(skimr)
```

```
##
## Attaching package: 'skimr'
```

```
## The following object is masked from 'package:knitr':
##
##     kable
```

```
skim(rental_tibble)
```

```
## Skim summary statistics
##  n obs: 16044
##  n variables: 7
##
## -- Variable type:integer ----------------------------------------------------------------------
##      variable missing complete     n    mean      sd p0   p25    p50
##   customer_id       0    16044 16044  297.14  172.45  1   148    296
##   inventory_id      0    16044 16044 2291.84 1322.21  1  1154   2291
```

```
##       rental_id         0     16044 16044 8025.37 4632.78  1 4013.75 8025.5
##        staff_id         0     16044 16044    1.5      0.5  1    1        1
##       p75   p100     hist
##   446       599
##  3433      4581
## 12037.25 16049
##     2         2
##
## -- Variable type:POSIXct -----------------------------------------------------------------
##      variable missing complete     n        min        max     median
##  last_update        0     16044 16044 2006-02-15 2006-02-23 2006-02-16
##  rental_date        0     16044 16044 2005-05-24 2006-02-14 2005-07-28
##  return_date      183     15861 16044 2005-05-25 2005-09-02 2005-08-01
##  n_unique
##         3
##     15815
##     15836
```
```r
wide_rental_skim <- skim_to_wide(rental_tibble)
```

10.2.6 Close the connection and shut down sql-pet

Where you place the `collect` function matters.

```r
dbDisconnect(con)
sp_docker_stop("sql-pet")
```

10.3 Additional reading

- Wickham (2018)
- Baumer (2018)

Chapter 11

Lazy Evaluation and Lazy Queries (11b)

11.1 This chapter:

- Reviews lazy evaluation and discusses its interaction with remote query execution on a dbms
- Demonstrates how `dplyr` queries behave in connection with several different functions
- Offers some further resources on lazy loading, evaluation, execution, etc.

11.1.1 Setup

The following packages are used in this chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
library(dbplyr)
require(knitr)
library(bookdown)
library(sqlpetr)
```

Assume that the Docker container with PostgreSQL and the dvdrental database are ready to go. If not go back to the previous Chapter

```
sp_docker_start("sql-pet")
```

Connect to the database:

```
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

11.2 R is lazy and comes with guardrails

By design, R is both a language and an interactive development environment (IDE). As a language, R tries to be as efficient as possible. As an IDE, R creates some guardrails to make it easy and safe to work with your data. For example `getOption("max.print")` prevents R from printing more rows of data than you can handle, with a nice default of 99999, which may or may not suit you.

On the other hand SQL is a *"Structured Query Language (SQL): a standard computer language for relational database management and data manipulation."*.[1] SQL has has various database-specific Interactive Development Environments (IDEs): for postgreSQL it's pgAdmin. Roger Peng explains in R Programming for Data Science that:

> R has maintained the original S philosophy, which is that it provides a language that is both useful for interactive work, but contains a powerful programming language for developing new tools.

This is complicated when R interacts with SQL. In the vignette for dbplyr Hadley Wikham explains:

> The most important difference between ordinary data frames and remote database queries is that your R code is translated into SQL and executed in the database on the remote server, not in R on your local machine. When working with databases, dplyr tries to be as lazy as possible:
>
> - It never pulls data into R unless you explicitly ask for it.
>
> - It delays doing any work until the last possible moment: it collects together everything you want to do and then sends it to the database in one step.

Exactly when, which and how much data is returned from the dbms is the topic of this chapter. Exactly how the data is represented in the dbms and then translated to a data frame is discussed in the DBI specification.

Eventually, if you are interacting with a dbms from R you will need to understand the differences between lazy loading, lazy evaluation, and lazy queries.

11.2.1 Lazy loading

"*Lazy loading is always used for code in packages but is optional (selected by the package maintainer) for datasets in packages.*"[2] Lazy loading means that the code for a particular function doesn't actually get loaded into memory until the last minute – when it's actually being used.

11.2.2 Lazy evaluation

Essentially "Lazy evaluation is a programming strategy that allows a symbol to be evaluated only when needed."[3] That means that lazy evaluation is about **symbols** such as function arguments[4] when they are evaluated. Tidy evaluation complicates lazy evaluation.[5]

11.2.3 Lazy Queries

"*When you create a "lazy" query, you're creating a pointer to a set of conditions on the database, but the query isn't actually run and the data isn't actually loaded until you call "next" or some similar method to*

---

[1] https://www.techopedia.com/definition/1245/structured-query-language-sql
[2] https://cran.r-project.org/doc/manuals/r-release/R-ints.html#Lazy-loading
[3] https://colinfay.me/lazyeval/
[4] http://adv-r.had.co.nz/Functions.html#function-arguments
[5] https://colinfay.me/tidyeval-1/

actually fetch the data and load it into an object."[6] The `collect()` function retrieves data into a local tibble.[7]

11.3 Lazy evaluation and lazy queries

11.3.1 Dplyr connection objects

As introduced in the previous chapter, the `dplyr::tbl` function creates an object that might **look** like a data frame in that when you enter it on the command line, it prints a bunch of rows from the dbms table. But is actually a **list** object that `dplyr` uses for constructing queries and retrieving data from the DBMS.

The following code illustrates these issues. The `dplyr::tbl` function creates the connection object that we store in an object named `rental_table`:

```
rental_table <- dplyr::tbl(con, "rental")
```

At first glance, it kind of **looks** like a data frame although it only prints 10 of the table's 16,044 rows:

```
rental_table
```

```
## # Source:   table<rental> [?? x 7]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##    rental_id rental_date         inventory_id customer_id
##        <int> <dttm>                     <int>       <int>
## 1          2 2005-05-24 22:54:33         1525         459
## 2          3 2005-05-24 23:03:39         1711         408
## 3          4 2005-05-24 23:04:41         2452         333
## 4          5 2005-05-24 23:05:21         2079         222
## 5          6 2005-05-24 23:08:07         2792         549
## 6          7 2005-05-24 23:11:53         3995         269
## 7          8 2005-05-24 23:31:46         2346         239
## 8          9 2005-05-25 00:00:40         2580         126
## 9         10 2005-05-25 00:02:21         1824         399
## 10        11 2005-05-25 00:09:02         4443         142
## # ... with more rows, and 3 more variables: return_date <dttm>,
## #   staff_id <int>, last_update <dttm>
```

But consider the structure of `rental_table`:

```
str(rental_table)
```

```
## List of 2
##  $ src:List of 2
##   ..$ con  :Formal class 'PqConnection' [package "RPostgres"] with 3 slots
##   .. .. ..@ ptr     :<externalptr>
##   .. .. ..@ bigint  : chr "integer64"
##   .. .. ..@ typnames:'data.frame':   437 obs. of  2 variables:
##   .. .. .. ..$ oid    : int [1:437] 16 17 18 19 20 21 22 23 24 25 ...
##   .. .. .. ..$ typname: chr [1:437] "bool" "bytea" "char" "name" ...
##   ..$ disco: NULL
##   ..- attr(*, "class")= chr [1:3] "src_dbi" "src_sql" "src"
##  $ ops:List of 2
##   ..$ x   : 'ident' chr "rental"
##   ..$ vars: chr [1:7] "rental_id" "rental_date" "inventory_id" "customer_id" ...
```

---

[6]https://www.quora.com/What-is-a-lazy-query
[7]https://dplyr.tidyverse.org/reference/compute.html

```
##   ..- attr(*, "class")= chr [1:3] "op_base_remote" "op_base" "op"
##  - attr(*, "class")= chr [1:4] "tbl_dbi" "tbl_sql" "tbl_lazy" "tbl"
```

It has two rows. The first row contains all the information in the `con` object, which contains information about all the tables and objects in the database:

```
rental_table$src$con@typnames$typname[380:437]
```

```
##  [1] "customer"                    "_customer"
##  [3] "actor_actor_id_seq"          "actor"
##  [5] "_actor"                      "category_category_id_seq"
##  [7] "category"                    "_category"
##  [9] "film_film_id_seq"            "film"
## [11] "_film"                       "pg_toast_16434"
## [13] "film_actor"                  "_film_actor"
## [15] "film_category"               "_film_category"
## [17] "actor_info"                  "_actor_info"
## [19] "address_address_id_seq"      "address"
## [21] "_address"                    "city_city_id_seq"
## [23] "city"                        "_city"
## [25] "country_country_id_seq"      "country"
## [27] "_country"                    "customer_list"
## [29] "_customer_list"              "film_list"
## [31] "_film_list"                  "inventory_inventory_id_seq"
## [33] "inventory"                   "_inventory"
## [35] "language_language_id_seq"    "language"
## [37] "_language"                   "nicer_but_slower_film_list"
## [39] "_nicer_but_slower_film_list" "payment_payment_id_seq"
## [41] "payment"                     "_payment"
## [43] "rental_rental_id_seq"        "rental"
## [45] "_rental"                     "sales_by_film_category"
## [47] "_sales_by_film_category"     "staff_staff_id_seq"
## [49] "staff"                       "_staff"
## [51] "pg_toast_16529"              "store_store_id_seq"
## [53] "store"                       "_store"
## [55] "sales_by_store"              "_sales_by_store"
## [57] "staff_list"                  "_staff_list"
```

The second row contains a list of the columns in the `rental` table, among other things:

```
rental_table$ops$vars
```

```
## [1] "rental_id"    "rental_date"  "inventory_id" "customer_id"
## [5] "return_date"  "staff_id"     "last_update"
```

To illustrate the different issues involved in data retrieval, we create equivalent connection objects to link to two other tables.

```
staff_table <- dplyr::tbl(con, "staff")
# the 'staff' table has 2 rows


customer_table <- dplyr::tbl(con, "customer")
# the 'customer' table has 599 rows
```

11.4 When does a lazy query trigger data retrieval?

11.4.1 Create a black box query for experimentation

Here is a typical string of dplyr verbs strung together with the magrittr `%>%` command that will be used to tease out the several different behaviors that a lazy query has when passed to different R functions. This query joins three connection objects into a query we'll call `Q`:

```
Q <- rental_table %>%
  left_join(staff_table, by = c("staff_id" = "staff_id")) %>%
  rename(staff_email = email) %>%
  left_join(customer_table, by = c("customer_id" = "customer_id")) %>%
  rename(customer_email = email) %>%
  select(rental_date, staff_email, customer_email)
```

11.4.2 Experiment overview

Think of `Q` as a black box for the moment. The following examples will show how `Q` is interpreted differently by different functions. In this table, a single green check indicates that some rows are returned, two green checks indicates that all the rows are returned, and the red X indicates that no rows have are returned.

R code	Result
`Q %>% print()`	Prints x rows; same as just entering `Q`
`Q %>% as.tibble()`	Forces `Q` to be a tibble
`Q %>% head()`	Prints the first 6 rows
`Q %>% length()`	Counts the rows in `Q`
`Q %>% str()`	Shows the top 3 levels of the **object Q**
`Q %>% nrow()`	**Attempts** to determine the number of rows
`Q %>% tally()`	Counts all the rows – on the dbms side
`Q %>% collect(n = 20)`	Prints 20 rows

R code	Result
`Q %>% collect(n = 20) %>%` `head()`	✔ Prints 6 rows
`Q %>% show_query()`	✘ **Translates** the lazy query object into SQL

`Qc <- Q %>%`
`count(customer_email, sort`
`= TRUE) Qc`
	✘ **Extends** the lazy query object

(The next chapter will discuss how to build queries and how to explore intermediate steps.)

11.4.3   Q %>% print()

Remember that `Q %>% print()` is equivalent to `print(Q)` and the same as just entering `Q` on the command line. We use the magrittr pipe operator here because chaining functions highlights how the same object behaves differently in each use.

```
Q %>% print()
```

```
## # Source:    lazy query [?? x 3]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##    rental_date          staff_email             customer_email
##    <dttm>               <chr>                   <chr>
##  1 2005-05-24 22:54:33 Mike.Hillyer@sakilasta~ tommy.collazo@sakilacustome~
##  2 2005-05-24 23:03:39 Mike.Hillyer@sakilasta~ manuel.murrell@sakilacustom~
##  3 2005-05-24 23:04:41 Jon.Stephens@sakilasta~ andrew.purdy@sakilacustomer~
##  4 2005-05-24 23:05:21 Mike.Hillyer@sakilasta~ delores.hansen@sakilacustom~
##  5 2005-05-24 23:08:07 Mike.Hillyer@sakilasta~ nelson.christenson@sakilacu~
##  6 2005-05-24 23:11:53 Jon.Stephens@sakilasta~ cassandra.walters@sakilacus~
##  7 2005-05-24 23:31:46 Jon.Stephens@sakilasta~ minnie.romero@sakilacustome~
##  8 2005-05-25 00:00:40 Mike.Hillyer@sakilasta~ ellen.simpson@sakilacustome~
##  9 2005-05-25 00:02:21 Jon.Stephens@sakilasta~ danny.isom@sakilacustomer.o~
## 10 2005-05-25 00:09:02 Jon.Stephens@sakilasta~ april.burns@sakilacustomer.~
## # ... with more rows
```

✔   R retrieves 10 observations and 3 columns. In its role as IDE, R has provided nicely formatted output that is similar to what it prints for a tibble, with descriptive information about the dataset and each column:

> # Source: lazy query [?? x 3] # Database: postgres [postgres@localhost:5432/dvdrental]
> rental_date staff_email customer_email <dttm> <chr> <chr>

R has not determined how many rows are left to retrieve as it notes `... with more rows`.

11.4.4 Q %>% as.tibble()

In contrast to `print()`, the `as.tibble()` function causes R to download the whole table, using tibble's default of displaying only the first 10 rows.

```
Q %>% as.tibble()
```

```
## # A tibble: 16,044 x 3
##    rental_date         staff_email             customer_email
##    <dttm>              <chr>                   <chr>
##  1 2005-05-24 22:54:33 Mike.Hillyer@sakilasta~ tommy.collazo@sakilacustome~
##  2 2005-05-24 23:03:39 Mike.Hillyer@sakilasta~ manuel.murrell@sakilacustom~
##  3 2005-05-24 23:04:41 Jon.Stephens@sakilasta~ andrew.purdy@sakilacustomer~
##  4 2005-05-24 23:05:21 Mike.Hillyer@sakilasta~ delores.hansen@sakilacustom~
##  5 2005-05-24 23:08:07 Mike.Hillyer@sakilasta~ nelson.christenson@sakilacu~
##  6 2005-05-24 23:11:53 Jon.Stephens@sakilasta~ cassandra.walters@sakilacus~
##  7 2005-05-24 23:31:46 Jon.Stephens@sakilasta~ minnie.romero@sakilacustome~
##  8 2005-05-25 00:00:40 Mike.Hillyer@sakilasta~ ellen.simpson@sakilacustome~
##  9 2005-05-25 00:02:21 Jon.Stephens@sakilasta~ danny.isom@sakilacustomer.o~
## 10 2005-05-25 00:09:02 Jon.Stephens@sakilasta~ april.burns@sakilacustomer.~
## # ... with 16,034 more rows
```

11.4.5 Q %>% head()

The `head()` function is very similar to print but has a different "`max.print`" value.

```
Q %>% head()
```

```
## # Source:   lazy query [?? x 3]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##   rental_date         staff_email             customer_email
##   <dttm>              <chr>                   <chr>
## 1 2005-05-24 22:54:33 Mike.Hillyer@sakilasta~ tommy.collazo@sakilacustomer~
## 2 2005-05-24 23:03:39 Mike.Hillyer@sakilasta~ manuel.murrell@sakilacustome~
## 3 2005-05-24 23:04:41 Jon.Stephens@sakilasta~ andrew.purdy@sakilacustomer.~
## 4 2005-05-24 23:05:21 Mike.Hillyer@sakilasta~ delores.hansen@sakilacustome~
## 5 2005-05-24 23:08:07 Mike.Hillyer@sakilasta~ nelson.christenson@sakilacus~
## 6 2005-05-24 23:11:53 Jon.Stephens@sakilasta~ cassandra.walters@sakilacust~
```

11.4.6 Q %>% length()

Because the `Q` object is relatively complex, using `str()` on it prints many lines. You can glimpse what's going on with `length()`:

```
Q %>% length()
```

```
## [1] 2
```

11.4.7   Q %>% str()

Looking inside shows some of what's going on (three levels deep):

```
Q %>% str(max.level = 3)
```

```
## List of 2
##  $ src:List of 2
##   ..$ con  :Formal class 'PqConnection' [package "RPostgres"] with 3 slots
##   ..$ disco: NULL
##   ..- attr(*, "class")= chr [1:3] "src_dbi" "src_sql" "src"
##  $ ops:List of 4
##   ..$ name: chr "select"
##   ..$ x    :List of 4
##   .. ..$ name: chr "rename"
##   .. ..$ x    :List of 4
##   .. .. ..- attr(*, "class")= chr [1:3] "op_join" "op_double" "op"
##   .. ..$ dots:List of 1
##   .. ..$ args: list()
##   .. ..- attr(*, "class")= chr [1:3] "op_rename" "op_single" "op"
##   ..$ dots:List of 3
##   .. ..$ : language ~rental_date
##   .. .. ..- attr(*, ".Environment")=<environment: 0x7ff56ad81e58>
##   .. ..$ : language ~staff_email
##   .. .. ..- attr(*, ".Environment")=<environment: 0x7ff56ad81e58>
##   .. ..$ : language ~customer_email
##   .. .. ..- attr(*, ".Environment")=<environment: 0x7ff56ad81e58>
##   .. ..- attr(*, "class")= chr "quosures"
##   ..$ args: list()
##   ..- attr(*, "class")= chr [1:3] "op_select" "op_single" "op"
##  - attr(*, "class")= chr [1:4] "tbl_dbi" "tbl_sql" "tbl_lazy" "tbl"
```

11.4.8   Q %>% nrow()

Notice the difference between `nrow()` and `tally()`. The `nrow` functions returns `NA` and does not execute a query:

```
Q %>% nrow()
```

```
## [1] NA
```

11.4.9   Q %>% tally()

The `tally` function actually counts all the rows.

```
Q %>% tally()
```

```
## # Source:    lazy query [?? x 1]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##    n
##    <S3: integer64>
## 1 16044
```

The `nrow()` function knows that `Q` is a list. On the other hand, the `tally()` function tells SQL to go count all the rows. Notice that `Q` results in 16,044 rows – the same number of rows as `rental`.

11.4.10 Q %>% collect()

The `dplyr::collect()` function triggers a dbFetch() function behind the scenes, which forces R to download a specified number of rows:

```
Q %>% collect(n = 20)
```

```
## # A tibble: 20 x 3
##    rental_date         staff_email          customer_email
##    <dttm>              <chr>                <chr>
##  1 2005-05-24 22:54:33 Mike.Hillyer@sakilasta~ tommy.collazo@sakilacustome~
##  2 2005-05-24 23:03:39 Mike.Hillyer@sakilasta~ manuel.murrell@sakilacustom~
##  3 2005-05-24 23:04:41 Jon.Stephens@sakilasta~ andrew.purdy@sakilacustomer~
##  4 2005-05-24 23:05:21 Mike.Hillyer@sakilasta~ delores.hansen@sakilacustom~
##  5 2005-05-24 23:08:07 Mike.Hillyer@sakilasta~ nelson.christenson@sakilacu~
##  6 2005-05-24 23:11:53 Jon.Stephens@sakilasta~ cassandra.walters@sakilacus~
##  7 2005-05-24 23:31:46 Jon.Stephens@sakilasta~ minnie.romero@sakilacustome~
##  8 2005-05-25 00:00:40 Mike.Hillyer@sakilasta~ ellen.simpson@sakilacustome~
##  9 2005-05-25 00:02:21 Jon.Stephens@sakilasta~ danny.isom@sakilacustomer.o~
## 10 2005-05-25 00:09:02 Jon.Stephens@sakilasta~ april.burns@sakilacustomer.~
## 11 2005-05-25 00:19:27 Jon.Stephens@sakilasta~ deanna.byrd@sakilacustomer.~
## 12 2005-05-25 00:22:55 Mike.Hillyer@sakilasta~ raymond.mcwhorter@sakilacus~
## 13 2005-05-25 00:31:15 Mike.Hillyer@sakilasta~ theodore.culp@sakilacustome~
## 14 2005-05-25 00:39:22 Mike.Hillyer@sakilasta~ ronald.weiner@sakilacustome~
## 15 2005-05-25 00:43:11 Jon.Stephens@sakilasta~ steven.curley@sakilacustome~
## 16 2005-05-25 01:06:36 Mike.Hillyer@sakilasta~ isaac.oglesby@sakilacustome~
## 17 2005-05-25 01:10:47 Jon.Stephens@sakilasta~ ruth.martinez@sakilacustome~
## 18 2005-05-25 01:17:24 Mike.Hillyer@sakilasta~ ronnie.ricketts@sakilacusto~
## 19 2005-05-25 01:48:41 Jon.Stephens@sakilasta~ roberta.harper@sakilacustom~
## 20 2005-05-25 01:59:46 Jon.Stephens@sakilasta~ craig.morrell@sakilacustome~
```

```
Q %>% collect(n = 20) %>% head()
```

```
## # A tibble: 6 x 3
##   rental_date         staff_email          customer_email
##   <dttm>              <chr>                <chr>
## 1 2005-05-24 22:54:33 Mike.Hillyer@sakilasta~ tommy.collazo@sakilacustomer~
## 2 2005-05-24 23:03:39 Mike.Hillyer@sakilasta~ manuel.murrell@sakilacustome~
## 3 2005-05-24 23:04:41 Jon.Stephens@sakilasta~ andrew.purdy@sakilacustomer.~
## 4 2005-05-24 23:05:21 Mike.Hillyer@sakilasta~ delores.hansen@sakilacustome~
## 5 2005-05-24 23:08:07 Mike.Hillyer@sakilasta~ nelson.christenson@sakilacus~
## 6 2005-05-24 23:11:53 Jon.Stephens@sakilasta~ cassandra.walters@sakilacust~
```

The `collect` function triggers the creation of a tibble and controls the number of rows that the DBMS sends

to R. Notice that `head` only prints 6 of the 25 rows that R has retrieved.

11.4.11   Q %>% show_query()

```
Q %>% show_query()
```

```
## <SQL>
## SELECT "rental_date", "staff_email", "customer_email"
## FROM (SELECT "rental_id", "rental_date", "inventory_id", "customer_id", "return_date", "staff_id", "las
## FROM (SELECT "TBL_LEFT"."rental_id" AS "rental_id", "TBL_LEFT"."rental_date" AS "rental_date", "TBL_LEF
##  FROM (SELECT "rental_id", "rental_date", "inventory_id", "customer_id", "return_date", "staff_id", "la
## FROM (SELECT "TBL_LEFT"."rental_id" AS "rental_id", "TBL_LEFT"."rental_date" AS "rental_date", "TBL_LEF
##   FROM "rental" AS "TBL_LEFT"
##   LEFT JOIN "staff" AS "TBL_RIGHT"
##   ON ("TBL_LEFT"."staff_id" = "TBL_RIGHT"."staff_id")
## ) "qzdwvvvtzq") "TBL_LEFT"
##   LEFT JOIN "customer" AS "TBL_RIGHT"
##   ON ("TBL_LEFT"."customer_id" = "TBL_RIGHT"."customer_id")
## ) "rnitnthkfz") "ohetoyqsmw"
```

Hand-written SQL code to do the same job will probably look a lot nicer and could be more efficient, but functionally dplyr does the job.

11.4.12   Qc <- Q %>% count(customer_email)

 Until `Q` is executed, we can add to it. This behavior is the basis for a useful debugging and development process where queries are built up incrementally.

```
Qc <- Q %>% count(customer_email, sort = TRUE)
```

 When all the accumulated `dplyr` verbs are executed, they are submitted to the dbms and the number of rows that are returned follow the same rules as discussed above.

```
Qc
```

```
## # Source:     lazy query [?? x 2]
## # Database:   postgres [postgres@localhost:5432/dvdrental]
## # Ordered by: desc(n)
##    customer_email                 n
##    <chr>                          <S3: integer64>
##  1 eleanor.hunt@sakilacustomer.org   46
##  2 karl.seal@sakilacustomer.org      45
##  3 clara.shaw@sakilacustomer.org     42
##  4 marcia.dean@sakilacustomer.org    42
##  5 tammy.sanders@sakilacustomer.org  41
##  6 wesley.bull@sakilacustomer.org    40
##  7 sue.peters@sakilacustomer.org     40
##  8 tim.cary@sakilacustomer.org       39
##  9 rhonda.kennedy@sakilacustomer.org 39
## 10 marion.snyder@sakilacustomer.org  39
```

```
## # ... with more rows
```

See more example of lazy execution can be found Here.

```
dbDisconnect(con)
sp_docker_stop("sql-pet")
```

11.5   Other resources

- Benjamin S. Baumer, A Grammar for Reproducible and Painless Extract-Transform-Load Operations on Medium Data: https://arxiv.org/pdf/1708.07073

Chapter 12

DBI and SQL (11c)

12.1 This chapter:

- Introduces more DBI functions and demonstrates techniques for submitting SQL to the dbms
- Illustrates some of the differences between writing `dplyr` commands and SQL
- Suggests some strategies for dividing the work between your local R session and the dbms

12.1.1 Setup

The following packages are used in this chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
library(dbplyr)
require(knitr)
library(bookdown)
library(sqlpetr)
```

Assume that the Docker container with PostgreSQL and the dvdrental database are ready to go. If not go back to the previous Chapter

```
sp_docker_start("sql-pet")
```

Connect to the database:

```
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

12.2 SQL in R Markdown

When you create a report to run repeatedly, you might want to put that query into R markdown. See the discussion of multiple language engines in R Markdown. That way you can also execute that SQL code in a

chunk with the following header:

```
{sql, connection=con, output.var = "query_results"}
```

```sql
SELECT "staff_id", COUNT(*) AS "n"
FROM "rental"
GROUP BY "staff_id";
```

Rmarkdown stored that query result in a tibble:

```
query_results
```

```
##   staff_id    n
## 1        2 8004
## 2        1 8040
```

12.3   DBI Package

In this chapter we touched on a number of functions from the DBI Package. The table in file 96b shows
other functions in the package. The Chapter column references a section in the book if we have used it.

```r
film_table <- tbl(con, "film")
```

12.3.1   Retrieve the whole table

SQL code that is submitted to a database is evaluated all at once[1]. To think through an SQL query, either
use dplyr to build it up step by step and then convert it to SQL code or an IDE such as pgAdmin. DBI
returns a data.frame, so you don't have dplyr's guardrails.

```r
res <- dbSendQuery(con, 'SELECT "title", "rental_duration", "length"
FROM "film"
WHERE ("rental_duration" > 5.0 AND "length" > 117.0)')

res_output <- dbFetch(res)
str(res_output)
```

```
## 'data.frame':    202 obs. of  3 variables:
## $ title         : chr  "African Egg" "Alamo Videotape" "Alaska Phantom" "Alley Evolution" ...
##  $ rental_duration: int  6 6 6 6 6 7 6 7 6 6 ...
##  $ length         : int  130 126 136 180 181 179 119 127 170 162 ...
```

```r
dbClearResult(res)
```

12.3.2   Or a chunk at a time

```r
res <- dbSendQuery(con, 'SELECT "title", "rental_duration", "length"
FROM "film"
WHERE ("rental_duration" > 5.0 AND "length" > 117.0)')

set.seed(5432)

chunk_num <- 0
while (!dbHasCompleted(res)) {
```

---

[1]From R's perspective. Actually there are 4 steps behind the scenes.

```r
  chunk_num <- chunk_num + 1
  chunk <- dbFetch(res, n = sample(7:13,1))
  # print(nrow(chunk))
  chunk$chunk_num <- chunk_num
  if (!chunk_num %% 9) {print(chunk)}
}
```

```
##                    title rental_duration length chunk_num
## 1          Grinch Massage               7    150         9
## 2         Groundhog Uncut               6    139         9
## 3           Half Outfield               6    146         9
## 4           Hamlet Wisdom               7    146         9
## 5           Harold French               6    168         9
## 6            Hedwig Alter               7    169         9
## 7         Holes Brannigan               7    128         9
## 8         Hollow Jeopardy               7    136         9
## 9      Holocaust Highball               6    149         9
## 10              Home Pity               7    185         9
## 11          Homicide Peach              6    141         9
## 12         Hotel Happiness              6    181         9
##                      title rental_duration length chunk_num
## 1           Towers Hurricane             7    144        18
## 2                   Town Ark             6    136        18
## 3          Trading Pinocchio             6    170        18
## 4      Trainspotting Strangers          7    132        18
## 5             Uncut Suicides             7    172        18
## 6         Unforgiven Zoolander           7    129        18
## 7            Uprising Uptown             6    174        18
## 8                Vanilla Day             7    122        18
## 9             Vietnam Smoochy            7    174        18
```

```r
dbClearResult(res)
```

12.4 Dividing the work between R on your machine and the DBMS

They work together.

12.4.1 Make the server do as much work as you can

- show_query as a first draft of SQL. May or may not use SQL code submitted directly.

12.4.2 Criteria for choosing between `dplyr` and native SQL

This probably belongs later in the book.

- performance considerations: first get the right data, then worry about performance
- Trade offs between leaving the data in PostgreSQL vs what's kept in R:
  - browsing the data
  - larger samples and complete tables
  - using what you know to write efficient queries that do most of the work on the server

Where you place the `collect` function matters. Here is a typical string of dplyr verbs strung together with the magrittr `%>%` command that will be used to tease out the several different behaviors that a lazy query has when passed to different R functions. This query joins three connection objects into a query we'll call `Q`:

```r
rental_table <- dplyr::tbl(con, "rental")
staff_table <- dplyr::tbl(con, "staff")
# the 'staff' table has 2 rows
customer_table <- dplyr::tbl(con, "customer")
# the 'customer' table has 599 rows

Q <- rental_table %>%
  left_join(staff_table, by = c("staff_id" = "staff_id")) %>%
  rename(staff_email = email) %>%
  left_join(customer_table, by = c("customer_id" = "customer_id")) %>%
  rename(customer_email = email) %>%
  select(rental_date, staff_email, customer_email)
```

```r
Q %>% show_query()
```

```
## <SQL>
## SELECT "rental_date", "staff_email", "customer_email"
## FROM (SELECT "rental_id", "rental_date", "inventory_id", "customer_id", "return_date", "staff_id", "last
## FROM (SELECT "TBL_LEFT"."rental_id" AS "rental_id", "TBL_LEFT"."rental_date" AS "rental_date", "TBL_LEF
##  FROM (SELECT "rental_id", "rental_date", "inventory_id", "customer_id", "return_date", "staff_id", "la
## FROM (SELECT "TBL_LEFT"."rental_id" AS "rental_id", "TBL_LEFT"."rental_date" AS "rental_date", "TBL_LEF
##   FROM "rental" AS "TBL_LEFT"
##   LEFT JOIN "staff" AS "TBL_RIGHT"
##   ON ("TBL_LEFT"."staff_id" = "TBL_RIGHT"."staff_id")
## ) "tvnvuviyiw") "TBL_LEFT"
##   LEFT JOIN "customer" AS "TBL_RIGHT"
##   ON ("TBL_LEFT"."customer_id" = "TBL_RIGHT"."customer_id")
## ) "dkimtwhtoo") "dkadgsqpgd"
```

Here is the SQL query formatted for readability:

```sql
SELECT "rental_date",
       "staff_email",
       "customer_email"
FROM   (SELECT "rental_id",
               "rental_date",
               "inventory_id",
               "customer_id",
               "return_date",
               "staff_id",
               "last_update.x",
               "first_name.x",
               "last_name.x",
               "address_id.x",
               "staff_email",
               "store_id.x",
               "active.x",
               "username",
               "password",
               "last_update.y",
               "picture",
               "store_id.y",
```

```
            "first_name.y",
            "last_name.y",
            "email" AS "customer_email",
            "address_id.y",
            "activebool",
            "create_date",
            "last_update",
            "active.y"
    FROM    (SELECT "TBL_LEFT"."rental_id"     AS "rental_id",
                    "TBL_LEFT"."rental_date"   AS "rental_date",
                    "TBL_LEFT"."inventory_id"  AS "inventory_id",
                    "TBL_LEFT"."customer_id"   AS "customer_id",
                    "TBL_LEFT"."return_date"   AS "return_date",
                    "TBL_LEFT"."staff_id"      AS "staff_id",
                    "TBL_LEFT"."last_update.x" AS "last_update.x",
                    "TBL_LEFT"."first_name"    AS "first_name.x",
                    "TBL_LEFT"."last_name"     AS "last_name.x",
                    "TBL_LEFT"."address_id"    AS "address_id.x",
                    "TBL_LEFT"."staff_email"   AS "staff_email",
                    "TBL_LEFT"."store_id"      AS "store_id.x",
                    "TBL_LEFT"."active"        AS "active.x",
                    "TBL_LEFT"."username"      AS "username",
                    "TBL_LEFT"."password"      AS "password",
                    "TBL_LEFT"."last_update.y" AS "last_update.y",
                    "TBL_LEFT"."picture"       AS "picture",
                    "TBL_RIGHT"."store_id"     AS "store_id.y",
                    "TBL_RIGHT"."first_name"   AS "first_name.y",
                    "TBL_RIGHT"."last_name"    AS "last_name.y",
                    "TBL_RIGHT"."email"        AS "email",
                    "TBL_RIGHT"."address_id"   AS "address_id.y",
                    "TBL_RIGHT"."activebool"   AS "activebool",
                    "TBL_RIGHT"."create_date"  AS "create_date",
                    "TBL_RIGHT"."last_update"  AS "last_update",
                    "TBL_RIGHT"."active"       AS "active.y"
             FROM   (SELECT "rental_id",
                            "rental_date",
                            "inventory_id",
                            "customer_id",
                            "return_date",
                            "staff_id",
                            "last_update.x",
                            "first_name",
                            "last_name",
                            "address_id",
                            "email" AS "staff_email",
                            "store_id",
                            "active",
                            "username",
                            "password",
                            "last_update.y",
                            "picture"
                     FROM   (SELECT "TBL_LEFT"."rental_id"    AS "rental_id",
                                    "TBL_LEFT"."rental_date"  AS
                                    "rental_date",
```

```
                                         "TBL_LEFT"."inventory_id" AS
                                         "inventory_id",
                                         "TBL_LEFT"."customer_id"  AS
                                         "customer_id",
                                         "TBL_LEFT"."return_date"  AS
                                         "return_date",
                                         "TBL_LEFT"."staff_id"     AS "staff_id",
                                         "TBL_LEFT"."last_update"  AS
                                         "last_update.x",
                                         "TBL_RIGHT"."first_name"  AS "first_name"
                                         ,
                        "TBL_RIGHT"."last_name"   AS "last_name",
                        "TBL_RIGHT"."address_id"  AS "address_id",
                        "TBL_RIGHT"."email"       AS "email",
                        "TBL_RIGHT"."store_id"    AS "store_id",
                        "TBL_RIGHT"."active"      AS "active",
                        "TBL_RIGHT"."username"    AS "username",
                        "TBL_RIGHT"."password"    AS "password",
                        "TBL_RIGHT"."last_update" AS "last_update.y",
                        "TBL_RIGHT"."picture"     AS "picture"
                               FROM   "rental" AS "TBL_LEFT"
                                   LEFT JOIN "staff" AS "TBL_RIGHT"
                                          ON ( "TBL_LEFT"."staff_id" =
                                               "TBL_RIGHT"."staff_id" ))
                            "ymdofxkiex") "TBL_LEFT"
                  LEFT JOIN "customer" AS "TBL_RIGHT"
                         ON ( "TBL_LEFT"."customer_id" =
                             "TBL_RIGHT"."customer_id" ))
             "exddcnhait") "aohfdiedlb"
```

Hand-written SQL code to do the same job will probably look a lot nicer and could be more efficient, but functionally dplyr does the job.

```
GQ <- dbGetQuery(
  con,
  "select r.rental_date, s.email staff_email,c.email customer_email
     from rental r
         left outer join staff s on r.staff_id = s.staff_id
         left outer join customer c on r.customer_id = c.customer_id
  "
)
```

But because `Q` hasn't been executed, we can add to it. This behavior is the basis for a useful debugging and development process where queries are built up incrementally.

Where you place the `collect` function matters.

```
dbDisconnect(con)
sp_docker_stop("sql-pet")
```

Chapter 13

Joins and complex queries (13)

This chapter demonstrates how to:

- Use prinmary and foreign keys to retrieve specific rows of a table
- do different kinds of queries
- Exercises
- Query the database to get basic information about each dvdrental story
- How to interact with the database using different strategies

Verify Docker is up and running:

```
sp_check_that_docker_is_up()
```

```
## [1] "Docker is up but running no containers"
```

verify pet DB is available, it may be stopped.

```
sp_show_all_docker_containers()
```

```
## CONTAINER ID     IMAGE              COMMAND               CREATED          STATUS             PORTS
## 5802b471a9cf     postgres-dvdrental  "docker-entrypoint.s…"  33 seconds ago   Exited (0) 2 seconds ag
```

Start up the `docker-pet` container

```
sp_docker_start("sql-pet")
```

now connect to the database with R

```
# need to wait for Docker & Postgres to come up before connecting.

con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

13.1   Database constraints

As a data analyst, you really do not have to worry about database constraints since you are primarily writing dplyr/SQL queries to pull data out of the database. Constraints can be enforced at multiple levels, column, table, multiple tables, or at the schema itself.

For this tutorial, we are primarily concerned with primary and foreign key constraints. If one looks at all the tables in the DVD Rental ERD, the first column is the name of the table followed by "id". This is the primary key on the table. In some of the tables, there are other columns that begin with the name of a different table, the foreign table, and end in "_id". These are foreign keys and the foreign key value is the primary key value on the foreign table. The DBA will index the primary and foreign key columns to speed up query performanace.

13.2   Making up data for Join Examples

13.2.1   insert yourself as a new customer

```
# Customer 600 should be the next customer.
# It gets deleted here just in case it was added in a different session.
dbExecute(
  con,
  "delete from customer
   where customer_id = 600;
   "
)
```

```
## [1] 0
```

```
# Now add yourself as the next customer.  Replace Sophie Yang with your name.
dbExecute(
  con,
  "insert into customer
  (customer_id,store_id,first_name,last_name,email,address_id
  ,activebool,create_date,last_update,active)
  values(600,2,'Sophie','Yang','email@email.com',1,TRUE,now()::date,now()::date,1)
  ;
  "
)
```

```
## [1] 1
```

The `film` table has a primary key, film_id, and a foreign key column, language_id. One cannot insert a new row into the film table with a language_id = 10 because of a constraint on the language_id column. The language_id value must already exist in the `language` table before the database will allow the new row to be inserted into the table.

To work around this inconvenience for the tutorial:

1. we drop the smy_film table if it exists from a previous session.

```
dbExecute(con, "drop table if exists smy_film;")
```

```
## [1] 0
```

2. we create a new table smy_film from the film table and add a new row with a language_id = 10;

```
dbExecute(con, "create table smy_film as select * from film;")
```

```
## [1] 1000
```

3. We create a film with language_id = 10;

```
dbExecute(
  con,
```

```
"insert into smy_film
(film_id,title,description,release_year,language_id
,rental_duration,rental_rate,length,replacement_cost,rating
 ,last_update,special_features,fulltext)
values(1001,'Sophie''s Choice','orphaned language_id=10',2018,10
      ,7,4.99,120,14.99,'PG'
      ,now()::date,'{Trailers}','')
;
"
)
```

[1] 1

4. Confirm that the new record exists.

```
dbGetQuery(
  con,
  "select film_id,title,description,language_id from smy_film where film_id = 1001;"
)
```

```
##   film_id          title               description language_id
## 1    1001 Sophie's Choice orphaned language_id=10          10
```

13.3 Joins

In section 'SQL Quick Start Simple Retrieval', there is a brief discussion of databases and 3NF. The goal of normalization is to push the data into separate tables at a very granular level.

Bill Kent famously summarized 3NF as every non-key column "must provide a fact about the key,the whole key, and nothing but the key, so help me Codd."

Normalization breaks data down and JOINs denormalizes the data and builds it back up.

The above diagram can be found here There are additional graphics at the link, but the explanations are poorly worded and hard to follow.

The diagram above shows nicely the hierarchy of different types of joins. For this tutorial, we can think of joins as either an Inner Join or an Outer Join.

Instead of showing standard Venn diagrams showing the different JOINS, we use an analogy. For those interested though, the typical Venn diagrams can be found here.

13.3.1 Valentines Party

Imagine you are at a large costume Valentine's Day dance party. The hostess of the party, a data scientist, would like to learn more about the people attending her party. She interrupts the music to let everyone know it is time for the judges to evaluate the winners for best costumes and associated prizes.

She requests the following:

1. All the couples at the party to line up in front of her with the men on the left and the women on the right, (inner join)

2. All the remaining men to form a second line two feet behind the married men,(left outer join)

3. Right Outer Join: All the remaining women to form a third line two feet in front of the married women, (right outer join, all couples + unattached women)

Figure 13.1: SQL_JOIN_TYPES

Full Outer Join – As our data scientist looks out at the three lines, she can clearly see the three distinct lines, her full outer join.

As the three judges start walking down the lines, she makes one more announcement.

4. There is a special prize for the man and woman who can guess the average age of the members of the opposite sex. To give everyone a chance to come up with an average age, she asks the men to stay in line and the women to move down the mens line in order circling back around until they get back to their starting point in line, (full outer join, every man seen by every woman and vice versa).

It is hard enough to tell someone's age when they don't have a mask, how do you get the average age when people have masks?

The hostess knows that there is usually some data anomolies. As she looks out she sees a small cluster of people who did not line up. Being the hostess with the mostest, she wants to get to know that small cluster better. Since they are far off and in costume, she cannot tell if they are men or women. More importantly, she does not know if they identify as a man or a woman, both – (kind of a stretch for a self join), neither, or something else. Ahh, the inquisitive mind wants to know.

13.3.2   Join Syntax

Join	dplyr	sql
inner	inner_join(customer_tbl, rental_tbl, by = 'customer_id', suffix = c(".c", ".r")) customer_tbl %>% inner_join(rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	from customer c join rental r on c.customer_id = r.customer_id
left	left_join(customer_tbl, rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	from customer c left outer join rental r on c.customer_id = r.customer_id

Join	dplyr	sql
	customer_tbl %>% left_join(rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	
right	right_join(customer_tbl, rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	from customer c right outer join rental r on c.customer_id = r.customer_id
	customer_tbl %>% right_join(rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	
full	full_join(customer_tbl, rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	from customer c full outer join rental r on c.customer_id = r.customer_id
	customer_tbl %>% full_join(rental_tbl, by = 'customer_id', suffix = c(".c", ".r"))	

13.3.3 Join Tables

The dplyr join documentation describes two different types of joins, `mutating` and `filtering` joins. For those coming to R with a SQL background, the mutating documentation is misleading in one respect. Here is the inner_join documentation.

```
inner_join()
```

`return all rows from x where there are matching values in y, and all columns from x and y. If there are multipl`

The misleading part is that all the columns from *x* and *y*. If the join column is `KEY`, SQL will return x.KEY and y.KEY. Dplyr retuns KEY. It appears that the KEY value comes from the key/driving table. This difference should become clear in the outer join examples.

In the next couple of examples, we will pull all the language and smy_film table data from the database into memory because the tables are small. In the *_join verbs, the `by` and `suffix` parameters are included because it helps document the actual join and the source of join columns.

13.4 Natural Join Time Bomb

The dplyr default join is a natural join, joining tables on common column names. One of many links why one should not use natural joins can be found here.

13.5 Join Templates

In this section we look at two tables, `language` and `smy_film` and various joins using dplyr and SQL. Each dplyr code block has three purposes.

1. Show a working join example.

2. The code blocks can be used as templates for beginning more complex dplyr pipes.
3. The code blocks show the number of joins performed.

In these examples, the join condition, the `by` parameter,

```
by = c('language_id','language_id')
```

the two columns are the same. In multi-column joins, each language_id would be replace with a vector of column names used in the join by position. Note the column names do not need to be identical by position.

The suffix parameter is a way to distinguish the same column name in the joined tables. The suffixes are usually an single letter to represent the name of the table.

```
language_table <- DBI::dbReadTable(con, "language")
film_table <- DBI::dbReadTable(con, "smy_film")
```

13.5.1   dplyr Inner Join Template

```r
languages_ij <- language_table %>%
  inner_join(film_table, by = c("language_id", "language_id"), suffix(c(".l", ".f"))) %>%
  group_by(language_id, name) %>%
  summarize(inner_joins = n())

languages_ij
```

```
## # A tibble: 1 x 3
## # Groups:   language_id [?]
##   language_id name                inner_joins
##         <int> <chr>                     <int>
## 1           1 "English          "        1000
```

13.5.1.1   SQL Inner Join

```r
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,count(*) n
   from language l join smy_film f on l.language_id = f.language_id
  group by l.language_id,l.name;"
)

rs
```

```
##   language_id          name    n
## 1           1 English       1000
```

The output tells us that there are 1000 inner joins between the language_table and the film_table.

13.5.2   dplyr Left Outer Join Template

```r
languages_loj <- language_table %>%
  left_join(film_table, by = c("language_id", "language_id"), suffix(c(".l", ".f"))) %>%
  mutate(
    join_type = "loj"
    , film_lang_id = if_else(is.na(film_id), film_id, language_id)
  ) %>%
  group_by(join_type, language_id, name, film_lang_id) %>%
  summarize(lojs = n()) %>%
  select(join_type, language_id, film_lang_id, name, lojs)
print(languages_loj)
```

```
## # A tibble: 6 x 5
## # Groups:   join_type, language_id, name [6]
```

```
##    join_type language_id film_lang_id name                    lojs
##    <chr>           <int>        <int> <chr>                   <int>
## 1 loj                 1            1 "English           "      1000
## 2 loj                 2           NA "Italian           "         1
## 3 loj                 3           NA "Japanese          "         1
## 4 loj                 4           NA "Mandarin          "         1
## 5 loj                 5           NA "French            "         1
## 6 loj                 6           NA "German            "         1
# View(languages_loj)
# sp_print_df(languages_loj)
```

Compare the mutate verb in the above code block with film_lang_id in the equivalent SQL code block below.

13.5.2.1 SQL Left Outer Join

```
rs <- dbGetQuery(
  con,
  "select l.language_id
       ,f.language_id film_lang_id
       ,trim(l.name) as name
       ,count(*) lojs
   from language l left outer join smy_film f
       on l.language_id = f.language_id
  group by l.language_id,l.name,f.language_id
order by l.language_id;"
)
# sp_print_df(rs)
rs
```

```
##   language_id film_lang_id     name lojs
## 1           1            1  English 1000
## 2           2           NA   Italian    1
## 3           3           NA  Japanese    1
## 4           4           NA  Mandarin    1
## 5           5           NA    French    1
## 6           6           NA    German    1
```

The lojs column returns the number of rows found on the keys from the left table, language, and the right table, the film table. For the "English" row, the language_id and film_lang_id match and a 1000 inner joins were performed. For all the other languages, there was only 1 join and they all came from the left outer table, the language table, language_id's 2 - 6. The right table, the film table returned NA, because no match was found.

1. The left outer join always returns all rows from the left table, the driving/key table, if not reduced via a filter()/where clause.

2. All rows that inner join returns all the columns/derived columns specified in the select clause from both the left and right tables.

3. All rows from the left table, the outer table, without a matching row on the right returns all the columns/derived column values specified in the select clause from the left, but the values from right table have all values of NA.

13.5.2.2   dplyr Right Outer Join

```
languages_roj <- language_table %>%
  right_join(film_table, by = c("language_id", "language_id"), suffix(c(".l", ".f")), all = film_table)
  mutate(
    lang_id = if_else(is.na(name), 0L, language_id)
    , join_type = "rojs"
  ) %>%
  group_by(join_type, language_id, name, lang_id) %>%
  summarize(rojs = n()) %>%
  select(join_type, lang_id, language_id, name, rojs)

sp_print_df(languages_roj)
```

join_type	lang_id	language_id	name	rojs
rojs	1	1	English	1000
rojs	0	10	NA	1

```
languages_roj
```

```
## # A tibble: 2 x 5
## # Groups:   join_type, language_id, name [2]
##   join_type lang_id language_id name                          rojs
##   <chr>       <int>       <int> <chr>                        <int>
## 1 rojs            1           1 "English                "     1000
## 2 rojs            0          10 <NA>                              1
```

Review the mutate above with l.language_id below.

13.5.2.3   SQL Right Outer Join

```
rs <- dbGetQuery(
  con,
  "select 'roj' join_type,l.language_id,f.language_id language_id_f,l.name,count(*) rojs
   from language l right outer join smy_film f on l.language_id = f.language_id
  group by l.language_id,l.name,f.language_id
order by l.language_id;"
)
sp_print_df(rs)
```

join_type	language_id	language_id_f	name	rojs
roj	1	1	English	1000
roj	NA	10	NA	1

```
rs
```

```
##   join_type language_id language_id_f              name rojs
## 1       roj           1             1 English             1000
## 2       roj          NA            10              <NA>     1
```

The rojs column returns the number of rows found on the keys from the right table, film, and the left table, the language table. For the "English" row, the language_id and film_lang_id match and a 1000 inner joins were performed. For language_id = 10 from the right table, there was only 1 join to a non-existant row in the language table on the left.

1. The right outer join always returns all rows from the right table, the driving/key table, if not reduced via a filter()/where clause.

2. All rows that inner join returns all the columns/derived columns specified in the select clause from both the left and right tables.

3. All rows from the right table, the outer table, without a matching row on the left returns all the columns/derived column values specified in the select clause from the right, but the values from left table have all values of NA.

13.5.2.4 dplyr Full Outer Join

```
languages_foj <- language_table %>%
  full_join(film_table, by = c("language_id", "language_id"), suffix(c(".l", ".f"))) %>%
  mutate(film_lang = if_else(is.na(film_id), paste0("No ", name, " films."), if_else(is.na(name), "Alien
  group_by(language_id, name, film_lang) %>%
  summarize(n = n())

sp_print_df(languages_foj)
```

language_id	name	film_lang	n
1	English	English	1000
2	Italian	No Italian films.	1
3	Japanese	No Japanese films.	1
4	Mandarin	No Mandarin films.	1
5	French	No French films.	1
6	German	No German films.	1
10	NA	Alien	1

```
languages_foj
```

```
## # A tibble: 7 x 4
## # Groups:   language_id, name [?]
##    language_id name                film_lang                        n
##          <int> <chr>               <chr>                        <int>
## 1            1 "English           " "English            "        1000
## 2            2 "Italian           " No Italian          films.      1
## 3            3 "Japanese          " No Japanese         films.      1
## 4            4 "Mandarin          " No Mandarin         films.      1
## 5            5 "French            " No French           films.      1
## 6            6 "German            " No German           films.      1
## 7           10 <NA>                Alien                           1
```

13.5.2.5 SQL full Outer Join

```
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,f.language_id language_id_f,count(*) fojs
   from language l full outer join smy_film f on l.language_id = f.language_id
  group by l.language_id,l.name,f.language_id
order by l.language_id;"
)
sp_print_df(rs)
```

language_id	name	language_id_f	fojs
1	English	1	1000
2	Italian	NA	1
3	Japanese	NA	1
4	Mandarin	NA	1
5	French	NA	1
6	German	NA	1
NA	NA	10	1

rs

```
##   language_id                name language_id_f fojs
## 1           1 English                         1 1000
## 2           2 Italian                        NA    1
## 3           3 Japanese                       NA    1
## 4           4 Mandarin                       NA    1
## 5           5 French                         NA    1
## 6           6 German                         NA    1
## 7          NA                 <NA>           10    1
```

Looking at the SQL output, the full outer join is the combination of the left and right outer joins.

1. Language_id = 1 is the inner join.
2. Language_id = 2 - 6 is the left outer join
3. Language_id = 10 is the right outer join.

One can also just look at the language_id on the left and language_id_f on the right for a non NA value to see which side is outer side/driving side of the join.

13.5.2.6   dplyr anti Join

The anti join is a left outer join without the inner joined rows. It only returns the rows from the left table that do not have a match from the right table.

```
languages_aj <- language_table %>%
  anti_join(film_table, by = c("language_id", "language_id"), suffix(c(".l", ".f"))) %>%
  mutate(type = "anti_join") %>%
  group_by(type, language_id, name) %>%
  summarize(anti_joins = n()) %>%
  select(type, language_id, name, anti_joins)
sp_print_df(languages_aj)
```

type	language_id	name	anti_joins
anti_join	2	Italian	1
anti_join	3	Japanese	1
anti_join	4	Mandarin	1
anti_join	5	French	1
anti_join	6	German	1

languages_aj

```
## # A tibble: 5 x 4
## # Groups:   type, language_id [5]
##   type      language_id name                anti_joins
##   <chr>           <int> <chr>                    <int>
## 1 anti_join           2 "Italian           "        1
## 2 anti_join           3 "Japanese          "        1
## 3 anti_join           4 "Mandarin          "        1
```

```
## 4 anti_join          5 "French              "        1
## 5 anti_join          6 "German              "        1
```

13.5.2.7 SQL anti Join 1, Left Outer Join where NULL on Right

SQL doesn't have an anti join key word. Here are three different ways to achieve the same result.

```
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,count(*) fojs
   from language l left outer join smy_film f on l.language_id = f.language_id
  where f.language_id is null
  group by l.language_id,l.name
order by l.language_id;"
)
sp_print_df(rs)
```

language_id	name	fojs
2	Italian	1
3	Japanese	1
4	Mandarin	1
5	French	1
6	German	1

rs

```
##   language_id           name fojs
## 1           2 Italian           1
## 2           3 Japanese          1
## 3           4 Mandarin          1
## 4           5 French            1
## 5           6 German            1
```

13.5.2.8 SQL anti Join 2, ID in driving table and NOT IN lookup table

```
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,count(*) fojs
   from language l
  where l.language_id NOT IN (select language_id from film)
  group by l.language_id,l.name
order by l.language_id;"
)
sp_print_df(rs)
```

language_id	name	fojs
2	Italian	1
3	Japanese	1
4	Mandarin	1
5	French	1
6	German	1

rs

```
##   language_id           name fojs
## 1           2 Italian           1
```

```
## 2           3 Japanese           1
## 3           4 Mandarin           1
## 4           5 French             1
## 5           6 German             1
```

13.5.2.9   SQL anti Join 3, NOT EXISTS and Correlated subquery

```
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,count(*) fojs
   from language l
  where not exists (select language_id from film f where f.language_id = l.language_id)
 group by l.language_id,l.name
"
)
sp_print_df(rs)
```

language_id	name	fojs
2	Italian	1
3	Japanese	1
4	Mandarin	1
5	French	1
6	German	1

```
rs
```

```
##   language_id            name fojs
## 1           2 Italian           1
## 2           3 Japanese          1
## 3           4 Mandarin          1
## 4           5 French            1
## 5           6 German            1
```

13.6   SQL anti join Costs

```
sql_aj1 <- dbGetQuery(
  con,
  "explain analyze select l.language_id,l.name,count(*) fojs
   from language l left outer join smy_film f on l.language_id = f.language_id
  where f.language_id is null
  group by l.language_id,l.name
"
)

sql_aj2 <- dbGetQuery(
  con,
  "explain analyze select l.language_id,l.name,count(*) fojs
   from language l
  where l.language_id NOT IN (select language_id from film)
  group by l.language_id,l.name
"
)
```

```
sql_aj3 <- dbGetQuery(
  con,
  "explain analyze select l.language_id,l.name,count(*) fojs
   from language l
  where not exists (select language_id from film f where f.language_id = l.language_id)
 group by l.language_id,l.name
"
)
```

13.6.0.0.1 SQL Costs

```
print(glue("sql_aj1 loj-null costs=", sql_aj1[1, 1]))
```

sql_aj1 loj-null costs=GroupAggregate (cost=68.56..68.61 rows=3 width=96) (actual time=20.655..20.790 r

```
print(glue("sql_aj2 not in costs=", sql_aj2[1, 1]))
```

sql_aj2 not in costs=GroupAggregate (cost=67.60..67.65 rows=3 width=96) (actual time=23.405..23.522 row

```
print(glue("sql_aj3 not exist costs=", sql_aj3[1, 1]))
```

sql_aj3 not exist costs=GroupAggregate (cost=24.24..24.30 rows=3 width=96) (actual time=0.430..0.543 ro

13.7 dplyr Anti joins

In this next section we look at two methods to implemnt an anti join in dplyr.

```
customer_table <- tbl(con, "customer") # DBI::dbReadTable(con, "customer")
rental_table <- tbl(con, "rental") # DBI::dbReadTable(con, "rental")

# Method 1.  dplyr anti_join
daj1 <-
  anti_join(customer_table, rental_table, by = "customer_id", suffix = c(".c", ".r")) %>%
  select(c("first_name", "last_name", "email")) %>%
  explain()
```

```
## <SQL>
## SELECT "first_name", "last_name", "email"
## FROM (SELECT * FROM "customer" AS "TBL_LEFT"
##
## WHERE NOT EXISTS (
##   SELECT 1 FROM "rental" AS "TBL_RIGHT"
##   WHERE ("TBL_LEFT"."customer_id" = "TBL_RIGHT"."customer_id")
## )) "lghnetfgzu"
```

```
##
## <PLAN>
## Hash Anti Join  (cost=510.99..552.63 rows=300 width=334)
##   Hash Cond: ("TBL_LEFT".customer_id = "TBL_RIGHT".customer_id)
##   -> Seq Scan on customer "TBL_LEFT"  (cost=0.00..14.99 rows=599 width=338)
##   -> Hash  (cost=310.44..310.44 rows=16044 width=2)
##       -> Seq Scan on rental "TBL_RIGHT"  (cost=0.00..310.44 rows=16044 width=2)
```

```
customer_table <- tbl(con, "customer") # DBI::dbReadTable(con, "customer")
rental_table <- tbl(con, "rental") # DBI::dbReadTable(con, "rental")
```

```
# Method 2.  dplyr loj with NA
daj2 <-
  left_join(customer_table, rental_table, by = c("customer_id", "customer_id"), suffix = c(".c", ".r"))
  filter(is.na(rental_id)) %>%
  select(c("first_name", "last_name", "email")) %>%
  explain()
```

```
## <SQL>
## SELECT "first_name", "last_name", "email"
## FROM (SELECT "TBL_LEFT"."customer_id" AS "customer_id", "TBL_LEFT"."store_id" AS "store_id", "TBL_LEFT"
##   FROM "customer" AS "TBL_LEFT"
##   LEFT JOIN "rental" AS "TBL_RIGHT"
##   ON ("TBL_LEFT"."customer_id" = "TBL_RIGHT"."customer_id")
## ) "bisndnmtew"
## WHERE ((("rental_id") IS NULL))
##
## <PLAN>
## Hash Right Join  (cost=22.48..375.33 rows=80 width=334)
##   Hash Cond: ("TBL_RIGHT".customer_id = "TBL_LEFT".customer_id)
##   Filter: ("TBL_RIGHT".rental_id IS NULL)
##   ->  Seq Scan on rental "TBL_RIGHT"  (cost=0.00..310.44 rows=16044 width=6)
##   ->  Hash  (cost=14.99..14.99 rows=599 width=338)
##         ->  Seq Scan on customer "TBL_LEFT"  (cost=0.00..14.99 rows=599 width=338)
```

13.7.1 dplyr Costs

```
<PLAN>
Hash Anti Join  (cost=510.99..529.72 rows=1 width=45)
  Hash Cond: ("TBL_LEFT".customer_id = "TBL_RIGHT".customer_id)
  ->  Seq Scan on customer "TBL_LEFT"  (cost=0.00..14.99 rows=599 width=49)
  ->  Hash  (cost=310.44..310.44 rows=16044 width=2)
        ->  Seq Scan on rental "TBL_RIGHT"  (cost=0.00..310.44 rows=16044 width=2)

<PLAN>
Hash Right Join  (cost=22.48..375.33 rows=1 width=45)
  Hash Cond: ("TBL_RIGHT".customer_id = "TBL_LEFT".customer_id)
  Filter: ("TBL_RIGHT".rental_id IS NULL)
  ->  Seq Scan on rental "TBL_RIGHT"  (cost=0.00..310.44 rows=16044 width=6)
  ->  Hash  (cost=14.99..14.99 rows=599 width=49)
        ->  Seq Scan on customer "TBL_LEFT"  (cost=0.00..14.99 rows=599 width=49)
```

In this example, the dplyr anti_join verb is *1.4113447 to 22.7308719* times more expensive than the left outer join with a null condition.

```
sql_aj1 <- dbGetQuery(
  con,
  "explain analyze select c.customer_id,count(*) lojs
   from customer c left outer join rental r on c.customer_id = r.customer_id
  where r.customer_id is null
  group by c.customer_id
order by c.customer_id;"
)
sp_print_df(sql_aj1)
```

QUERY PLAN
GroupAggregate (cost=564.97..570.22 rows=300 width=12) (actual time=270.692..270.708 rows=1 loops=1)
Group Key: c.customer_id
-> Sort (cost=564.97..565.72 rows=300 width=4) (actual time=270.657..270.673 rows=1 loops=1)
Sort Key: c.customer_id
Sort Method: quicksort Memory: 25kB
-> Hash Anti Join (cost=510.99..552.63 rows=300 width=4) (actual time=270.597..270.629 rows=1 loops=1)
Hash Cond: (c.customer_id = r.customer_id)
-> Seq Scan on customer c (cost=0.00..14.99 rows=599 width=4) (actual time=0.140..5.128 rows=600 loops=1)
-> Hash (cost=310.44..310.44 rows=16044 width=2) (actual time=260.198..260.205 rows=16044 loops=1)
Buckets: 16384 Batches: 1 Memory Usage: 661kB
-> Seq Scan on rental r (cost=0.00..310.44 rows=16044 width=2) (actual time=0.144..129.586 rows=16044 loops=1)
Planning time: 0.361 ms
Execution time: 271.018 ms

```
sql_aj1
```

```
##                                                         QUERY PLAN
## 1                 GroupAggregate  (cost=564.97..570.22 rows=300 width=12) (actual time=270.692..27
## 2                                                        Group Key: c.customer_id
## 3                       -> Sort  (cost=564.97..565.72 rows=300 width=4) (actual time=270.657..270.6
## 4                                                        Sort Key: c.customer_id
## 5                                                        Sort Method: quicksort  Memory: 25kB
## 6                 -> Hash Anti Join  (cost=510.99..552.63 rows=300 width=4) (actual time=270.597..27
## 7                                                        Hash Cond: (c.customer_id = r.customer_i
## 8               -> Seq Scan on customer c  (cost=0.00..14.99 rows=599 width=4) (actual time=0.140..5.
## 9                  -> Hash  (cost=310.44..310.44 rows=16044 width=2) (actual time=260.198..260.205
## 10                                                       Buckets: 16384  Batches: 1  Memory Usage: 66
## 11            -> Seq Scan on rental r  (cost=0.00..310.44 rows=16044 width=2) (actual time=0.144..129.
## 12                                                       Planning time: 0.361 ms
## 13                                                       Execution time: 271.018 ms
```

```r
sql_aj3 <- dbGetQuery(
  con,
  "explain analyze
select c.customer_id,count(*) lojs
  from customer c
 where not exists (select customer_id from rental r where c.customer_id = r.customer_id)
 group by c.customer_id
"
)

print(glue("sql_aj1 loj-null costs=", sql_aj1[1, 1]))
```

```
## sql_aj1 loj-null costs=GroupAggregate  (cost=564.97..570.22 rows=300 width=12) (actual time=270.692..270
```

```r
print(glue("sql_aj3 not exist costs=", sql_aj3[1, 1]))
```

```
## sql_aj3 not exist costs=HashAggregate  (cost=554.13..557.13 rows=300 width=12) (actual time=287.783..287
```

13.8 Exercises

13.8.1 Anti joins – Find customers who have never rented a movie, take 2.

This is a left outer join from customer to the rental table with an NA rental_id.

13.8.1.1   SQL Anti-Join

```
rs <- dbGetQuery(
  con,
  "select c.first_name
                       ,c.last_name
                       ,c.email
                 from customer c
                     left outer join rental r
                          on c.customer_id = r.customer_id
                 where r.rental_id is null;
               "
)
sp_print_df(head(rs))
```

first_name	last_name	email
Sophie	Yang	email@email.com

<– Add dplyr semi-join example –>

13.8.2   SQL Rows Per Table

In the examples above, we looked at how many rows were involved in each of the join examples and which side of the join they came from. It is often helpful to know how many rows are in each table as a sanity check on the joins.

Below is the SQL version to return all the row counts from each table in the DVD Rental System.

```
rs <- dbGetQuery(
  con,
  "select *
              from (      select 'actor' tbl_name,count(*) from actor
                     union select 'category' tbl_name,count(*) from category
                     union select 'film' tbl_name,count(*) from film
                     union select 'film_actor' tbl_name,count(*) from film_actor
                     union select 'film_category' tbl_name,count(*) from film_category
                     union select 'language' tbl_name,count(*) from language
                     union select 'inventory' tbl_name,count(*) from inventory
                     union select 'rental' tbl_name,count(*) from rental
                     union select 'payment' tbl_name,count(*) from payment
                     union select 'staff' tbl_name,count(*) from staff
                     union select 'customer' tbl_name,count(*) from customer
                     union select 'address' tbl_name,count(*) from address
                     union select 'city' tbl_name,count(*) from city
                     union select 'country' tbl_name,count(*) from country
                     union select 'store' tbl_name,count(*) from store
                     union select 'smy_film' tbl_name,count(*) from smy_film
                     ) counts
               order by tbl_name
              ;
              "
)
sp_print_df(head(rs))
```

tbl_name	count
actor	200
address	603
category	16
city	600
country	109
customer	600

rs

```
##          tbl_name count
## 1           actor   200
## 2         address   603
## 3        category    16
## 4            city   600
## 5         country   109
## 6        customer   600
## 7            film  1000
## 8      film_actor  5462
## 9   film_category  1000
## 10      inventory  4581
## 11       language     6
## 12        payment 14596
## 13         rental 16044
## 14       smy_film  1001
## 15          staff     2
## 16          store     2
```

13.8.2.1 Exercise dplyr Rows Per Table

In the code block below

1. Get the row counts for a couple more tables
2. What is the structure of film_table object?

```r
film_table <- tbl(con, "film") # DBI::dbReadTable(con, "customer")
language_table <- tbl(con, "language") # DBI::dbReadTable(con, "rental")

film_rows <- film_table %>% mutate(name = "film") %>% group_by(name) %>% summarize(rows = n())
language_rows <- language_table %>%
  mutate(name = "language") %>%
  group_by(name) %>%
  summarize(rows = n())
rows_per_table <- rbind(as.data.frame(film_rows), as.data.frame(language_rows))
rows_per_table
```

```
##       name rows
## 1     film 1000
## 2 language    6
```

13.8.2.2 SQL film distribution based on language

The SQL below is very similar to the `SQL full Outer Join` above. Instead of counting the joins, it counts the number films associated with each language.

```r
rs <- dbGetQuery(
  con,
  "select l.language_id id
                  ,l.name
                  ,sum(case when f.language_id is not null then 1 else 0 end) total
           from language l
                  full outer join film f
                       on l.language_id = f.language_id
           group by l.language_id,l.name
           order by l.name;
           ;
          "
)
sp_print_df(head(rs))
```

id	name	total
1	English	1000
5	French	0
6	German	0
2	Italian	0
3	Japanese	0
4	Mandarin	0

```r
rs
```

```
##   id              name total
## 1  1 English            1000
## 2  5 French                0
## 3  6 German                0
## 4  2 Italian               0
## 5  3 Japanese              0
## 6  4 Mandarin              0
```

13.8.2.3   Exercise dplyr film distribution based on language

Below is the code block from the `dplyr Full Outer Join` section above. Modify the code block to match the output from the SQL version.

```r
rs <- dbGetQuery(
  con,
  "select l.language_id,l.name,f.language_id language_id_f,count(*) fojs
   from language l full outer join smy_film f on l.language_id = f.language_id
  group by l.language_id,l.name,f.language_id
order by l.language_id;"
)
sp_print_df(rs)
```

language_id	name	language_id_f	fojs
1	English	1	1000
2	Italian	NA	1
3	Japanese	NA	1
4	Mandarin	NA	1
5	French	NA	1
6	German	NA	1
NA	NA	10	1

```
rs
```

```
##   language_id               name language_id_f fojs
## 1           1 English                        1 1000
## 2           2 Italian                       NA    1
## 3           3 Japanese                      NA    1
## 4           4 Mandarin                      NA    1
## 5           5 French                        NA    1
## 6           6 German                        NA    1
## 7          NA            <NA>                10    1
```

13.9   Store analysis

How are the stores performing.

13.9.1   SQL store revenue stream

How are the stores performing? The SQL code shows the payments made to each store in the business.

```r
rs <- dbGetQuery(
  con,
  "select store_id,sum(p.amount) amt,count(*) cnt
                 from payment p
                      join staff s
                        on p.staff_id = s.staff_id
              group by store_id order by 2 desc
               ;
                "
)
sp_print_df(head(rs))
```

store_id	amt	cnt
2	31059.92	7304
1	30252.12	7292

13.9.1.1   Exercise dplyr store revenue stream

Complete the following code block to return the payments made to each store.

```r
payment_table <- tbl(con, "payment") # DBI::dbReadTable(con, "payment")
staff_table <- tbl(con, "staff") # DBI::dbReadTable(con, "staff")

store_revenue <- payment_table %>%
  inner_join(staff_table, by = "staff_id", suffix = c(".p", ".s")) %>%
  head()

store_revenue
```

```
## # Source:   lazy query [?? x 16]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##   payment_id customer_id staff_id rental_id amount payment_date
##        <int>       <int>    <int>     <int>  <dbl> <dttm>
## 1      17503         341        2      1520   7.99 2007-02-15 22:25:46
```

```
## 2         17504            341             1        1778    1.99 2007-02-16 17:23:14
## 3         17505            341             1        1849    7.99 2007-02-16 22:41:45
## 4         17506            341             2        2829    2.99 2007-02-19 19:39:56
## 5         17507            341             2        3130    7.99 2007-02-20 17:31:48
## 6         17508            341             1        3382    5.99 2007-02-21 12:33:49
## # ... with 10 more variables: first_name <chr>, last_name <chr>,
## #   address_id <int>, email <chr>, store_id <int>, active <lgl>,
## #   username <chr>, password <chr>, last_update <dttm>, picture <blob>
```

13.9.2   SQL:Estimate Outstanding Balance

The following SQL code calculates for each store

1. the number of payments still open and closed from the DVD Rental Stores customer base.
2. the total amount that their customers have paid
3. the average price per/movie based off of the movies that have been paid.
4. the estimated outstanding balance based off the open unpaid rentals * the average price per paid movie.

```r
rs <- dbGetQuery(
  con,
  "SELECT s.store_id store,sum(CASE WHEN payment_id IS NULL THEN 1 ELSE 0 END) open
    ,sum(CASE WHEN payment_id IS NOT NULL THEN 1 ELSE 0 END) paid
    ,sum(p.amount) paid_amt
    ,count(*) rentals
    ,round(sum(p.amount) / sum(CASE WHEN payment_id IS NOT NULL
                                    THEN 1
                                    ELSE 0
                               END), 2) avg_price
    ,round(round(sum(p.amount) / sum(CASE WHEN payment_id IS NOT NULL
                                    THEN 1
                                    ELSE 0
                               END), 2) * sum(CASE WHEN payment_id IS NULL
                                                   THEN 1
                                                   ELSE 0
                               END), 2) est_balance
FROM rental r
LEFT JOIN payment p
    ON r.rental_id = p.rental_id
JOIN staff s
    ON r.staff_id = s.staff_id
group by s.store_id;
"
)
sp_print_df(head(rs))
```

store	open	paid	paid_amt	rentals	avg_price	est_balance
1	713	7331	30498.71	8044	4.16	2966.08
2	739	7265	30813.33	8004	4.24	3133.36

```r
rs
```

```
##   store open paid paid_amt rentals avg_price est_balance
## 1     1  713 7331 30498.71    8044      4.16     2966.08
## 2     2  739 7265 30813.33    8004      4.24     3133.36
```

13.9.2.1   Exercise Dplyr Modify the following dplyr code to match the SQL output from above.

```r
payment_table <- tbl(con, "payment") # DBI::dbReadTable(con, "payment")
rental_table <- tbl(con, "rental") # DBI::dbReadTable(con, "rental")

est_bal <- rental_table %>%
  left_join(payment_table, by = c("rental_id", "rental_id"), suffix = c(".r", ".p")) %>%
  mutate(
    missing = ifelse(is.na(payment_id), 1, 0)
    , found = ifelse(!is.na(payment_id), 1, 0)
  ) %>%
  summarize(
    open = sum(missing, na.rm = TRUE)
    , paid = sum(found, na.rm = TRUE)
    , paid_amt = sum(amount, na.rm = TRUE)
    , rentals = n()
  ) %>%
  summarize(
    open = open
    , paid = paid
    , paid_amt = paid_amt
    , rentals = rentals
    , avg_price = paid_amt / paid
    , est_balance = paid_amt / paid * open
  )
est_bal
```

```
## # Source:   lazy query [?? x 6]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##    open  paid paid_amt rentals         avg_price est_balance
##   <dbl> <dbl>    <dbl> <S3: integer64>     <dbl>       <dbl>
## 1  1452 14596   61312. 16048               4.20       6099.
```

13.9.3   SQL actual outstanding balance

In the previous exercise, we estimated the outstanding amount. After reviewing the rental table, the actual movie rental rate is in the table. We use that to calculate the outstanding balance below.

```r
rs <- dbGetQuery(
  con,
  "SELECT sum(f.rental_rate) open_amt
     ,count(*) count
FROM rental r
LEFT JOIN payment p
    ON r.rental_id = p.rental_id
INNER JOIN inventory i
    ON r.inventory_id = i.inventory_id
INNER JOIN film f
    ON i.film_id = f.film_id
WHERE p.rental_id IS NULL
;"
)
sp_print_df(head(rs))
```

open_amt	count
4297.48	1452

```
rs
```

```
##    open_amt count
## 1  4297.48  1452
```

```r
payment_table <- tbl(con, "payment") # DBI::dbReadTable(con, "payment")
rental_table <- tbl(con, "rental") # DBI::dbReadTable(con, "rental")
inventory_table <- tbl(con, "inventory") # DBI::dbReadTable(con, "inventory")
film_table <- tbl(con, "film") # DBI::dbReadTable(con, "film")

act_bal <- rental_table %>%
  left_join(payment_table, by = c("rental_id", "rental_id"), suffix = c(".r", ".p")) %>%
  inner_join(inventory_table, by = c("inventory_id", "inventory_id"), suffix = c(".r", ".i")) %>%
  inner_join(film_table, by = c("film_id", "film_id"), suffix = c(".i", ".f")) %>%
  head()

act_bal
```

```
## # Source:   lazy query [?? x 27]
## # Database: postgres [postgres@localhost:5432/dvdrental]
##   rental_id rental_date          inventory_id customer_id.r
##       <int> <dttm>                      <int>         <int>
## 1         1 2005-05-24 22:53:30           367           130
## 2         2 2005-05-24 22:54:33          1525           459
## 3         3 2005-05-24 23:03:39          1711           408
## 4         4 2005-05-24 23:04:41          2452           333
## 5         5 2005-05-24 23:05:21          2079           222
## 6         6 2005-05-24 23:08:07          2792           549
## # ... with 23 more variables: return_date <dttm>, staff_id.r <int>,
## #   last_update.r <dttm>, payment_id <int>, customer_id.p <int>,
## #   staff_id.p <int>, amount <dbl>, payment_date <dttm>, film_id <int>,
## #   store_id <int>, last_update.i <dttm>, title <chr>, description <chr>,
## #   release_year <int>, language_id <int>, rental_duration <int>,
## #   rental_rate <dbl>, length <int>, replacement_cost <dbl>, rating <S3:
## #   pq_mpaa_rating>, last_update <dttm>, special_features <S3: pq__text>,
## #   fulltext <S3: pq_tsvector>
```

13.9.4 Rank customers with highest open amounts

```r
rs <- dbGetQuery(
  con,
  "select c.customer_id,c.first_name,c.last_name,sum(f.rental_rate) open_amt,count(*) count
                from rental r
                    left outer join payment p
                      on r.rental_id = p.rental_id
                    join inventory i
                      on r.inventory_id = i.inventory_id
                    join film f
                      on i.film_id = f.film_id
                    join customer c
                      on r.customer_id = c.customer_id
                where p.rental_id is null
```

```
                   group by c.customer_id,c.first_name,c.last_name
                   order by open_amt desc
                   limit 25
                   ;"
)
sp_print_df(head(rs))
```

customer_id	first_name	last_name	open_amt	count
293	Mae	Fletcher	35.90	10
307	Joseph	Joy	31.90	10
316	Steven	Curley	31.90	10
299	James	Gannon	30.91	9
274	Naomi	Jennings	29.92	8
326	Jose	Andrew	28.93	7

```
rs
```

```
##    customer_id first_name  last_name open_amt count
## 1          293        Mae   Fletcher    35.90    10
## 2          307     Joseph        Joy    31.90    10
## 3          316     Steven     Curley    31.90    10
## 4          299      James     Gannon    30.91     9
## 5          274      Naomi   Jennings    29.92     8
## 6          326       Jose     Andrew    28.93     7
## 7          338     Dennis     Gilman    27.92     8
## 8          277       Olga    Jimenez    27.92     8
## 9          327      Larry   Thrasher    26.93     7
## 10         330      Scott    Shelley    26.93     7
## 11         322      Jason  Morrissey    26.91     9
## 12         340    Patrick     Newsom    25.92     8
## 13         336     Joshua       Mark    25.92     8
## 14         304      David      Royal    24.93     7
## 15         339     Walter   Perryman    23.94     6
## 16         239     Minnie     Romero    23.94     6
## 17         310     Daniel     Cabral    22.93     7
## 18         296     Ramona       Hale    22.93     7
## 19         313     Donald      Mahon    22.93     7
## 20         287      Becky      Miles    22.93     7
## 21         272        Kay   Caldwell    22.93     7
## 22         303    William Satterfield  22.93     7
## 23         329      Frank   Waggoner    22.91     9
## 24         311       Paul      Trout    21.92     8
## 25         109       Edna       West    20.93     7
```

13.9.5   what film has been rented the most

```
rs <- dbGetQuery(
  con,
  "SELECT i.film_id
    ,f.title
    ,rental_rate
    ,sum(rental_rate) revenue
    ,count(*) count --16044
FROM rental r
```

```
INNER JOIN inventory i
    ON r.inventory_id = i.inventory_id
INNER JOIN film f
    ON i.film_id = f.film_id
GROUP BY i.film_id
    ,f.title
    ,rental_rate
ORDER BY count DESC
LIMIT 25
;"
)
sp_print_df(head(rs))
```

film_id	title	rental_rate	revenue	count
103	Bucket Brotherhood	4.99	169.66	34
738	Rocketeer Mother	0.99	32.67	33
489	Juggler Hardly	0.99	31.68	32
730	Ridgemont Submarine	0.99	31.68	32
767	Scalawag Duck	4.99	159.68	32
331	Forward Temple	2.99	95.68	32

rs

```
##     film_id                title rental_rate revenue count
## 1       103  Bucket Brotherhood        4.99  169.66    34
## 2       738     Rocketeer Mother        0.99   32.67    33
## 3       489       Juggler Hardly        0.99   31.68    32
## 4       730 Ridgemont Submarine        0.99   31.68    32
## 5       767        Scalawag Duck        4.99  159.68    32
## 6       331       Forward Temple        2.99   95.68    32
## 7       382       Grit Clockwork        0.99   31.68    32
## 8       735         Robbers Joon        2.99   92.69    31
## 9       973            Wife Turn        4.99  154.69    31
## 10      621         Network Peak        2.99   92.69    31
## 11     1000            Zorro Ark        4.99  154.69    31
## 12       31        Apache Divine        4.99  154.69    31
## 13      369    Goodfellas Salute        4.99  154.69    31
## 14      753      Rush Goodfellas        0.99   30.69    31
## 15      891       Timberland Sky        0.99   30.69    31
## 16      418         Hobbit Alien        0.99   30.69    31
## 17      127        Cat Coneheads        4.99  149.70    30
## 18      559           Married Go        2.99   89.70    30
## 19      374        Graffiti Love        0.99   29.70    30
## 20      748 Rugrats Shakespeare        0.99   29.70    30
## 21      239         Dogma Family        4.99  149.70    30
## 22      285      English Bulworth        0.99   29.70    30
## 23      109    Butterfly Chocolat        0.99   29.70    30
## 24      450       Idols Snatchers        2.99   89.70    30
## 25      609         Muscle Bright        2.99   89.70    30
```

13.9.6 what film has been generated the most revenue assuming all amounts are collected

```
rs <- dbGetQuery(
  con,
  "select i.film_id,f.title,rental_rate
                   ,sum(rental_rate) revenue,count(*) count  --16044
                from rental r
                    join inventory i
                      on r.inventory_id = i.inventory_id
                    join film f
                      on i.film_id = f.film_id
              group by i.film_id,f.title,rental_rate
              order by revenue desc
              ;"
)
sp_print_df(head(rs))
```

film_id	title	rental_rate	revenue	count
103	Bucket Brotherhood	4.99	169.66	34
767	Scalawag Duck	4.99	159.68	32
973	Wife Turn	4.99	154.69	31
31	Apache Divine	4.99	154.69	31
369	Goodfellas Salute	4.99	154.69	31
1000	Zorro Ark	4.99	154.69	31

13.9.7 which films are in one store but not the other.

```
rs <- dbGetQuery(
  con,
  "select coalesce(i1.film_id,i2.film_id) film_id
                   ,f.title,f.rental_rate,i1.store_id,i1.count,i2.store_id,i2.count
                from     (select film_id,store_id,count(*) count
                            from inventory where store_id = 1
                           group by film_id,store_id) as i1
                       full outer join
                         (select film_id,store_id,count(*) count
                            from inventory where store_id = 2
                           group by film_id,store_id
                          ) as i2
                        on i1.film_id = i2.film_id
                       join film f
                         on coalesce(i1.film_id,i2.film_id) = f.film_id
                 where i1.film_id is null or i2.film_id is null
               order by f.title  ;
              "
)
sp_print_df(head(rs))
```

film_id	title	rental_rate	store_id	count	store_id..6	count..7
2	Ace Goldfinger	4.99	NA	NA	2	3
3	Adaptation Holes	2.99	NA	NA	2	4
5	African Egg	2.99	NA	NA	2	3
8	Airport Pollock	4.99	NA	NA	2	4
13	Ali Forever	4.99	NA	NA	2	4
20	Amelie Hellfighters	4.99	1	3	NA	NA

13.9.8   Compute the outstanding balance.

```
rs <- dbGetQuery(
  con,
  "select sum(f.rental_rate) open_amt,count(*) count
                 from rental r
                     left outer join payment p
                       on r.rental_id = p.rental_id
                     join inventory i
                       on r.inventory_id = i.inventory_id
                     join film f
                       on i.film_id = f.film_id
                where p.rental_id is null
                ;"
)
sp_print_df(head(rs))
```

open_amt	count
4297.48	1452

13.10   Different strategies for interacting with the database

select examples dbGetQuery returns the entire result set as a data frame.
For large returned datasets, complex or inefficient SQL statements, this may take a long time.

```
   dbSendQuery: parses, compiles, creates the optimized execution plan.
       dbFetch: Execute optimzed execution plan and return the dataset.
dbClearResult: remove pending query results from the database to your R environment
```

13.10.1   Use dbGetQuery

How many customers are there in the DVD Rental System

```
rs1 <- dbGetQuery(con, "select * from customer;")
sp_print_df(head(rs1))
```

customer_id	store_id	first_name	last_name	email	address_id	activebool	crea
524	1	Jared	Ely	jared.ely@sakilacustomer.org	530	TRUE	200
1	1	Mary	Smith	mary.smith@sakilacustomer.org	5	TRUE	200
2	1	Patricia	Johnson	patricia.johnson@sakilacustomer.org	6	TRUE	200
3	1	Linda	Williams	linda.williams@sakilacustomer.org	7	TRUE	200
4	2	Barbara	Jones	barbara.jones@sakilacustomer.org	8	TRUE	200
5	1	Elizabeth	Brown	elizabeth.brown@sakilacustomer.org	9	TRUE	200

```
pco <- dbSendQuery(con, "select * from customer;")
rs2 <- dbFetch(pco)
dbClearResult(pco)
sp_print_df(head(rs2))
```

customer_id	store_id	first_name	last_name	email	address_id	activebool	crea
524	1	Jared	Ely	jared.ely@sakilacustomer.org	530	TRUE	200
1	1	Mary	Smith	mary.smith@sakilacustomer.org	5	TRUE	200
2	1	Patricia	Johnson	patricia.johnson@sakilacustomer.org	6	TRUE	200
3	1	Linda	Williams	linda.williams@sakilacustomer.org	7	TRUE	200
4	2	Barbara	Jones	barbara.jones@sakilacustomer.org	8	TRUE	200
5	1	Elizabeth	Brown	elizabeth.brown@sakilacustomer.org	9	TRUE	200

13.10.2 Use dbExecute

13.10.3 Anti join – Find Sophie who has never rented a movie.

```
customer_table <- DBI::dbReadTable(con, "customer")
rental_table <- DBI::dbReadTable(con, "rental")

customer_tbl <- dplyr::tbl(con, "customer")
rental_tbl <- dplyr::tbl(con, "rental")

dplyr_tbl_loj <-
  left_join(customer_tbl, rental_tbl, by = "customer_id", suffix = c(".c", ".r")) %>%
  filter(is.na(rental_id)) %>%
  select(c("first_name", "last_name", "email"))

rs <- dbGetQuery(
  con,
  "select c.first_name
                  ,c.last_name
                  ,c.email
              from customer c
                  left outer join rental r
                      on c.customer_id = r.customer_id
              where r.rental_id is null;
          "
)
sp_print_df(head(rs))
```

first_name	last_name	email
Sophie	Yang	email@email.com

```
View(dplyr_tbl_loj)
```

country

fk_city

city

fk_address_city

address

store_address_id_fkey

customer_address_id_fkey

store

customer

staff_address_id_fkey

store_manager_staff_id_fkey

payment_customer_id_fkey

staff

payment   payment_staff_id_fkey
               rental_customer_id_fkey
               rental_staff_id_key

payment_rental_id_fkey

rental

rental_inventory_id_fkey

inventory

inventory_film_id_fkey

film   film_language_id_fkey

film_actor_film_id_fkey   language

film_actor   film_category_film_id_fkey

film_actor_actor_id_fkey

actor

film_category

film_category_category_id_fkey

category

```r
# diconnect from the db
dbDisconnect(con)

sp_docker_stop("sql-pet")
```

```r
knitr::knit_exit()
```

Chapter 14

SQL Quick start - simple retrieval (15)

This chapter demonstrates:

- Several elementary SQL statements
- SQL databases and 3rd normal form

14.1 Intro

- Coverage in this book. There are many SQL tutorials that are available. For example, we are drawing some materials from a tutorial we recommend. In particular, we will not replicate the lessons there, which you might want to complete. Instead, we are showing strategies that are recommended for R users. That will include some translations of queries that are discussed there.

- https://datacarpentry.org/R-ecology-lesson/05-r-and-databases.html Very good intro. How is ours different?

Start up the `docker-pet` container

```
sp_docker_start("sql-pet")
```

Now connect to the `dvdrental` database with R

```
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password =  Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10)
con
```

```
## <PqConnection> dvdrental@localhost:5432
```

```
colFmt <- function(x,color)
{
  # x string
  # color
  outputFormat = knitr::opts_knit$get("rmarkdown.pandoc.to")
  if(outputFormat == 'latex')
    paste("\\textcolor{",color,"}{",x,"}",sep="")
  else if(outputFormat == 'html')
```

```
    paste("<font color='",color,"'>",x,"</font>",sep="")
  else
    x
}

# sample call
# * `r colFmt('Cover inline tables in future section','red')`
```

Moved this from 11-elementary-queries

```
dplyr_summary_df <-
    read.delim(
    '11-dplyr_sql_summary_table.tsv',
    header = TRUE,
    sep = '\t',
    as.is = TRUE
    )

head(dplyr_summary_df)
```

```
##   In          Dplyr_Function
## 1  Y                arrange()
## 2 Y?              distinct()
## 3  Y        select() rename()
## 4  N                  pull()
## 5  Y     mutate() transmute()
## 6  Y summarise() summarize()
##                                    description
## 1                  Arrange rows by variables
## 2           Return rows with matching conditions
## 3              Select/rename variables by name
## 4                  Pull out a single variable
## 5                          Add new variables
## 6 Reduces multiple values down to a single value
##                          SQL_Clause Notes              Category
## 1                          ORDER BY    NA Basic single-table verbs
## 2                  SELECT distinct *    NA Basic single-table verbs
## 3        SELECT column_name alias_name    NA Basic single-table verbs
## 4                SELECT column_name;    NA Basic single-table verbs
## 5 SELECT computed_value computed_name    NA Basic single-table verbs
## 6 SELECT aggregate_functions GROUP BY    NA Basic single-table verbs
```

14.2   Databases and Third Normal Form - 3NF

Most relational database applications are designed to be third normal form "like", 3NF. The key benefits of 3NF are

1. speedy on-line transactional processing, OLTP.
2. improved referential integrity, reduce modification anomalies that can occur during an insert, update, or delete operation.
3. reduced storage, elimination of redundant data.

3NF is great for database application input performance, but not so great for getting the data back out for the data analyst or report writer. As a data analyst, you might get the ubiquitous Excel spreadsheet
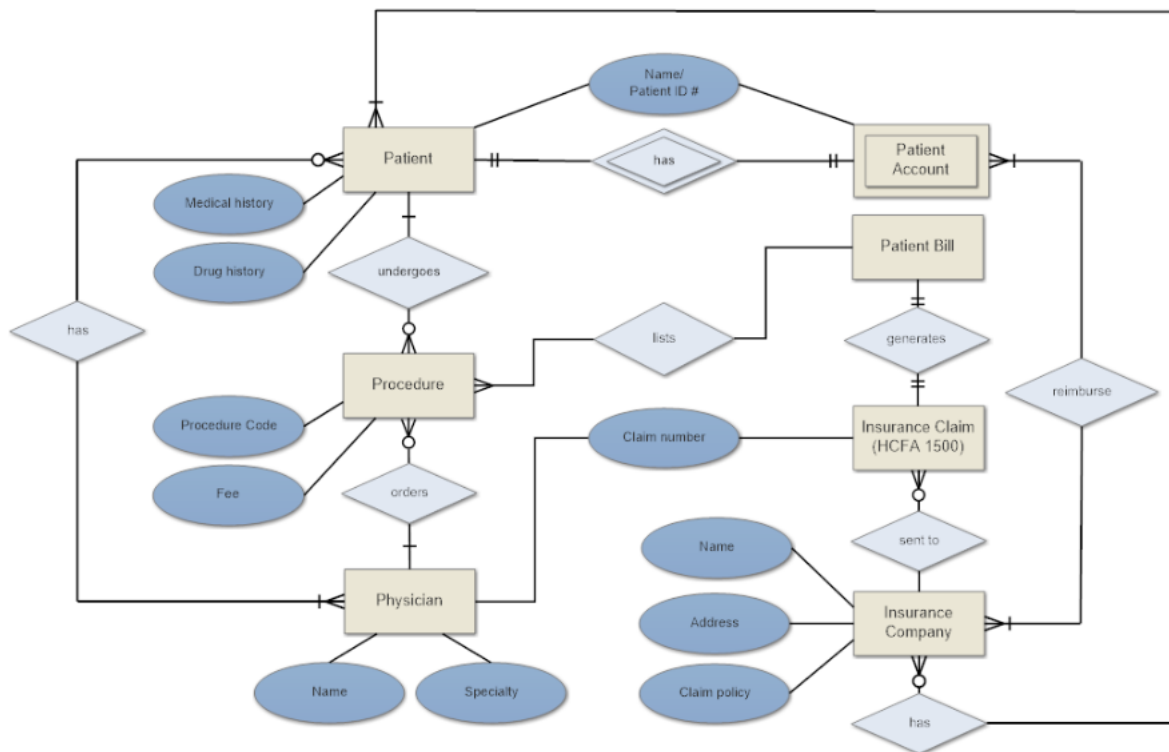
Figure 14.1: hospital-billing-erd

with all the information needed to start an Exploratory Data Analysis, EDA. The spreadsheet may have provider, patient, diagnosis, procedure, and insurance information all "neatly" arranged on a single row. At least "neatly" when compared to the same information stored in the database, in at least 5 tables.

For this tutorial, the most important thing to know about 3NF is that the data you are looking for gets spread across many many tables. Working in a relational database requires you to

1. find the many many different tables that contains your data.

2. Understand the relationships that tie the tables together correctly to ensure that data is not dropped or duplicated. Data that is dropped or duplicated can either over or understate your aggregated numeric values.

https://www.smartdraw.com/entity-relationship-diagram/examples/hospital-billing-entity-relationship-diagram/

Real life applications have 100's or even 1000's of tables supporting the application. The goal is to transform the application data model into a useful data analysis model using the DDL and DML SQL statements.

14.3 SQL Commands

SQL commands fall into four categories.

SQL Category	Definition
DDL:Data Definition Language	DBA's execute these commands to define objects in the database.

SQL Category	Definition
DML:Data Manipulation Language	Users and developers execute these commands to investigate data.
DCL:Data Control Language	DBA's execute these commands to grant/revoke access to
TCL:Transaction Control Language	Developers execute these commands when developing applications.

Data analysts use the SELECT DML command to learn interesting things about the data stored in the database. Applications are used to control the insert, update, and deletion of data in the database. Data users can update the database objects via the application which enforces referential integrity in the database. Data users should never directly update data application database objects. Leave this task to the developers and DBA's.

DBA's can setup a sandbox within the database for a data analyst. The application(s) do not maintain the data in the sandbox.

The `sql-pet` database is tiny, but for the purposes of these exercises, we assume that data so large that it will not easily fit into the memory of your laptop.

This tutorial focuses on the most frequently used SQL statement, the SQL SELECT statement.

A SQL SELECT statement consists of 1 to 6 clauses.

SQL Clause	DPLYR Verb	SQL Description
SELECT	SELECT()	Contains a list of column names from an object or a derived value.
	mutate()	
FROM		Contains a list of related tables from which the SELECT list of columns is derived.
WHERE	filter()	Provides the filter conditions the objects in the FROM clause must meet.
GROUP BY	group_by()	Contains a list rollup aggregation columns.
HAVING		Provides the filter condition on the the GROUP BY clause.
ORDER BY	arrange()	Contains a list of column names indicating the order of the column value. Each column can be either ASCending or DEScending.

The foundation of the SQL language is based set theory and the result of a SQL SELECT statement is referred to as a result set. A SQL SELECT statement is "guaranteed" to return the same set of data, but not necessarily in the same order. However, in practice, the result set is usually in the same order.

SQL SELECT statements can be broken up into two categories, SELECT detail statements and SELECT aggregate statements.

SELECT DETAIL	SELECT AGGREGATE
select det_col1...det_coln	select det_agg1..., agg1,...,aggn
from same	from same
where same	where same
	group by det_agg1
	having
order by same	order by same

Table 14.4: select all columns

store_id	manager_staff_id	address_id	last_update
1	1	1	2006-02-15 09:57:12
2	2	2	2006-02-15 09:57:12

SELECT DETAIL	SELECT AGGREGATE

The difference between the two statements is the AGGREGATE has

1. select clause has one or more detail columns, det_agg1..., on which values get aggregated against/rolled up to.
2. select clause zero or more aggregated values, agg1, ..., aggn
3. group by clause is required and matches the one or more detail columns, det_agg1.
4. having clause is optional and adds a filter condition on one or more agg1 ... aggn values.

14.4 SQL SELECT Quick Start

This section focuses on getting new SQL users familiar with the six SQL query clauses and a single table. SQL queries from multiple tables are discussed in the JOIN section of this tutorial. The JOIN section resolves the issue introduced with 3NF, the splitting of data into many many tables, back into a denormalaized format similar to the Excel spreadsheet.

The DBI::dbGetQuery function is used to submit SQL SELECT statements to the Postgres database. At a minimum it requires two parameters, a connection object and a SQL SELECT statement.

In the following section we only look at SELECT DETAIL statements.

14.4.1 SELECT Clause: Column Selection – Vertical Partioning of Data

14.4.1.1 1. Simplest SQL query: All rows and all columns from a single table.

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store;
  ")
kable(rs,caption = 'select all columns')
```

14.4.1.2 2. Same Query as 1, but only show first two columns;

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select STORE_ID, manager_staff_id from store;
  ")
kable(rs,caption = 'select first two columns only')
```

Table 14.5: select first two columns only

store_id	manager_staff_id
1	1
2	2

Table 14.6: reverse the column order

manager_staff_id	store_id
1	1
2	2

14.4.1.3   3. Same Query as 2, but reverse the column order

dvdrental=# select manager_staff_id,store_id from store;

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select manager_staff_id,store_id from store;
  ")
kable(rs,caption = 'reverse the column order')
```

14.4.1.4   4. Rename Columns – SQL column alias in the result set

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select manager_staff_id mgr_sid,store_id st_id from store;
  ")
kable(rs,caption = 'Rename Columns')
```

```
The manager_staff_id has changed to mgr_sid.
store_id has changed to st_id.
```

```
Note that the column names have changed in the result set only, not in the actual database table.
The DBA's will not allow a space or other special characters in a database table column name.
```

```
Some motivations for aliasing the result set column names are
```

```
   1.   Some database table column names are not user friendly.
   2.  When multiple tables are joined, the column names may be the same in one or more tables and one needs to d
```

14.4.1.5   5. Adding Meta Data Columns to the Result Set

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select 'derived column' showing
```

Table 14.7: Rename Columns

mgr__sid	st__id
1	1
2	2

Table 14.8: Adding Meta Data Columns

showing	store__id	manager__staff__id	address__id	last__update	db	user	dtts
derived column	1	1	1	2006-02-15 09:57:12	dvdrental	postgres	2018/12/18 2
derived column	2	2	2	2006-02-15 09:57:12	dvdrental	postgres	2018/12/18 2

```
      ,*
      ,current_database() db
      ,user
      ,to_char(now(),'YYYY/MM/DD HH24:MI:SS') dtts
   from store;
  ")
kable(rs,caption = 'Adding Meta Data Columns')
```

All the previous examples easily fit on a single line. This one is longer. Each column is entered on its own

1. The showing column is a hard coded string surrounded by single quotes. Note that single quotes are for ha
2. The db and dtts, date timestamp, are new columns generated from Postgres System Information Functions.
3. Note that `user` is not a function call, no parenthesis.

14.4.2 SQL Comments

SQL supports both a single line comment, preceed the line with two dashes, `--`, and a C like block comment, * … */.

14.4.2.1 6. Single line comment –

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select 'single line comment, dtts' showing
        ,*
        ,current_database() db
        ,user
    --   ,to_char(now(),'YYYY/MM/DD HH24:MI:SS') dtts
    from store;
  ")
kable(rs,caption = 'Sincle line comment')
```

The dtts line is commented out with the two dashes and is dropped from the end of the result set columns.

14.4.2.2 7. Multi-line comment /*…*/

Table 14.9: Sincle line comment

showing	store_id	manager_staff_id	address_id	last_update	db	user
single line comment, dtts	1	1	1	2006-02-15 09:57:12	dvdrental	postgres
single line comment, dtts	2	2	2	2006-02-15 09:57:12	dvdrental	postgres

Table 14.10: Multi-line comment

showing	store_id	manager_staff_id	address_id	last_update
block comment drop db, user, and dtts	1	1	1	2006-02-15 09:57:12
block comment drop db, user, and dtts	2	2	2	2006-02-15 09:57:12

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select 'block comment drop db, user, and dtts' showing
        ,*
        /*
        ,current_database() db
        ,user
        ,to_char(now(),'YYYY/MM/DD HH24:MI:SS') dtts
        */
    from store;
  ")
kable(rs,caption = 'Multi-line comment')
```

The three columns db, user, and dtts, between the /* and */ have been commented and no longer appear as the

14.4.3   FROM Clause

The `FROM` clause contains one or more datasets, usually database tables/views, from which the `SELECT` columns are derived. For now, in the examples, we are only using a single table. If the database reflects a relational model, your data is likely spread out over several tables. The key take away when beginning your analysis is to pick the table that has most of the data that you need for your analysis. This table becomes your main or driving table to build your SQL query statement around. After identifying your driving table, potentially save yourself a lot of time and heart ache, review any view that is built on your driving table. If one or more exist, especially, if vendor built, may already have the additional information needed for your analysis.

Insert SQL here or link to Views dependent on what

In this tutorial, there is only a single user hitting the database and row/table locking is not necessary and considered out of scope.

14.4.3.1   Table Uses

- A table can be used more than once in a FROM clause. These are self-referencing tables. An example is an EMPLOYEE table which contains a foriegn key to her manager. Her manager also has a foriegn key to her manager, etc up the corporate ladder.

- In the example above, the EMPLOYEE table plays two roles, employee and manager. The next line shows the FROM clause showing the same table used twice.

FROM EMPLOYEE EE, EMPLOYEE MGR

- The EE and MGR are aliases for the EMPLOYEE table and represent the different roles the EMPLOYEE table plays.

- Since all the column names are exactly the same for the EE and MGR role, the column names need to be prefixed with their role alias, e.g., SELECT MGR.EE_NAME, EE.EE_NAME … shows the manager name and her employee name(s) who work for her.

- It is a good habit to always alias your tables and prefix your column names with the table alias to eliminate any ambiguity as to where the column came from. This is critical where there is inconsistent table column naming convention. It also helps when debugging larger SQL queries.

- Cover inline tables in future section

Side Note: Do not create an unintended Cartesian join. If one has more than one table in the FROM clause, make

14.4.4   WHERE Clause: Row Selection – Horizontal Partitioning of Data

In the previous SELECT clause section, the SELECT statement either partitioned data vertically across the table columns or derived vertical column values. This section provides examples that partitions the table data across rows in the table.

The WHERE clause defines all the conditions the data must meet to be included or excluded in the final result set. If all the conditions are met data is returned or it is rejected. This is commonly referred to as the data set filter condition.

Side Note: For performance optimization reasons, the WHERE clause should reduce the dataset down to the small

The WHERE condition(s) can be simple or complex, but in the end are the appliction of the logic rules shown in the table below.

p	q	p and q	p or q
T	T	T	T
T	F	F	T
T	N	N	T
F	F	F	F
F	N	F	T
N	N	N	N

When the filter logic is complex, it is sometimes easier to represent the where clause symbollically and apply a version of DeMorgan's law which is shown below.

1. (A and B)' = A' or B'
2. (A or B)' = A' and B'

14.4.4.1   Examples Continued

We begin with `1`, our simplest SQL query.

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store;
```

Table 14.12: select all columns

store_id	manager_staff_id	address_id	last_update
1	1	1	2006-02-15 09:57:12
2	2	2	2006-02-15 09:57:12

Table 14.13: WHERE always FALSE

store_id	manager_staff_id	address_id	last_update

```
  ")
kable(rs,caption = 'select all columns')
```

14.4.4.2   8 WHERE condition logically never TRUE.

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store where 1 = 0;
  ")
kable(rs,caption = 'WHERE always FALSE')
```

Since 1 = 0 is always false, no rows are ever returned. Initially this construct seems useless, but actually

14.4.4.3   9 WHERE condition logically always TRUE.

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store where 1 = 1;
  ")
kable(rs,caption = 'WHERE always TRUE')
```

Since 1 = 1 is always true, all rows are always returned. Initially this construct seems useless, but actuall

14.4.4.4   10 WHERE equality condition

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store where store_id = 2;
  ")
kable(rs,caption = 'WHERE EQUAL')
```

The only row where the store_id = 2 is row 2 and it is the only row returned.

Table 14.14: WHERE always TRUE

store_id	manager_staff_id	address_id	last_update
1	1	1	2006-02-15 09:57:12
2	2	2	2006-02-15 09:57:12

Table 14.15: WHERE EQUAL

store_id	manager_staff_id	address_id	last_update
2	2	2	2006-02-15 09:57:12

14.4.4.5 11 WHERE NOT equal conditions

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store where store_id <> 2;
  ")
kable(rs,caption = 'WHERE NOT EQUAL')
```

`<> is syntactically the same as !=`

The only row where the `store_id <> 2` is row 1 and only row 1 is returned.

14.4.4.6 12 WHERE OR condition

```
rs <-
  DBI::dbGetQuery(
  con,
  "
  select * from store where manager_staff_id = 1 or store_id < 3;
  ")
kable(rs,caption = 'WHERE OR condition')
```

The first condition `manager_staff_id = 1` returns a single row and the second condition `store_id < 3` returns tw

Following table is modified from http://www.tutorialspoint.com/sql/sql-operators

SQL Comparison Operators

Operator	Description	example
=	Checks if the values of two operands are equal or not, if yes then condition becomes true.	(a = b) is not true.
!=	Checks if the values of two operands are equal or not, if values are not equal then condition becomes true.	(a != b) is true.
<>	Checks if the values of two operands are equal or not, if values are not equal then condition becomes true.	(a <> b) is true.
>	Checks if the value of left operand is greater than the value of right operand, if yes then condition becomes true.	(a > b) is not true.
<	Checks if the value of left operand is less than the value of right operand, if yes then condition becomes true.	(a < b) is true.
>=	Checks if the value of left operand is greater than or equal to the value of right operand, if yes then condition becomes true.	(a >= b) is not true.

Operator	Description	example
<=	Checks if the value of left operand is less than or equal to the value of right operand, if yes then condition becomes true.	(a <= b) is true.
!<	Checks if the value of left operand is not less than the value of right operand, if yes then condition becomes true.	(a !< b) is false.
!>	Checks if the value of left operand is not greater than the value of right operand, if yes then condition becomes true.	(a !> b) is true.

Operator	Description
ALL	The ALL operator is used to compare a value to all values in another value set.
AND	The AND operator allows the existence of multiple conditions in an SQL statement's WHERE clause.
ANY	The ANY operator is used to compare a value to any applicable value in the list as per the condition.
BETWEEN	The BETWEEN operator is used to search for values that are within a set of values, given the minimum value and the maximum value.
EXISTS	The EXISTS operator is used to search for the presence of a row in a specified table that meets a certain criterion.
IN	The IN operator is used to compare a value to a list of literal values that have been specified.
LIKE	The LIKE operator is used to compare a value to similar values using wildcard operators.
NOT	The NOT operator reverses the meaning of the logical operator with which it is used. Eg: NOT EXISTS, NOT BETWEEN, NOT IN, etc. This is a negate operator.
OR	The OR operator is used to combine multiple conditions in an SQL statement's WHERE clause.
IS NULL	The NULL operator is used to compare a value with a NULL value.
UNIQUE	The UNIQUE operator searches every row of a specified table for uniqueness (no duplicates).

https://pgexercises.com/questions/basic

TO-DO's

1. inline tables
2. correlated subqueries

14.5   Paradigm Shift from R-Dplyr to SQL

Paraphrasing what some have said with an R dplyr background and no SQL experience, "It is like working from the inside out." This sentiment occurs because

1. The SQL SELECT statement begins at the end, the SELECT clause, and drills backwards, loosely speaking, to derive the desired result set.
2. SQL SELECT statements are an all or nothing proposition. One gets nothing if there is any kind of syntax error.

3. SQL SELECT result sets can be quite opaque. The WHERE clause can be very dense and difficult to trace through. It is rarely ever linear in nature.

4. Validating all the permutations in the where clause can be tough and tedious.

Table 14.16: WHERE NOT EQUAL

store_id	manager_staff_id	address_id	last_update
1	1	1	2006-02-15 09:57:12

Table 14.17: WHERE OR condition

store_id	manager_staff_id	address_id	last_update
1	1	1	2006-02-15 09:57:12
2	2	2	2006-02-15 09:57:12

14.5.1 Big bang versus piped incremental steps.

1. Dplyr starts with one or more sources joined together in a conceptually similar way that SQL joins sources.
2. The pipe and filter() function breaks down the filter conditions into small managable logical steps. This makes it much easier to understand what is happening in the derivation of the final tibble. Adding tees through out the pipe line gives one full trace back of all the data transformations at every pipe.

Helpful tidyverse functions that output tibbles: tbl_module function in https://github.com/nhemerson/tibbleColumns package;

Mental picture: SQL approach: Imagine a data lake named Niagera Falls and drinking from it without drowning. R-Dplyr approach: Imagine a resturant at the bottom of the Niagera Falls data lake and having a refreshing dring out of the water faucet.

14.5.2 SQL Execution Order

The table below is derived from this site. https://www.periscopedata.com/blog/sql-query-order-of-operations It shows what goes on under the hood SQL SELECT hood.

SEQ	SQL	Function	Dplyr
1	WITH	Common Table expression, CTE, one or more datasets/tables used FROM clause.	.data parameter in dplyr functions
2	FROM	Choose and join tables to get base data	.data parameter in dplyr functions
3	ON	Choose and join tables to get base data	dplyr join family of functions
4	JOIN	Choose and join tables to get base data	dplyr join family of functions
5	WHERE	filters the base data	dplyr filter()
6	GROUP BY	aggregates the base data	dplyr group_by family of functions
7	WITH CUBE/ROLLUP	aggregates the base data	is this part of the dplyr grammar
8	HAVING	filters aggregated data	dplyr filter()
9	SELECT	Returns final data set	dplyr select()
10	DISTINCT	Dedupe the final data set	dplyr distinct()
11	ORDER BY	Sorts the final data set	arrange()
12	TOP/LIMIT	Limits the number of rows in data set	
13	OFFSET/FETCH	Limits the number of rows in data set	

The SEQ column shows the standard order of SQL execution. One take away for this tutorial is that the SELECT clause actually executes late in the process, even though it is the first clause in the entire SELECT statement. A second take away is that SQL execution order, or tweaked order, plays a critical role in SQL query tuning.

6. SQL for View table dependencies.
7. Add cartesian join exercise.

Chapter 15

Getting metadata about and from the database (21)

This chapter demonstrates:

- What kind of data about the database is contained in a dbms
- Several methods for obtaining metadata from the dbms

The following packages are used in this chapter:

```
library(tidyverse)
library(DBI)
library(RPostgres)
library(glue)
library(here)
require(knitr)
library(dbplyr)
library(sqlpetr)
```

Assume that the Docker container with PostgreSQL and the dvdrental database are ready to go.

```
sp_docker_start("sql-pet")
```

Connect to the database:

```
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
```

15.1 Database contents and structure

After just looking at the data you seek, it might be worthwhile stepping back and looking at the big picture.

15.1.1   Database structure

For large or complex databases you need to use both the available documentation for your database (e.g., the dvdrental database) and the other empirical tools that are available. For example it's worth learning to interpret the symbols in an Entity Relationship Diagram:
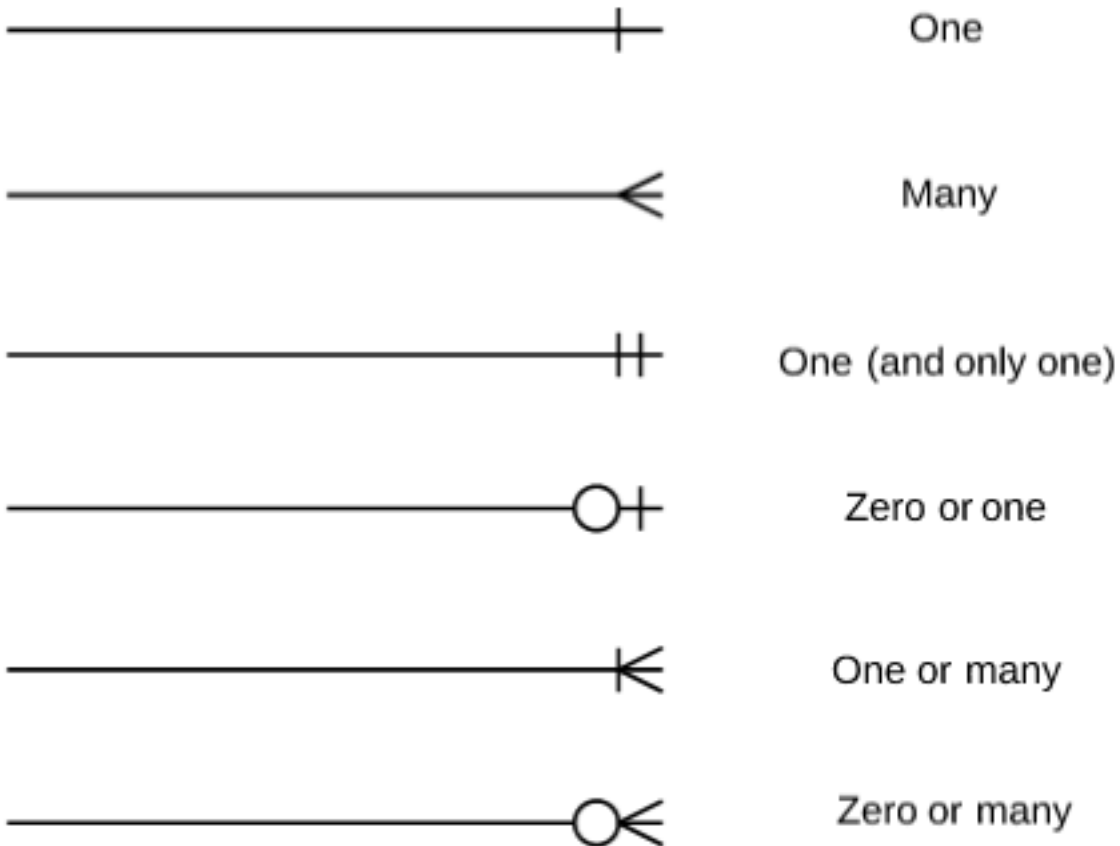
The `information_schema` is a trove of information *about* the database. Its format is more or less consistent across the different SQL implementations that are available. Here we explore some of what's available using several different methods. Postgres stores a lot of metadata.

15.1.2   Contents of the `information_schema`

For this chapter R needs the `dbplyr` package to access alternate schemas. A schema is an object that contains one or more tables. Most often there will be a default schema, but to access the metadata, you need to explicitly specify which schema contains the data you want.

15.1.3   What tables are in the database?

The simplest way to get a list of tables is with

```
table_list <- DBI::dbListTables(con)


kable(table_list)
```

x
actor_info
customer_list
film_list
nicer_but_slower_film_list
sales_by_film_category
staff
sales_by_store
staff_list
category
film_category
country
actor
language
inventory
payment
rental
city
store
film
address
film_actor
customer
smy_film

15.1.4 Digging into the `information_schema`

We usually need more detail than just a list of tables. Most SQL databases have an `information_schema` that has a standard structure to describe and control the database.

The `information_schema` is in a different schema from the default, so to connect to the `tables` table in the `information_schema` we connect to the database in a different way:

```
table_info_schema_table <- tbl(con, dbplyr::in_schema("information_schema", "tables"))
```

The `information_schema` is large and complex and contains 211 tables. So it's easy to get lost in it.

This query retrieves a list of the tables in the database that includes additional detail, not just the name of the table.

```
table_info <- table_info_schema_table %>%
  filter(table_schema == "public") %>%
  select(table_catalog, table_schema, table_name, table_type) %>%
  arrange(table_type, table_name) %>%
  collect()

kable(table_info)
```

table_catalog	table_schema	table_name	table_type
dvdrental	public	actor	BASE TABLE
dvdrental	public	address	BASE TABLE
dvdrental	public	category	BASE TABLE
dvdrental	public	city	BASE TABLE
dvdrental	public	country	BASE TABLE
dvdrental	public	customer	BASE TABLE
dvdrental	public	film	BASE TABLE
dvdrental	public	film_actor	BASE TABLE
dvdrental	public	film_category	BASE TABLE
dvdrental	public	inventory	BASE TABLE
dvdrental	public	language	BASE TABLE
dvdrental	public	payment	BASE TABLE
dvdrental	public	rental	BASE TABLE
dvdrental	public	smy_film	BASE TABLE
dvdrental	public	staff	BASE TABLE
dvdrental	public	store	BASE TABLE
dvdrental	public	actor_info	VIEW
dvdrental	public	customer_list	VIEW
dvdrental	public	film_list	VIEW
dvdrental	public	nicer_but_slower_film_list	VIEW
dvdrental	public	sales_by_film_category	VIEW
dvdrental	public	sales_by_store	VIEW
dvdrental	public	staff_list	VIEW

In this context `table_catalog` is synonymous with `database`.

Notice that *VIEWS* are composites made up of one or more *BASE TABLES*.

The SQL world has its own terminology. For example `rs` is shorthand for `result set`. That's equivalent to using `df` for a `data frame`. The following SQL query returns the same information as the previous one.

```
rs <- dbGetQuery(
  con,
  "select table_catalog, table_schema, table_name, table_type
  from information_schema.tables
  where table_schema not in ('pg_catalog','information_schema')
  order by table_type, table_name
  ;"
)
kable(rs)
```

table_catalog	table_schema	table_name	table_type
dvdrental	public	actor	BASE TABLE
dvdrental	public	address	BASE TABLE
dvdrental	public	category	BASE TABLE
dvdrental	public	city	BASE TABLE
dvdrental	public	country	BASE TABLE
dvdrental	public	customer	BASE TABLE
dvdrental	public	film	BASE TABLE
dvdrental	public	film_actor	BASE TABLE
dvdrental	public	film_category	BASE TABLE
dvdrental	public	inventory	BASE TABLE
dvdrental	public	language	BASE TABLE
dvdrental	public	payment	BASE TABLE
dvdrental	public	rental	BASE TABLE
dvdrental	public	smy_film	BASE TABLE
dvdrental	public	staff	BASE TABLE
dvdrental	public	store	BASE TABLE
dvdrental	public	actor_info	VIEW
dvdrental	public	customer_list	VIEW
dvdrental	public	film_list	VIEW
dvdrental	public	nicer_but_slower_film_list	VIEW
dvdrental	public	sales_by_film_category	VIEW
dvdrental	public	sales_by_store	VIEW
dvdrental	public	staff_list	VIEW

15.2 What columns do those tables contain?

Of course, the `DBI` package has a `dbListFields` function that provides the simplest way to get the minimum, a list of column names:

```
DBI::dbListFields(con, "rental")
```

```
## [1] "rental_id"    "rental_date"  "inventory_id" "customer_id"
## [5] "return_date"  "staff_id"     "last_update"
```

But the `information_schema` has a lot more useful information that we can use.

```
columns_info_schema_table <- tbl(con, dbplyr::in_schema("information_schema", "columns"))
```

Since the `information_schema` contains 1868 columns, we are narrowing our focus to just one table. This query retrieves more information about the `rental` table:

```
columns_info_schema_info <- columns_info_schema_table %>%
  filter(table_schema == "public") %>%
  select(
    table_catalog, table_schema, table_name, column_name, data_type, ordinal_position,
    character_maximum_length, column_default, numeric_precision, numeric_precision_radix
  ) %>%
  collect(n = Inf) %>%
  mutate(data_type = case_when(
    data_type == "character varying" ~ paste0(data_type, " (", character_maximum_length, ")"),
    data_type == "real" ~ paste0(data_type, " (", numeric_precision, ",", numeric_precision_radix, ")")
    TRUE ~ data_type
  )) %>%
  filter(table_name == "rental") %>%
```

```
  select(-table_schema, -numeric_precision, -numeric_precision_radix)

glimpse(columns_info_schema_info)

## Observations: 7
## Variables: 7
## $ table_catalog           <chr> "dvdrental", "dvdrental", "dvdrental"...
## $ table_name              <chr> "rental", "rental", "rental", "rental...
## $ column_name             <chr> "rental_id", "rental_date", "inventor...
## $ data_type               <chr> "integer", "timestamp without time zo...
## $ ordinal_position        <int> 1, 2, 3, 4, 5, 6, 7
## $ character_maximum_length <int> NA, NA, NA, NA, NA, NA, NA
## $ column_default          <chr> "nextval('rental_rental_id_seq'::regc...

kable(columns_info_schema_info)
```

table_catalog	table_name	column_name	data_type	ordinal_position	character_maximum_
dvdrental	rental	rental_id	integer	1	
dvdrental	rental	rental_date	timestamp without time zone	2	
dvdrental	rental	inventory_id	integer	3	
dvdrental	rental	customer_id	smallint	4	
dvdrental	rental	return_date	timestamp without time zone	5	
dvdrental	rental	staff_id	smallint	6	
dvdrental	rental	last_update	timestamp without time zone	7	

15.2.1   What is the difference between a `VIEW` and a `BASE TABLE`?

The `BASE TABLE` has the underlying data in the database

```
table_info_schema_table %>%
  filter(table_schema == "public" & table_type == "BASE TABLE") %>%
  select(table_name, table_type) %>%
  left_join(columns_info_schema_table, by = c("table_name" = "table_name")) %>%
  select(
    table_type, table_name, column_name, data_type, ordinal_position,
    column_default
  ) %>%
  collect(n = Inf) %>%
  filter(str_detect(table_name, "cust")) %>%
  kable()
```

table_type	table_name	column_name	data_type	ordinal_position	column_default
BASE TABLE	customer	store_id	smallint	2	NA
BASE TABLE	customer	first_name	character varying	3	NA
BASE TABLE	customer	last_name	character varying	4	NA
BASE TABLE	customer	email	character varying	5	NA
BASE TABLE	customer	address_id	smallint	6	NA
BASE TABLE	customer	active	integer	10	NA
BASE TABLE	customer	customer_id	integer	1	nextval('customer_cu
BASE TABLE	customer	activebool	boolean	7	true
BASE TABLE	customer	create_date	date	8	('now'::text)::date
BASE TABLE	customer	last_update	timestamp without time zone	9	now()

Probably should explore how the `VIEW` is made up of data from BASE TABLEs.

```r
table_info_schema_table %>%
  filter(table_schema == "public" & table_type == "VIEW") %>%
  select(table_name, table_type) %>%
  left_join(columns_info_schema_table, by = c("table_name" = "table_name")) %>%
  select(
    table_type, table_name, column_name, data_type, ordinal_position,
    column_default
  ) %>%
  collect(n = Inf) %>%
  filter(str_detect(table_name, "cust")) %>%
  kable()
```

table_type	table_name	column_name	data_type	ordinal_position	column_default
VIEW	customer_list	id	integer	1	NA
VIEW	customer_list	name	text	2	NA
VIEW	customer_list	address	character varying	3	NA
VIEW	customer_list	zip code	character varying	4	NA
VIEW	customer_list	phone	character varying	5	NA
VIEW	customer_list	city	character varying	6	NA
VIEW	customer_list	country	character varying	7	NA
VIEW	customer_list	notes	text	8	NA
VIEW	customer_list	sid	smallint	9	NA

15.2.2 What data types are found in the database?

```r
columns_info_schema_info %>% count(data_type)
```

```
## # A tibble: 3 x 2
##   data_type                     n
##   <chr>                     <int>
## 1 integer                       2
## 2 smallint                      2
## 3 timestamp without time zone   3
```

15.3 Characterizing how things are named

Names are the handle for accessing the data. Tables and columns may or may not be named consistently or in a way that makes sense to you. You should look at these names *as data*.

15.3.1 Counting columns and name reuse

Pull out some rough-and-ready but useful statistics about your database. Since we are in SQL-land we talk about variables as `columns`.

```r
public_tables <- columns_info_schema_table %>%
  filter(table_schema == "public") %>%
  collect()

public_tables %>%
  count(table_name, sort = TRUE) %>% head(n = 15) %>%
  kable()
```

table_name	n
film	13
smy_film	13
staff	11
customer	10
customer_list	9
address	8
film_list	8
nicer_but_slower_film_list	8
staff_list	8
rental	7
payment	6
actor	4
actor_info	4
city	4
inventory	4

How many *column names* are shared across tables (or duplicated)?

```
public_tables %>% count(column_name, sort = TRUE) %>% filter(n > 1)
```

```
## # A tibble: 40 x 2
##    column_name      n
##    <chr>        <int>
##  1 last_update     15
##  2 film_id          5
##  3 address_id       4
##  4 description      4
##  5 first_name       4
##  6 last_name        4
##  7 length           4
##  8 name             4
##  9 rating           4
## 10 store_id         4
## # ... with 30 more rows
```

How many column names are unique?

```
public_tables %>% count(column_name) %>% filter(n == 1) %>% count()
```

```
## # A tibble: 1 x 1
##       nn
##    <int>
## 1     18
```

15.4 Database keys

15.4.1 Direct SQL

How do we use this output? Could it be generated by dplyr?

```
rs <- dbGetQuery(
  con,
  "
--SELECT conrelid::regclass as table_from
```

```
select table_catalog||'.'||table_schema||'.'||table_name table_name
, conname, pg_catalog.pg_get_constraintdef(r.oid, true) as condef
FROM information_schema.columns c,pg_catalog.pg_constraint r
WHERE 1 = 1 --r.conrelid = '16485'
  AND r.contype  in ('f','p') ORDER BY 1
;"
)
glimpse(rs)
```

```
## Observations: 61,644
## Variables: 3
## $ table_name <chr> "dvdrental.information_schema.administrable_role_au...
## $ conname    <chr> "actor_pkey", "actor_pkey", "actor_pkey", "country_...
## $ condef     <chr> "PRIMARY KEY (actor_id)", "PRIMARY KEY (actor_id)",...
```

```
kable(head(rs))
```

table_name	conname	condef
dvdrental.information_schema.administrable_role_authorizations	actor_pkey	PRIMARY KEY (actor_id)
dvdrental.information_schema.administrable_role_authorizations	actor_pkey	PRIMARY KEY (actor_id)
dvdrental.information_schema.administrable_role_authorizations	actor_pkey	PRIMARY KEY (actor_id)
dvdrental.information_schema.administrable_role_authorizations	country_pkey	PRIMARY KEY (country_id)
dvdrental.information_schema.administrable_role_authorizations	country_pkey	PRIMARY KEY (country_id)
dvdrental.information_schema.administrable_role_authorizations	country_pkey	PRIMARY KEY (country_id)

The following is more compact and looks more useful. What is the difference between the two?

```
rs <- dbGetQuery(
  con,
  "select conrelid::regclass as table_from
      ,c.conname
      ,pg_get_constraintdef(c.oid)
  from pg_constraint c
  join pg_namespace n on n.oid = c.connamespace
 where c.contype in ('f','p')
   and n.nspname = 'public'
order by conrelid::regclass::text, contype DESC;
"
)
glimpse(rs)
```

```
## Observations: 33
## Variables: 3
## $ table_from           <chr> "actor", "address", "address", "category"...
## $ conname              <chr> "actor_pkey", "address_pkey", "fk_address...
## $ pg_get_constraintdef <chr> "PRIMARY KEY (actor_id)", "PRIMARY KEY (a...
```

```
kable(head(rs))
```

table_from	conname	pg_get_constraintdef
actor	actor_pkey	PRIMARY KEY (actor_id)
address	address_pkey	PRIMARY KEY (address_id)
address	fk_address_city	FOREIGN KEY (city_id) REFERENCES city(city_id)
category	category_pkey	PRIMARY KEY (category_id)
city	city_pkey	PRIMARY KEY (city_id)
city	fk_city	FOREIGN KEY (country_id) REFERENCES country(country_id)

```
dim(rs)[1]
```

```
## [1] 33
```

15.4.2   Database keys with dplyr

This query shows the primary and foreign keys in the database.

```
tables <- tbl(con, dbplyr::in_schema("information_schema", "tables"))
table_constraints <- tbl(con, dbplyr::in_schema("information_schema", "table_constraints"))
key_column_usage <- tbl(con, dbplyr::in_schema("information_schema", "key_column_usage"))
referential_constraints <- tbl(con, dbplyr::in_schema("information_schema", "referential_constraints"))
constraint_column_usage <- tbl(con, dbplyr::in_schema("information_schema", "constraint_column_usage"))

keys <- tables %>%
  left_join(table_constraints, by = c(
    "table_catalog" = "table_catalog",
    "table_schema" = "table_schema",
    "table_name" = "table_name"
  )) %>%
  # table_constraints %>%
  filter(constraint_type %in% c("FOREIGN KEY", "PRIMARY KEY")) %>%
  left_join(key_column_usage,
    by = c(
      "table_catalog" = "table_catalog",
      "constraint_catalog" = "constraint_catalog",
      "constraint_schema" = "constraint_schema",
      "table_name" = "table_name",
      "table_schema" = "table_schema",
      "constraint_name" = "constraint_name"
    )
  ) %>%
  # left_join(constraint_column_usage) %>% # does this table add anything useful?
  select(table_name, table_type, constraint_name, constraint_type, column_name, ordinal_position) %>%
  arrange(table_name) %>%
  collect()
glimpse(keys)
```

```
## Observations: 35
## Variables: 6
## $ table_name       <chr> "actor", "address", "address", "category", "c...
## $ table_type       <chr> "BASE TABLE", "BASE TABLE", "BASE TABLE", "BA...
## $ constraint_name  <chr> "actor_pkey", "address_pkey", "fk_address_cit...
## $ constraint_type  <chr> "PRIMARY KEY", "PRIMARY KEY", "FOREIGN KEY", ...
## $ column_name      <chr> "actor_id", "address_id", "city_id", "categor...
## $ ordinal_position <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, ...
```

```
kable(keys)
```

table_name	table_type	constraint_name	constraint_type	column_name	ordinal_pos
actor	BASE TABLE	actor_pkey	PRIMARY KEY	actor_id	
address	BASE TABLE	address_pkey	PRIMARY KEY	address_id	
address	BASE TABLE	fk_address_city	FOREIGN KEY	city_id	
category	BASE TABLE	category_pkey	PRIMARY KEY	category_id	
city	BASE TABLE	city_pkey	PRIMARY KEY	city_id	
city	BASE TABLE	fk_city	FOREIGN KEY	country_id	
country	BASE TABLE	country_pkey	PRIMARY KEY	country_id	
customer	BASE TABLE	customer_address_id_fkey	FOREIGN KEY	address_id	
customer	BASE TABLE	customer_pkey	PRIMARY KEY	customer_id	
film	BASE TABLE	film_language_id_fkey	FOREIGN KEY	language_id	
film	BASE TABLE	film_pkey	PRIMARY KEY	film_id	
film_actor	BASE TABLE	film_actor_actor_id_fkey	FOREIGN KEY	actor_id	
film_actor	BASE TABLE	film_actor_film_id_fkey	FOREIGN KEY	film_id	
film_actor	BASE TABLE	film_actor_pkey	PRIMARY KEY	actor_id	
film_actor	BASE TABLE	film_actor_pkey	PRIMARY KEY	film_id	
film_category	BASE TABLE	film_category_category_id_fkey	FOREIGN KEY	category_id	
film_category	BASE TABLE	film_category_film_id_fkey	FOREIGN KEY	film_id	
film_category	BASE TABLE	film_category_pkey	PRIMARY KEY	film_id	
film_category	BASE TABLE	film_category_pkey	PRIMARY KEY	category_id	
inventory	BASE TABLE	inventory_film_id_fkey	FOREIGN KEY	film_id	
inventory	BASE TABLE	inventory_pkey	PRIMARY KEY	inventory_id	
language	BASE TABLE	language_pkey	PRIMARY KEY	language_id	
payment	BASE TABLE	payment_customer_id_fkey	FOREIGN KEY	customer_id	
payment	BASE TABLE	payment_pkey	PRIMARY KEY	payment_id	
payment	BASE TABLE	payment_rental_id_fkey	FOREIGN KEY	rental_id	
payment	BASE TABLE	payment_staff_id_fkey	FOREIGN KEY	staff_id	
rental	BASE TABLE	rental_customer_id_fkey	FOREIGN KEY	customer_id	
rental	BASE TABLE	rental_inventory_id_fkey	FOREIGN KEY	inventory_id	
rental	BASE TABLE	rental_pkey	PRIMARY KEY	rental_id	
rental	BASE TABLE	rental_staff_id_key	FOREIGN KEY	staff_id	
staff	BASE TABLE	staff_address_id_fkey	FOREIGN KEY	address_id	
staff	BASE TABLE	staff_pkey	PRIMARY KEY	staff_id	
store	BASE TABLE	store_address_id_fkey	FOREIGN KEY	address_id	
store	BASE TABLE	store_manager_staff_id_fkey	FOREIGN KEY	manager_staff_id	
store	BASE TABLE	store_pkey	PRIMARY KEY	store_id	

What do we learn from the following query? How is it useful?

```
rs <- dbGetQuery(
  con,
  "SELECT r.*,
  pg_catalog.pg_get_constraintdef(r.oid, true) as condef
  FROM pg_catalog.pg_constraint r
  WHERE 1=1 --r.conrelid = '16485' AND r.contype = 'f' ORDER BY 1;
  "
)


head(rs)


##                       conname connamespace contype condeferrable
## 1 cardinal_number_domain_check        12703       c         FALSE
## 2              yes_or_no_check        12703       c         FALSE
## 3                   year_check         2200       c         FALSE
```

```
## 4                      actor_pkey          2200         p          FALSE
## 5                    address_pkey          2200         p          FALSE
## 6                   category_pkey          2200         p          FALSE
##    condeferred convalidated conrelid contypid conindid confrelid
## 1       FALSE          TRUE        0    12716        0         0
## 2       FALSE          TRUE        0    12724        0         0
## 3       FALSE          TRUE        0    16397        0         0
## 4       FALSE          TRUE    16420        0    16555         0
## 5       FALSE          TRUE    16461        0    16557         0
## 6       FALSE          TRUE    16427        0    16559         0
##    confupdtype confdeltype confmatchtype conislocal coninhcount
## 1                                             TRUE           0
## 2                                             TRUE           0
## 3                                             TRUE           0
## 4                                             TRUE           0
## 5                                             TRUE           0
## 6                                             TRUE           0
##    connoinherit conkey confkey conpfeqop conppeqop conffeqop conexclop
## 1         FALSE   <NA>    <NA>      <NA>      <NA>      <NA>      <NA>
## 2         FALSE   <NA>    <NA>      <NA>      <NA>      <NA>      <NA>
## 3         FALSE   <NA>    <NA>      <NA>      <NA>      <NA>      <NA>
## 4          TRUE    {1}    <NA>      <NA>      <NA>      <NA>      <NA>
## 5          TRUE    {1}    <NA>      <NA>      <NA>      <NA>      <NA>
## 6          TRUE    {1}    <NA>      <NA>      <NA>      <NA>      <NA>
##
## 1
## 2 {SCALARARRAYOPEXPR :opno 98 :opfuncid 67 :useOr true :inputcollid 100 :args ({RELABELTYPE :arg {COERCET
## 3                                                                       {BOOLEXPR :boolop and :args
## 4
## 5
## 6
##                                                                                consrc
## 1                                                                           (VALUE >= 0)
## 2 ((VALUE)::text = ANY ((ARRAY['YES'::character varying, 'NO'::character varying])::text[]))
## 3                                                           ((VALUE >= 1901) AND (VALUE <= 2155))
## 4                                                                                 <NA>
## 5                                                                                 <NA>
## 6                                                                                 <NA>
##                                                                                condef
## 1                                                                     CHECK (VALUE >= 0)
## 2 CHECK (VALUE::text = ANY (ARRAY['YES'::character varying, 'NO'::character varying]::text[]))
## 3                                                       CHECK (VALUE >= 1901 AND VALUE <= 2155)
## 4                                                                      PRIMARY KEY (actor_id)
## 5                                                                    PRIMARY KEY (address_id)
## 6                                                                   PRIMARY KEY (category_id)
```

15.5   Creating your own data dictionary

If you are going to work with a database for an extended period it can be useful to create your own data dictionary. This can take the form of keeping detaild notes as well as extracting metadata from the dbms. Here is an illustration of the idea.

```r
some_tables <- c("rental", "city", "store")

all_meta <- map_df(some_tables, sp_get_dbms_data_dictionary, con = con)

all_meta
```

```
## # A tibble: 15 x 11
##    table_name var_name var_type num_rows num_blank num_unique min   q_25
##    <chr>      <chr>    <chr>        <int>     <int>      <int> <chr> <chr>
##  1 rental     rental_~ integer      16044         0      16044 1     4013
##  2 rental     rental_~ double       16044         0      15815 2005~ 2005~
##  3 rental     invento~ integer      16044         0       4580 1     1154
##  4 rental     custome~ integer      16044         0        599 1     148
##  5 rental     return_~ double       16044       183      15836 2005~ 2005~
##  6 rental     staff_id integer      16044         0          2 1     1
##  7 rental     last_up~ double       16044         0          3 2006~ 2006~
##  8 city       city_id  integer        600         0        600 1     150
##  9 city       city     charact~       600         0        599 A Co~ Dzer~
## 10 city       country~ integer        600         0        109 1     28
## 11 city       last_up~ double         600         0          1 2006~ 2006~
## 12 store      store_id integer          2         0          2 1     1
## 13 store      manager~ integer          2         0          2 1     1
## 14 store      address~ integer          2         0          2 1     1
## 15 store      last_up~ double           2         0          1 2006~ 2006~
## # ... with 3 more variables: q_50 <chr>, q_75 <chr>, max <chr>
```

```r
glimpse(all_meta)
```

```
## Observations: 15
## Variables: 11
## $ table_name <chr> "rental", "rental", "rental", "rental", "rental", "...
## $ var_name   <chr> "rental_id", "rental_date", "inventory_id", "custom...
## $ var_type   <chr> "integer", "double", "integer", "integer", "double"...
## $ num_rows   <int> 16044, 16044, 16044, 16044, 16044, 16044, 16044, 60...
## $ num_blank  <int> 0, 0, 0, 0, 183, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0
## $ num_unique <int> 16044, 15815, 4580, 599, 15836, 2, 3, 600, 599, 109...
## $ min        <chr> "1", "2005-05-24 22:53:30", "1", "1", "2005-05-25 2...
## $ q_25       <chr> "4013", "2005-07-07 00:58:00", "1154", "148", "2005...
## $ q_50       <chr> "8025", "2005-07-28 16:03:27", "2291", "296", "2005...
## $ q_75       <chr> "12037", "2005-08-17 21:13:35", "3433", "446", "200...
## $ max        <chr> "16049", "2006-02-14 15:16:03", "4581", "599", "200...
```

```r
kable(head(all_meta))
```

table_name	var_name	var_type	num_rows	num_blank	num_unique	min	q_25
rental	rental_id	integer	16044	0	16044	1	4013
rental	rental_date	double	16044	0	15815	2005-05-24 22:53:30	2005-07-07 00:
rental	inventory_id	integer	16044	0	4580	1	1154
rental	customer_id	integer	16044	0	599	1	148
rental	return_date	double	16044	183	15836	2005-05-25 23:55:21	2005-07-10 15:
rental	staff_id	integer	16044	0	2	1	1

15.6   Save your work!

The work you do to understand the structure and contents of a database can be useful for others (including future-you). So at the end of a session, you might look at all the data frames you want to save. Consider saving them in a form where you can add notes at the appropriate level (as in a Google Doc representing table or columns that you annotate over time).

```
ls()
```

```
##  [1] "all_meta"                "columns_info_schema_info"
##  [3] "columns_info_schema_table" "con"
##  [5] "constraint_column_usage" "cranex"
##  [7] "key_column_usage"        "keys"
##  [9] "public_tables"           "referential_constraints"
## [11] "rs"                      "some_tables"
## [13] "table_constraints"       "table_info"
## [15] "table_info_schema_table" "table_list"
## [17] "tables"
```

Chapter 16

Drilling into your DBMS environment (22)

This chapter investigates:

- Elements of the database environment
- Differences between a database, a schema, and other objects
- Exercises

The following packages are used in this chapter:

```r
# These packages are called in almost every chapter of the book:
library(tidyverse)
library(DBI)
library(RPostgres)
require(knitr)
library(dbplyr)
library(sqlpetr)

display_rows <-  15 # as a default, show 15 rows
```

Start up the `docker-pet` container

```r
sp_docker_start("sql-pet")
```

Now connect to the `dvdrental` database with R

```r
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
con
```

```
## <PqConnection> dvdrental@localhost:5432
```

16.1   Which database?

Your DBA will create your user accounts and priviledges for the database(s) that you can access.

One of the challenges when working with a database(s) is finding where your data actually resides. Your best resources will be one or more subject matter experts, `SME`, and your DBA. Your data may actually reside in multiple databases, e.g., a detail and summary databases. In our tutorial, we focus on the one database, `dvdrental`. Database names usually reflect something about the data that they contain.

Your laptop is a server for the Docker Postgres databases. A database is a collection of files that Postgres manages in the background.

16.2   How many databases reside in the Docker Container?

```
rs <-
  DBI::dbGetQuery(
    con,
    "SELECT 'DB Names in Docker' showing
           ,datname DB
     FROM pg_database
    WHERE datistemplate = false;
    "
  )
kable(rs)
```

showing	db
DB Names in Docker	postgres
DB Names in Docker	dvdrental

Which databases are available?

Modify the connection call to connect to the `postgres` database.

```
# this code chunk is not evaluated because the `dbname` is not valid!
con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "your code goes here",
  seconds_to_test = 10
)

con
if (con != "There is no connection") {
  dbDisconnect(con)
}

# Answer: con <PqConnection> postgres@localhost:5432
```

```
# Reconnect to dvdrental

con <- sp_get_postgres_connection(
  user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
  dbname = "dvdrental",
  seconds_to_test = 10
)
con
```

```
## <PqConnection> dvdrental@localhost:5432
```

Note that the two Sys.getenv function calls work in this tutorial because both the user and password are available in both databases. This is a common practice in organinzations that have implemented single sign on across their organization.

```
Gotcha:
```

```
If one has data in multiple databases or multiple environments, Development,
Integration, and Prodution, it is very easy to connect to the wrong database in
the wrong environment.  Always double check your connection information when
logging in and before performing any inserts, updates, or deletes against the
database.
```

The following code block should be used to reduce propagating the above gotcha. Current_database(), CURRENT_DATE or CURRENT_TIMESTAMP, and 'result set' are the most useful and last three not so much. Instead of the host IP address having the actual hostname would be a nice addition.

```r
rs1 <-
  DBI::dbGetQuery(
    con,
    "SELECT current_database() DB
         ,CURRENT_DATE
         ,CURRENT_TIMESTAMP
         ,'result set description' showing
         ,session_user
         ,inet_server_addr() host
         ,inet_server_port() port
    "
  )
kable(rs1)
```

db	current_date	current_timestamp	showing	session_user	host	port
dvdrental	2018-12-18	2018-12-18 12:40:06	result set description	postgres	172.17.0.2	5432

Since we will only be working in the `dvdrental` database in this tutorial and reduce the number of output columns shown, only the 'result set description' will be used.

16.3 Which Schema?

In the code block below, we look at the `information_schema.table` which contains information about all the schemas and table/views within our dvdrental database. Databases can have one or more schemas, containers that hold tables or views. Schemas partition the database into big logical blocks of related data. Schema names usually reflect an application or logically related datasets. Occasionally a DBA will set up a new schema and use a users name.

What schemas are in the `dvdrental` database? How many entries are in each schema?

```r
## Database Schemas
#
rs1 <-
  DBI::dbGetQuery(
    con,
    "SELECT 'DB Schemas' showing,t.table_catalog DB,t.table_schema,COUNT(*) tbl_vws
     FROM information_schema.tables t
    GROUP BY t.table_catalog,t.table_schema
    "
  )
```

```
kable(rs1)
```

showing	db	table_schema	tbl_vws
DB Schemas	dvdrental	pg_catalog	121
DB Schemas	dvdrental	public	23
DB Schemas	dvdrental	information_schema	67

We see that there are three schemas. The pg_catalog is the standard PostgreSQL meta data and core schema. Postgres uses this schema to manage the internal workings of the database. DBA's are the primary users of pg_catalog. We used the pg_catalog schema to answer the question 'How many databases reside in the Docker Container?', but normally the data analyst is not interested in analyzing database data.

The information_schema contains ANSI standardized views used across the different SQL vendors, (Oracle, Sysbase, MS SQL Server, IBM DB2, etc). The information_schema contains a plethora of metadata that will help you locate your data tables, understand the relationships between the tables, and write efficient SQL queries.

16.4   Exercises

```
#
# Add an order by clause to order the output by the table catalog.
rs1 <- DBI::dbGetQuery(con, "SELECT '1. ORDER BY table_catalog' showing
                          ,t.table_catalog DB,t.table_schema,COUNT(*) tbl_vws
                         FROM information_schema.tables t
                         GROUP BY t.table_catalog,t.table_schema
                         ")
kable(rs1)
```

showing	db	table_schema	tbl_vws
1. ORDER BY table_catalog	dvdrental	pg_catalog	121
1. ORDER BY table_catalog	dvdrental	public	23
1. ORDER BY table_catalog	dvdrental	information_schema	67

```
# Add an order by clause to order the output by tbl_vws in descending order.
rs2 <- DBI::dbGetQuery(con, "SELECT '2. ORDER BY tbl_vws desc' showing
                          ,t.table_catalog DB,t.table_schema,COUNT(*) tbl_vws
                         FROM information_schema.tables t
                         GROUP BY t.table_catalog,t.table_schema
                         ")
kable(rs2)
```

showing	db	table_schema	tbl_vws
2. ORDER BY tbl_vws desc	dvdrental	pg_catalog	121
2. ORDER BY tbl_vws desc	dvdrental	public	23
2. ORDER BY tbl_vws desc	dvdrental	information_schema	67

```
# Complete the SQL statement to show everything about all the tables.

rs3 <- DBI::dbGetQuery(con, "SELECT '3. all information_schema tables' showing
                          ,'your code goes here'
                         FROM information_schema.tables t
                         ")
kable(head(rs3, display_rows))
```

showing	?column?
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here
3. all information_schema tables	your code goes here

```r
# Use the results from above to pull interesting columns from just the information_schema
rs4 <- DBI::dbGetQuery(con, "SELECT '4. information_schema.tables' showing
                                    ,'your code goes here'
                              FROM information_schema.tables t
                              where 'your code goes here' = 'your code goes here'
                              ")
head(rs4, display_rows)
```

```
##                        showing            ?column?
## 1  4. information_schema.tables your code goes here
## 2  4. information_schema.tables your code goes here
## 3  4. information_schema.tables your code goes here
## 4  4. information_schema.tables your code goes here
## 5  4. information_schema.tables your code goes here
## 6  4. information_schema.tables your code goes here
## 7  4. information_schema.tables your code goes here
## 8  4. information_schema.tables your code goes here
## 9  4. information_schema.tables your code goes here
## 10 4. information_schema.tables your code goes here
## 11 4. information_schema.tables your code goes here
## 12 4. information_schema.tables your code goes here
## 13 4. information_schema.tables your code goes here
## 14 4. information_schema.tables your code goes here
## 15 4. information_schema.tables your code goes here
```

```r
# Modify the SQL below with your interesting column names.
# Update the where clause to return only rows from the information schema and begin with 'tab'
rs5 <- DBI::dbGetQuery(con, "SELECT '5. information_schema.tables' showing
                                    ,'your code goes here'
                              FROM information_schema.tables t
                              where 'your code goes here' = 'your code goes here'
                              ")
kable(head(rs5, display_rows))
```

showing	?column?
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here
5. information_schema.tables	your code goes here

```r
# Modify the SQL below with your interesting column names.
# Update the where clause to return only rows from the information schema and begin with 'col'
rs6 <- DBI::dbGetQuery(con, "SELECT '6. information_schema.tables' showing
                                  ,'your code goes here'
                              FROM information_schema.tables t
                             where 'your code goes here' = 'your code goes here'
                            ")
kable(head(rs6, display_rows))
```

showing	?column?
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here
6. information_schema.tables	your code goes here

In the next exercise we combine both the table and column output from the previous exercises. Review the following code block. The last two lines of the WHERE clause are swithced. Will the result set be the same or different? Execute the code block and review the two datasets.

```r
rs7 <- DBI::dbGetQuery(con, "SELECT '7. information_schema.tables' showing
                                  ,table_catalog||'.'||table_schema db_info, table_name, table_type
                              FROM information_schema.tables t
                             where table_schema = 'information_schema'
                               and table_name like 'table%' OR table_name like '%col%'
                               and table_type = 'VIEW'
                            ")
kable(head(rs7, display_rows))
```

showing	db_info	table_name	table_type
7. information_schema.tables	dvdrental.information_schema	collations	VIEW
7. information_schema.tables	dvdrental.information_schema	collation_character_set_applicability	VIEW
7. information_schema.tables	dvdrental.information_schema	column_domain_usage	VIEW
7. information_schema.tables	dvdrental.information_schema	column_privileges	VIEW
7. information_schema.tables	dvdrental.information_schema	column_udt_usage	VIEW
7. information_schema.tables	dvdrental.information_schema	columns	VIEW
7. information_schema.tables	dvdrental.information_schema	constraint_column_usage	VIEW
7. information_schema.tables	dvdrental.information_schema	key_column_usage	VIEW
7. information_schema.tables	dvdrental.information_schema	role_column_grants	VIEW
7. information_schema.tables	dvdrental.information_schema	table_constraints	VIEW
7. information_schema.tables	dvdrental.information_schema	table_privileges	VIEW
7. information_schema.tables	dvdrental.information_schema	tables	VIEW
7. information_schema.tables	dvdrental.information_schema	triggered_update_columns	VIEW
7. information_schema.tables	dvdrental.information_schema	view_column_usage	VIEW
7. information_schema.tables	dvdrental.information_schema	_pg_foreign_table_columns	VIEW

```
rs8 <- DBI::dbGetQuery(con, "SELECT '8. information_schema.tables' showing
                    ,table_catalog||'.'||table_schema db_info, table_name, table_type
                 FROM information_schema.tables t
               where table_schema = 'information_schema'
                 and table_type = 'VIEW'
                 and table_name like 'table%' OR table_name like '%col%'
               ")
kable(head(rs8, display_rows))
```

showing	db_info	table_name	table_type
8. information_schema.tables	dvdrental.information_schema	column_options	VIEW
8. information_schema.tables	dvdrental.information_schema	_pg_foreign_table_columns	VIEW
8. information_schema.tables	dvdrental.information_schema	view_column_usage	VIEW
8. information_schema.tables	dvdrental.information_schema	triggered_update_columns	VIEW
8. information_schema.tables	dvdrental.information_schema	tables	VIEW
8. information_schema.tables	dvdrental.information_schema	table_privileges	VIEW
8. information_schema.tables	dvdrental.information_schema	table_constraints	VIEW
8. information_schema.tables	dvdrental.information_schema	role_column_grants	VIEW
8. information_schema.tables	dvdrental.information_schema	key_column_usage	VIEW
8. information_schema.tables	dvdrental.information_schema	constraint_column_usage	VIEW
8. information_schema.tables	dvdrental.information_schema	columns	VIEW
8. information_schema.tables	dvdrental.information_schema	column_udt_usage	VIEW
8. information_schema.tables	dvdrental.information_schema	column_privileges	VIEW
8. information_schema.tables	dvdrental.information_schema	column_domain_usage	VIEW
8. information_schema.tables	dvdrental.information_schema	collation_character_set_applicability	VIEW

Operator/Element	Associativity	Description
.	left	table/column name separator
::	left	PostgreSQL-style typecast
[ ]	left	array element selection
-	right	unary minus
^	left	exponentiation
/ %	left	multiplication, division, modulo
+ -	left	addition, subtraction
IS		IS TRUE, IS FALSE, IS UNKNOWN, IS NULL
ISNULL		test for null
NOTNULL		test for not null

Operator/Element	Associativity	Description
(any other)	left	all other native and user-defined operators
IN		set membership
BETWEEN		range containment
OVERLAPS		time interval overlap
LIKE ILIKE SIMILAR		string pattern matching
< >		less than, greater than
=	right	equality, assignment
NOT	right	logical negation
AND	left	logical conjunction
OR	left	logical disjunction

```r
rs1 <- DBI::dbGetQuery(con, "SELECT t.table_catalog DB ,t.table_schema
                      ,t.table_name,t.table_type
                    FROM information_schema.tables t")

rs2 <- DBI::dbGetQuery(con, "SELECT t.table_catalog DB ,t.table_schema
                      ,t.table_type,COUNT(*) tbls
                    FROM information_schema.tables t
                  group by t.table_catalog ,t.table_schema
                      ,t.table_type
                  ")

rs3 <- DBI::dbGetQuery(con, "SELECT distinct t.table_catalog DB ,t.table_schema
                      ,t.table_type tbls
                    FROM information_schema.tables t
                  ")
```

```r
# kable(head(rs1 %>% arrange (table_name)))
# View(rs1)
# View(rs2)
# View(rs3)
kable(head(rs1))
```

db	table_schema	table_name	table_type
dvdrental	public	actor_info	VIEW
dvdrental	public	customer_list	VIEW
dvdrental	public	film_list	VIEW
dvdrental	public	nicer_but_slower_film_list	VIEW
dvdrental	public	sales_by_film_category	VIEW
dvdrental	public	staff	BASE TABLE

```r
kable(head(rs2))
```

db	table_schema	table_type	tbls
dvdrental	information_schema	BASE TABLE	7
dvdrental	information_schema	VIEW	60
dvdrental	pg_catalog	BASE TABLE	62
dvdrental	public	BASE TABLE	16
dvdrental	public	VIEW	7
dvdrental	pg_catalog	VIEW	59

```
kable(head(rs3))
```

db	table_schema	tbls
dvdrental	information_schema	BASE TABLE
dvdrental	information_schema	VIEW
dvdrental	pg_catalog	BASE TABLE
dvdrental	public	BASE TABLE
dvdrental	public	VIEW
dvdrental	pg_catalog	VIEW

www.dataquest.io/blog/postgres-internals

Comment on the practice of putting a comma at the beginning of a line in SQL code.

```
## Explain a `dplyr::join

tbl_pk_fk_df <- DBI::dbGetQuery(
  con,
  "
SELECT --t.table_catalog,t.table_schema,
   c.table_name
    ,kcu.column_name
    ,c.constraint_name
    ,c.constraint_type
    ,coalesce(c2.table_name, '') ref_table
    ,coalesce(kcu2.column_name, '') ref_table_col
FROM information_schema.tables t
LEFT JOIN information_schema.table_constraints c
  ON t.table_catalog = c.table_catalog
    AND t.table_schema = c.table_schema
    AND t.table_name = c.table_name
LEFT JOIN information_schema.key_column_usage kcu
    ON c.constraint_schema = kcu.constraint_schema
        AND c.constraint_name = kcu.constraint_name
LEFT JOIN information_schema.referential_constraints rc
    ON c.constraint_schema = rc.constraint_schema
        AND c.constraint_name = rc.constraint_name
LEFT JOIN information_schema.table_constraints c2
    ON rc.unique_constraint_schema = c2.constraint_schema
        AND rc.unique_constraint_name = c2.constraint_name
LEFT JOIN information_schema.key_column_usage kcu2
    ON c2.constraint_schema = kcu2.constraint_schema
        AND c2.constraint_name = kcu2.constraint_name
        AND kcu.ordinal_position = kcu2.ordinal_position
WHERE c.constraint_type IN ('PRIMARY KEY', 'FOREIGN KEY')
  AND c.table_catalog = 'dvdrental'
    AND c.table_schema = 'public'
ORDER BY c.table_name;
"
)

# View(tbl_pk_fk_df)

tables_df <- tbl_pk_fk_df %>% distinct(table_name)
# View(tables_df)
```

```r
library(DiagrammeR)

table_nodes_ndf <- create_node_df(
  n <- nrow(tables_df)
  , type <- "table"
  , label <- tables_df$table_name
  ,
  shape = "rectangle"
  , width = 1
  , height = .5
  , fontsize = 18
)

tbl_pk_fk_ids_df <- inner_join(tbl_pk_fk_df, table_nodes_ndf
  ,
  by = c("table_name" = "label")
  , suffix(c("st", "s"))
) %>%
  rename("src_tbl_id" = id) %>%
  left_join(table_nodes_ndf
    ,
    by = c("ref_table" = "label")
    , suffix(c("st", "t"))
  ) %>%
  rename("fk_tbl_id" = id)

tbl_fk_df <- tbl_pk_fk_ids_df %>% filter(constraint_type == "FOREIGN KEY")
tbl_pk_df <- tbl_pk_fk_ids_df %>% filter(constraint_type == "PRIMARY KEY")
# View(tbl_pk_fk_ids_df)
# View(tbl_fk_df)
# View(tbl_pk_df)
kable(head(tbl_fk_df))
```

table_name	column_name	constraint_name	constraint_type	ref_table	ref_table_col	src_tbl_id
address	city_id	fk_address_city	FOREIGN KEY	city	city_id	2
city	country_id	fk_city	FOREIGN KEY	country	country_id	4
customer	address_id	customer_address_id_fkey	FOREIGN KEY	address	address_id	6
film	language_id	film_language_id_fkey	FOREIGN KEY	language	language_id	7
film_actor	actor_id	film_actor_actor_id_fkey	FOREIGN KEY	actor	actor_id	8
film_actor	film_id	film_actor_film_id_fkey	FOREIGN KEY	film	film_id	8

```r
kable(head(tbl_pk_df))
```

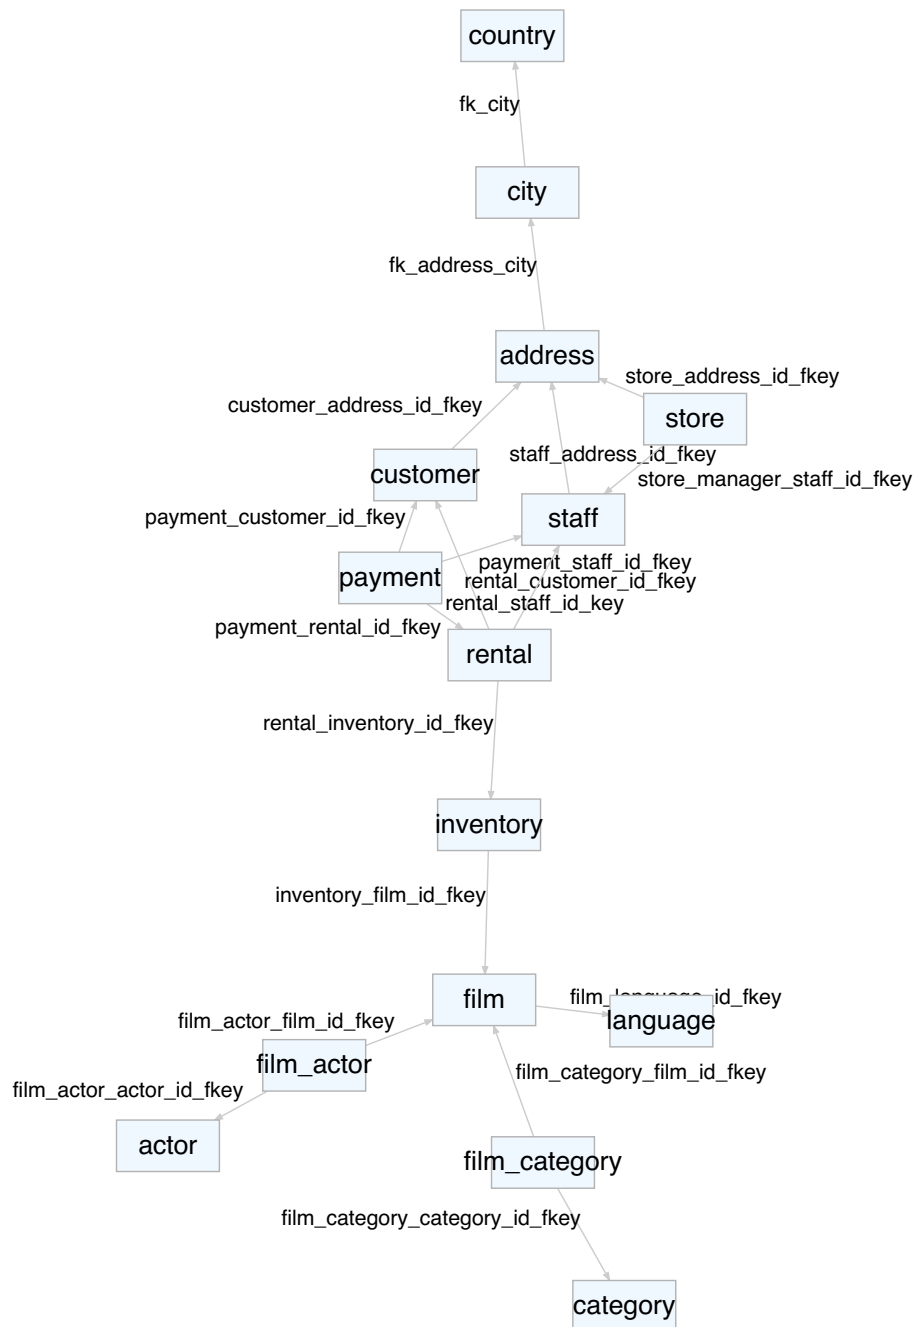| table_name	column_name	constraint_name	constraint_type	ref_table	ref_table_col	src_tbl_id	type.x	
actor	actor_id	actor_pkey	PRIMARY KEY			1	table	
address	address_id	address_pkey	PRIMARY KEY			2	table	
category	category_id	category_pkey	PRIMARY KEY			3	table	
city	city_id	city_pkey	PRIMARY KEY			4	table	
country	country_id	country_pkey	PRIMARY KEY			5	table	
customer	customer_id	customer_pkey	PRIMARY KEY			6	table	

```r
# Create an edge data frame, edf

fk_edf <-
  create_edge_df(
```

```r
    from = tbl_fk_df$src_tbl_id,
    to = tbl_fk_df$fk_tbl_id,
    rel = "fk",
    label = tbl_fk_df$constraint_name,
    fontsize = 15
  )
# View(fk_edf)
```

```r
create_graph(
  nodes_df = table_nodes_ndf,
  edges_df = fk_edf,
  graph_name = "Simple FK Graph"
) %>%
render_graph()
```

country

fk_city

city

fk_address_city

address

store_address_id_fkey

store

customer_address_id_fkey

customer

staff_address_id_fkey

store_manager_staff_id_fkey

staff

payment_customer_id_fkey

payment

payment_staff_id_fkey
rental_customer_id_fkey
rental_staff_id_key

payment_rental_id_fkey

rental

rental_inventory_id_fkey

inventory

inventory_film_id_fkey

film

film_language_id_fkey

language

film_actor_film_id_fkey

film_actor

film_category_film_id_fkey

film_actor_actor_id_fkey

actor

film_category

film_category_category_id_fkey

category

```
dbDisconnect(con)
# system2('docker','stop sql-pet')
```

Chapter 17

Explain queries (71)

This chapter demonstrates:

- How to investigate SQL query performance

```r
# These packages are called in almost every chapter of the book:
library(tidyverse)
library(DBI)
library(RPostgres)
library(glue)
library(here)
require(knitr)
library(dbplyr)
library(sqlpetr)
```

- examining `dplyr` queries (`dplyr::show_query` on the R side v EXPLAIN on the PostgreSQL side)

Start up the `docker-pet` container

```r
sp_docker_start("sql-pet")
```

now connect to the database with R

```r
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                         password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                         dbname = "dvdrental",
                         seconds_to_test = 10)
```

17.1   Performance considerations

```r
## Explain a `dplyr::join`

## Explain the quivalent SQL join
rs1 <- DBI::dbGetQuery(con
                ,"SELECT c.*
                    FROM pg_catalog.pg_class c
                    JOIN pg_catalog.pg_namespace n ON n.oid = c.relnamespace
                   WHERE  n.nspname = 'public'
                     AND  c.relname = 'cust_movies'
                     AND  c.relkind = 'r'
```

```
                      ;
                      "
                  )
head(rs1)
```

```
##   [1] relname            relnamespace      reltype
##   [4] reloftype          relowner          relam
##   [7] relfilenode        reltablespace     relpages
##  [10] reltuples          relallvisible     reltoastrelid
##  [13] relhasindex        relisshared       relpersistence
##  [16] relkind            relnatts          relchecks
##  [19] relhasoids         relhaspkey        relhasrules
##  [22] relhastriggers     relhassubclass    relrowsecurity
##  [25] relforcerowsecurity relispopulated   relreplident
##  [28] relispartition     relfrozenxid      relminmxid
##  [31] relacl             reloptions        relpartbound
## <0 rows> (or 0-length row.names)
```

This came from 14-sql_pet-examples-part-b.Rmd

```
rs1 <- DBI::dbGetQuery(con,
            "explain select r.*
              from rental r
            ;"
            )
head(rs1)
```

```
##                                                 QUERY PLAN
## 1 Seq Scan on rental r  (cost=0.00..310.44 rows=16044 width=36)
```

```
rs2 <- DBI::dbGetQuery(con,
            "explain select count(*) count
              from rental r
                  left outer join payment p
                    on r.rental_id = p.rental_id
              where p.rental_id is null
            ;")
head(rs2)
```

```
##                                                            QUERY PLAN
## 1                          Aggregate  (cost=2086.78..2086.80 rows=1 width=8)
## 2                        -> Merge Anti Join  (cost=0.57..2066.73 rows=8022 width=0)
## 3                              Merge Cond: (r.rental_id = p.rental_id)
## 4       -> Index Only Scan using rental_pkey on rental r  (cost=0.29..1024.95 rows=16044 width=4)
## 5     -> Index Only Scan using idx_fk_rental_id on payment p  (cost=0.29..819.23 rows=14596 width=4)
```

```
rs3 <- DBI::dbGetQuery(con,
            "explain select sum(f.rental_rate) open_amt,count(*) count
              from rental r
                  left outer join payment p
                    on r.rental_id = p.rental_id
                  join inventory i
                    on r.inventory_id = i.inventory_id
                  join film f
                    on i.film_id = f.film_id
              where p.rental_id is null
            ;")
```

```
head(rs3)
```

```
##                                              QUERY PLAN
## 1                     Aggregate  (cost=2353.64..2353.65 rows=1 width=40)
## 2                 -> Hash Join  (cost=205.14..2313.53 rows=8022 width=12)
## 3                         Hash Cond: (i.film_id = f.film_id)
## 4               -> Hash Join  (cost=128.64..2215.88 rows=8022 width=2)
## 5                     Hash Cond: (r.inventory_id = i.inventory_id)
## 6           -> Merge Anti Join  (cost=0.57..2066.73 rows=8022 width=4)
```

```r
rs4 <- DBI::dbGetQuery(con,
              "explain select c.customer_id,c.first_name,c.last_name,sum(f.rental_rate) open_amt,count
                from rental r
                    left outer join payment p
                      on r.rental_id = p.rental_id
                    join inventory i
                      on r.inventory_id = i.inventory_id
                    join film f
                      on i.film_id = f.film_id
                    join customer c
                      on r.customer_id = c.customer_id
                 where p.rental_id is null
                 group by c.customer_id,c.first_name,c.last_name
                 order by open_amt desc
                ;"
              )
head(rs4)
```

```
##                                              QUERY PLAN
## 1              Sort  (cost=2452.49..2453.99 rows=599 width=260)
## 2                         Sort Key: (sum(f.rental_rate)) DESC
## 3     -> HashAggregate  (cost=2417.37..2424.86 rows=599 width=260)
## 4                             Group Key: c.customer_id
## 5         -> Hash Join  (cost=227.62..2357.21 rows=8022 width=232)
## 6                     Hash Cond: (r.customer_id = c.customer_id)
```

17.2 Clean up

```r
# dbRemoveTable(con, "cars")
# dbRemoveTable(con, "mtcars")
# dbRemoveTable(con, "cust_movies")

# diconnect from the db
dbDisconnect(con)

sp_docker_stop("sql-pet")
```

Chapter 18

SQL queries behind the scenes (72)

This chapter explains:

- Some details about how SQL queries work behind the scenes
- SQL queries are executed behind the scenes
- You can pass values to SQL queries

```r
# These packages are called in almost every chapter of the book:
library(tidyverse)
library(DBI)
library(RPostgres)
library(glue)
library(here)
require(knitr)
library(dbplyr)
library(sqlpetr)
```

Start up the `docker-pet` container

```r
sp_docker_start("sql-pet")
```

now connect to the database with R

```r
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                         password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                         dbname = "dvdrental",
                         seconds_to_test = 10)
```

18.1 SQL Execution Steps

- Parse the incoming SQL query
- Compile the SQL query
- Plan/optimize the data acquisition path
- Execute the optimized query / acquire and return data

```r
dbWriteTable(con, "mtcars", mtcars, overwrite = TRUE)
rs <- dbSendQuery(con, "SELECT * FROM mtcars WHERE cyl = 4")
dbFetch(rs)
```

```
##     mpg cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1
```

```
## 2   24.4    4 146.7   62 3.69 3.190 20.00   1  0     4     2
## 3   22.8    4 140.8   95 3.92 3.150 22.90   1  0     4     2
## 4   32.4    4  78.7   66 4.08 2.200 19.47   1  1     4     1
## 5   30.4    4  75.7   52 4.93 1.615 18.52   1  1     4     2
## 6   33.9    4  71.1   65 4.22 1.835 19.90   1  1     4     1
## 7   21.5    4 120.1   97 3.70 2.465 20.01   1  0     3     1
## 8   27.3    4  79.0   66 4.08 1.935 18.90   1  1     4     1
## 9   26.0    4 120.3   91 4.43 2.140 16.70   0  1     5     2
## 10  30.4    4  95.1  113 3.77 1.513 16.90   1  1     5     2
## 11  21.4    4 121.0  109 4.11 2.780 18.60   1  1     4     2
```

```
dbClearResult(rs)
```

18.2   Passing values to SQL statements

```
#Pass one set of values with the param argument:
rs <- dbSendQuery(con,"SELECT * FROM mtcars WHERE cyl = 4")
dbFetch(rs)
```

```
##      mpg cyl  disp  hp drat    wt  qsec vs am gear carb
## 1   22.8    4 108.0   93 3.85 2.320 18.61   1  1     4     1
## 2   24.4    4 146.7   62 3.69 3.190 20.00   1  0     4     2
## 3   22.8    4 140.8   95 3.92 3.150 22.90   1  0     4     2
## 4   32.4    4  78.7   66 4.08 2.200 19.47   1  1     4     1
## 5   30.4    4  75.7   52 4.93 1.615 18.52   1  1     4     2
## 6   33.9    4  71.1   65 4.22 1.835 19.90   1  1     4     1
## 7   21.5    4 120.1   97 3.70 2.465 20.01   1  0     3     1
## 8   27.3    4  79.0   66 4.08 1.935 18.90   1  1     4     1
## 9   26.0    4 120.3   91 4.43 2.140 16.70   0  1     5     2
## 10  30.4    4  95.1  113 3.77 1.513 16.90   1  1     5     2
## 11  21.4    4 121.0  109 4.11 2.780 18.60   1  1     4     2
```

```
dbClearResult(rs)
```

18.3   Pass multiple sets of values with dbBind():

```
rs <- dbSendQuery(con, "SELECT * FROM mtcars WHERE cyl = $1")
dbBind(rs, list(6L)) # cyl = 6
dbFetch(rs)
```

```
##     mpg cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  21.0    6 160.0  110 3.90 2.620 16.46   0  1     4     4
## 2  21.0    6 160.0  110 3.90 2.875 17.02   0  1     4     4
## 3  21.4    6 258.0  110 3.08 3.215 19.44   1  0     3     1
## 4  18.1    6 225.0  105 2.76 3.460 20.22   1  0     3     1
## 5  19.2    6 167.6  123 3.92 3.440 18.30   1  0     4     4
## 6  17.8    6 167.6  123 3.92 3.440 18.90   1  0     4     4
## 7  19.7    6 145.0  175 3.62 2.770 15.50   0  1     5     6
```

```
dbBind(rs, list(8L)) # cyl = 8
dbFetch(rs)
```

```
##      mpg cyl  disp  hp drat    wt  qsec vs am gear carb
```

```
## 1   18.7    8 360.0 175 3.15 3.440 17.02  0  0     3     2
## 2   14.3    8 360.0 245 3.21 3.570 15.84  0  0     3     4
## 3   16.4    8 275.8 180 3.07 4.070 17.40  0  0     3     3
## 4   17.3    8 275.8 180 3.07 3.730 17.60  0  0     3     3
## 5   15.2    8 275.8 180 3.07 3.780 18.00  0  0     3     3
## 6   10.4    8 472.0 205 2.93 5.250 17.98  0  0     3     4
## 7   10.4    8 460.0 215 3.00 5.424 17.82  0  0     3     4
## 8   14.7    8 440.0 230 3.23 5.345 17.42  0  0     3     4
## 9   15.5    8 318.0 150 2.76 3.520 16.87  0  0     3     2
## 10 15.2    8 304.0 150 3.15 3.435 17.30  0  0     3     2
## 11 13.3    8 350.0 245 3.73 3.840 15.41  0  0     3     4
## 12 19.2    8 400.0 175 3.08 3.845 17.05  0  0     3     2
## 13 15.8    8 351.0 264 4.22 3.170 14.50  0  1     5     4
## 14 15.0    8 301.0 335 3.54 3.570 14.60  0  1     5     8
```

```r
dbClearResult(rs)
```

18.4   Clean up

```r
# dbRemoveTable(con, "cars")
dbRemoveTable(con, "mtcars")
# dbRemoveTable(con, "cust_movies")

# diconnect from the db
dbDisconnect(con)

sp_docker_stop("sql-pet")
```

Chapter 19

Writing to the DBMS (73)

At the end of this chapter, you will be able to

- Write queries in R using docker container.
- Start and connect to the database with R.
- Create, Modify, and remove the table.

Start up the `docker-pet` container:

```
sp_docker_start("sql-pet")
```

Now connect to the database with R using your login info:

```
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                        password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                        dbname = "dvdrental",
                        seconds_to_test = 10)
```

19.1 Create a new table

This is an example from the DBI help file.

```
dbWriteTable(con, "cars", head(cars, 3)) # "cars" is a built-in dataset, not to be confused with mtcars

dbReadTable(con, "cars")     # there are 3 rows
```

```
##   speed dist
## 1     4    2
## 2     4   10
## 3     7    4
```

19.2 Modify an existing table

To add additional rows or instances to the "cars" table, we will use INSERT command with their values.

There are two different ways of adding values: list them or pass values using the param argument.

```
dbExecute(
  con,
```

```
  "INSERT INTO cars (speed, dist) VALUES (1, 1), (2, 2), (3, 3)"
)
```

```
## [1] 3
```

```r
dbReadTable(con, "cars")    # there are now 6 rows
```

```
##   speed dist
## 1     4    2
## 2     4   10
## 3     7    4
## 4     1    1
## 5     2    2
## 6     3    3
```

```r
# Pass values using the param argument:
dbExecute(
  con,
  "INSERT INTO cars (speed, dist) VALUES ($1, $2)",
  param = list(4:7, 5:8)
)
```

```
## [1] 4
```

```r
dbReadTable(con, "cars")    # there are now 10 rows
```

```
##    speed dist
## 1      4    2
## 2      4   10
## 3      7    4
## 4      1    1
## 5      2    2
## 6      3    3
## 7      4    5
## 8      5    6
## 9      6    7
## 10     7    8
```

19.3   Remove table and Clean up

Here you will remove the table "cars", disconnect from the database and exit docker.

```r
dbRemoveTable(con, "cars")

# diconnect from the db
dbDisconnect(con)

sp_docker_stop("sql-pet")
```

Appendix A

Other resources (89)

A.1 Editing this book

- Here are instructions for editing this tutorial

A.2 Docker alternatives

- Choosing between Docker and Vagrant

A.3 Docker and R

- Noam Ross' talk on Docker for the UseR and his Slides give a lot of context and tips.
- Good Docker tutorials
    - An introductory Docker tutorial
    - A Docker curriculum
- Scott Came's materials about Docker and R on his website and at the 2018 UseR Conference focus on **R inside Docker**.
- It's worth studying the ROpensci Docker tutorial

A.4 Documentation for Docker and Postgres

- The Postgres image documentation
- Dockerize PostgreSQL
- Postgres & Docker documentation
- Usage examples of Postgres with Docker

A.5 SQL and dplyr

- Why SQL is not for analysis but dplyr is
- Data Manipulation with dplyr (With 50 Examples)

A.6   More Resources

- David Severski describes some key elements of connecting to databases with R for MacOS users
- This tutorial picks up ideas and tips from Ed Borasky's Data Science pet containers, which creates a framework based on that Hack Oregon example and explains why this repo is named pet-sql.

Appendix B

Mapping your local environment (92)

B.1 Environment Tools Used in this Chapter

Note that `tidyverse`, `DBI`, `RPostgres`, `glue`, and `knitr` are loaded. Also, we've sourced the `[db-login-batch-code.R]('r-database-docker/book-src/db-login-batch-code.R')` file which is used to log in to PostgreSQL.

library(rstudioapi)

The following code block defines Tool and versions for the graph that follows. The information order corresponds to the order shown in the graph.

```
library(DiagrammeR)

## OS information
os_lbl <- .Platform$OS.type
os_ver <- 0
if (os_lbl == 'windows') {
  os_ver <- system2('cmd',stdout = TRUE) %>%
    grep(x = .,pattern = 'Microsoft Windows \\[',value = TRUE) %>%
    gsub(x = .,pattern = "^Microsoft.+Version |\\]", replace = '')
}

if (os_lbl == 'unix' || os_lbl == 'Linux' || os_lbl == 'Mac') {
  os_ver <- system2('uname', '-r', stdout = TRUE)
}

## Command line interface into Docker Apps
## CLI/system2
cli <- array(dim = 3)
cli[1] <- "docker [OPTIONS] COMMAND ARGUMENTS\n\nsystem2(docker,[OPTIONS,]\n, COMMAND,ARGUMENTS)"
cli[2] <- 'docker exec -it sql-pet bash\n\nsystem2(docker,exec -it sql-pet bash)'
cli[3] <- 'docker exec -ti sql-pet psql -a \n-p 5432 -d dvdrental -U postgres\n\nsystem2(docker,exec -t

# R Information
r_lbl       <- names(R.Version())[1:7]
r_ver       <- R.Version()[1:7]

# RStudio Information
rstudio_lbl <- c('RStudio version','Current program mode')
```

```r
rstudio_ver <- c(as.character(rstudioapi::versionInfo()$version),rstudioapi::versionInfo()$mode)

# Docker Information
docker_lbl <- c('client version','server version')
docker_ver <- system2("docker", "version", stdout = TRUE) %>%
    grep(x = ., pattern = 'Version',value = TRUE) %>%
    gsub(x = ., pattern = ' +Version: +', replacement = '')

# Linux Information
linux_lbl <- 'Linux Version'
linux_ver <- system2('docker', 'exec -i sql-pet /bin/uname -r', stdout = TRUE)

# Postgres Information
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                        password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                        dbname = "dvdrental",
                        seconds_to_test = 10)

postgres_ver <- dbGetQuery(con,"select version()") %>%
  gsub(x = ., pattern = '\\(.*$', replacement = '')
```

The following code block uses the data generated from the previous code block as input to the subgraphs, the ones outlined in red. The application nodes are the parents of the subgraphs and are not outlined in red. The `Environment` application node represents the machine you are running the tutorial on and hosts the sub-applications.

Note that the '@@' variables are populated at the end of the `Environment` definition following the `## @@1 - @@5` source data comment.

```r
grViz("
digraph Envgraph {

  # graph, node, and edge definitions
  graph [compound = true, nodesep = .5, ranksep = .25,
         color = red]

  node [fontname = Helvetica, fontcolor = darkslategray,
        shape = rectangle, fixedsize = true, width = 1,
        color = darkslategray]

  edge [color = grey, arrowhead = none, arrowtail = none]

  # subgraph for Environment information
  subgraph cluster1 {
    node [fixedsize = true, width = 3]
    '@@1-1'
  }

  # subgraph for R information
  subgraph cluster2 {
    node [fixedsize = true, width = 3]
    '@@2-1' -> '@@2-2' -> '@@2-3' -> '@@2-4'
    '@@2-4' -> '@@2-5' -> '@@2-6' -> '@@2-7'
  }
```

```
  # subgraph for RStudio information
  subgraph cluster3 {
    node [fixedsize = true, width = 3]
    '@@3-1' -> '@@3-2'
  }

  # subgraph for Docker information
  subgraph cluster4 {
    node [fixedsize = true, width = 3]
    '@@4-1' -> '@@4-2'
  }

  # subgraph for Docker-Linux information
  subgraph cluster5 {
    node [fixedsize = true, width = 3]
    '@@5-1'
  }

  # subgraph for Docker-Postgres information
  subgraph cluster6 {
    node [fixedsize = true, width = 3]
    '@@6-1'
  }

  # subgraph for Docker-Postgres information
  subgraph cluster7 {
    node [fixedsize = true, height = 1.25, width = 4.0]
    '@@7-1' -> '@@7-2' -> '@@7-3'
  }

  CLI [label='CLI\nRStudio system2',height = .75,width=3.0, color = 'blue' ]
  Environment            [label = 'Linux,Mac,Windows',width = 2.5]
  Environment -> R
  Environment -> RStudio
  Environment -> Docker

  Environment -> '@@1'    [lhead = cluster1] # Environment Information
  R           -> '@@2-1'  [lhead = cluster2] # R Information
  RStudio     -> '@@3'    [lhead = cluster3] # RStudio Information
  Docker      -> '@@4'    [lhead = cluster4] # Docker Information
  Docker      -> '@@5'    [lhead = cluster5] # Docker-Linux Information
  Docker      -> '@@6'    [lhead = cluster6] # Docker-Postgres Information

  '@@1' -> CLI
  CLI         -> '@@7'    [lhead = cluster7] # CLI
  '@@7-2'     -> '@@5'
  '@@7-3'     -> '@@6'
}
[1]: paste0(os_lbl,      ':\\n', os_ver)
[2]: paste0(r_lbl,       ':\\n', r_ver)
[3]: paste0(rstudio_lbl,':\\n', rstudio_ver)
[4]: paste0(docker_lbl, ':\\n', docker_ver)
[5]: paste0(linux_lbl,  ':\\n', linux_ver)
```

```
[6]: paste0('PostgreSQL:\\n', postgres_ver)
[7]: cli
")
```

One sub-application not shown above is your local console/terminal/CLI application. In the tutorial, fully constructed docker commands are printed out and then executed. If for some reason the executed docker command fails, one can copy and paste it into your local terminal window to see additional error information. Failures seem more prevalent in the Windows environment.

B.2   Communicating with Docker Applications

In this tutorial, the two main ways to interface with the applications in the Docker container are through the CLI or the RStudio `system2` command. The blue box in the diagram above represents these two interfaces.

Appendix C

Creating the sql-pet Docker container one step at a time (93)

Step-by-step Docker container setup with dvdrental database installed This needs to run *outside a project* to compile correctly because of the complexities of how knitr sets working directories (or because we don't really understand how it works!) The purpose of this code is to

- Replicate the docker container generated in Chapter 5 of the book, but in a step-by-step fashion
- Show that the `dvdrental` database persists when stopped and started up again.

C.1 Overview

Doing all of this in a step-by-step way that might be useful to understand how each of the steps involved in setting up a persistent PostgreSQL database works. If you are satisfied with the method shown in Chapter 5, skip this and only come back if you're interested in picking apart the steps.

```r
library(tidyverse)
```

```
## -- Attaching packages ------------------------------------------------------------------ tidyvers
## v ggplot2 3.1.0     v purrr   0.2.5
## v tibble  1.4.2     v dplyr   0.7.8
## v tidyr   0.8.2     v stringr 1.3.1
## v readr   1.3.0     v forcats 0.3.0
## -- Conflicts --------------------------------------------------------------------------- tidyverse_conf
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```r
library(DBI)
library(RPostgres)
library(glue)
```

```
##
## Attaching package: 'glue'
## The following object is masked from 'package:dplyr':
##
##     collapse
```

```r
require(knitr)
```

```
## Loading required package: knitr
```

```r
library(dbplyr)
```

```
##
## Attaching package: 'dbplyr'
```

```
## The following objects are masked from 'package:dplyr':
##
##     ident, sql
```

```r
library(sqlpetr)
library(here)
```

```
## here() starts at /Users/jds/Documents/Library/R/r-system/sql-pet
```

C.2 Download the `dvdrental` backup file

The first step is to get a local copy of the `dvdrental` PostgreSQL **restore file**. It comes in a zip format and needs to be un-zipped.

```r
opts_knit$set(root.dir = normalizePath('../'))
if (!require(downloader)) install.packages("downloader")
```

```
## Loading required package: downloader
```

```r
library(downloader)

download("http://www.postgresqltutorial.com/wp-content/uploads/2017/10/dvdrental.zip", destfile = glue(

unzip("dvdrental.zip", exdir = here()) # creates a tar archhive named "dvdrental.tar"
```

Check on where we are and what we have in this directory:

```r
dir(path = here(), pattern = "^dvdrental(.tar|.zip)")
```

```
## [1] "dvdrental.tar" "dvdrental.zip"
```

```r
sp_show_all_docker_containers()
```

```
## CONTAINER ID     IMAGE              COMMAND              CREATED           STATUS              PORTS
## 5802b471a9cf     postgres-dvdrental "docker-entrypoint.s…" About a minute ago Exited (0) 5 seconds a
```

Remove the `sql-pet` container if it exists (e.g., from a prior run)

```r
if (system2("docker", "ps -a", stdout = TRUE) %>%
    grepl(x = ., pattern = 'sql-pet') %>%
    any()) {
  sp_docker_remove_container("sql-pet")
}
```

```
## [1] 0
```

C.3   Build the Docker Container

Build an image that derives from postgres:10. Connect the local and Docker directories that need to be shared. Expose the standard PostgreSQL port 5432.

```r
wd <- here()
wd
```

```
## [1] "/Users/jds/Documents/Library/R/r-system/sql-pet"
```

```r
docker_cmd <- glue(
  "run ",          # Run is the Docker command.  Everything that follows are `run` parameters.
  "--detach ", # (or `-d`) tells Docker to disconnect from the terminal / program issuing the command
  " --name sql-pet ",      # tells Docker to give the container a name: `sql-pet`
  "--publish 5432:5432 ", # tells Docker to expose the Postgres port 5432 to the local network with 543.
  "--mount ", # tells Docker to mount a volume -- mapping Docker's internal file structure to the host
  'type=bind,source="', wd, '",target=/petdir',
  " postgres:10 " # tells Docker the image that is to be run (after downloading if necessary)
)

docker_cmd
```

```
## run --detach  --name sql-pet --publish 5432:5432 --mount type=bind,source="/Users/jds/Documents/Library,
```

```r
system2("docker", docker_cmd, stdout = TRUE, stderr = TRUE)
```

```
## [1] "3c362230c39e14629c8889662ec2a22843c0e085f0842aba7b28d931abbeb5d5"
```

Peek inside the docker container and list the files in the `petdir` directory. Notice that `dvdrental.tar` is in both.

```r
# local file system:
dir(path = here(), pattern = "^dvdrental.tar")
```

```
## [1] "dvdrental.tar"
```

```r
# inside docker
system2('docker', 'exec sql-pet ls petdir | grep "dvdrental.tar" ',
        stdout = TRUE, stderr = TRUE)
```

```
## [1] "dvdrental.tar"
```

```r
Sys.sleep(3)
```

C.4   Create the database and restore from the backup

We can execute programs inside the Docker container with the `exec` command. In this case we tell Docker to execute the `psql` program inside the `sql-pet` container and pass it some commands as follows.

```r
sp_show_all_docker_containers()
```

```
## CONTAINER ID      IMAGE          COMMAND           CREATED          STATUS          PORTS            
## 3c362230c39e      postgres:10      "docker-entrypoint.s…"  4 seconds ago    Up 3 seconds    0.0.0.0:5
```

inside Docker, execute the postgress SQL command-line program to create the dvdrental database:

```r
system2('docker', 'exec sql-pet psql -U postgres -c "CREATE DATABASE dvdrental;"',
        stdout = TRUE, stderr = TRUE)
```

```
## [1] "CREATE DATABASE"
```

```r
Sys.sleep(3)
```

The **psql** program repeats back to us what it has done, e.g., to create a database named `dvdrental`. Next we execute a different program in the Docker container, **pg_restore**, and tell it where the restore file is located. If successful, the **pg_restore** just responds with a very laconic `character(0)`.

```r
system2("docker", "exec sql-pet pg_restore -U postgres -d dvdrental petdir/dvdrental.tar", stdout = TRUE
```

```
## character(0)
```

```r
Sys.sleep(3)
```

C.5   Connect to the database with R

If you are interested take a look inside the `sp_get_postgres_connection` function to see how the DBI package is beingcused.

```r
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                                  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                                  dbname = "dvdrental",
                                  seconds_to_test = 20)
```

```r
dbListTables(con)
```

```
##  [1] "actor_info"             "customer_list"
##  [3] "film_list"              "nicer_but_slower_film_list"
##  [5] "sales_by_film_category" "staff"
##  [7] "sales_by_store"         "staff_list"
##  [9] "category"               "film_category"
## [11] "country"                "actor"
## [13] "language"               "inventory"
## [15] "payment"                "rental"
## [17] "city"                   "store"
## [19] "film"                   "address"
## [21] "film_actor"             "customer"
```

```r
dbDisconnect(con)
```

```r
# Stop and start to demonstrate persistence
```

Stop the container

```r
sp_docker_stop("sql-pet")
```

Restart the container and verify that the dvdrental tables are still there

```r
sp_docker_start("sql-pet")
```

```r
con <- sp_get_postgres_connection(user = Sys.getenv("DEFAULT_POSTGRES_USER_NAME"),
                                  password = Sys.getenv("DEFAULT_POSTGRES_PASSWORD"),
                                  dbname = "dvdrental",
                                  seconds_to_test = 10)
```

```r
dbListTables(con)
```

```
##  [1] "actor_info"             "customer_list"
##  [3] "film_list"              "nicer_but_slower_film_list"
##  [5] "sales_by_film_category" "staff"
```

```
##  [7] "sales_by_store"           "staff_list"
##  [9] "category"                 "film_category"
## [11] "country"                  "actor"
## [13] "language"                 "inventory"
## [15] "payment"                  "rental"
## [17] "city"                     "store"
## [19] "film"                     "address"
## [21] "film_actor"               "customer"
```

C.6 Cleaning up

It's always good to have R disconnect from the database

```r
dbDisconnect(con)
```

Stop the container and show that the container is still there, so can be started again.

```r
sp_docker_stop("sql-pet")
```

show that the container still exists even though it's not running

```r
sp_show_all_docker_containers()
```

```
## CONTAINER ID     IMAGE         COMMAND              CREATED         STATUS                     PORTS
## 3c362230c39e     postgres:10     "docker-entrypoint.s…"  15 seconds ago     Exited (0) Less than a seco
```

We are leaving the `sql-pet` container intact so it can be used in running the rest of the examples and book.

Clean up by removing the local files used in creating the database:

```r
file.remove(here("dvdrental.zip"))
```

```
## [1] TRUE
```

```r
file.remove(here("dvdrental.tar"))
```

```
## [1] TRUE
```

Appendix D

APPENDIX D - Quick Guide to SQL (94)

SQL stands for Structured Query Language. It is a database language where we can perform certain operations on the existing database and we can use it create a new database. There are four main categories where the SQL commands fall into: DDL, DML, DCL, and TCL.

##Data Definition Langauge (DDL)

It consists of the SQL commands that can be used to define database schema. The DDL commands include:

```
1. CREATE
2. ALTER
3. TRUNCATE
4. COMMENT
5. RENAME
6. DROP
```

##Data Manipulation Langauge (DML)

These four SQL commands deals with the manipulation of data in the database.

```
1. SELECT
2. INSERT
3. UPDATE
4. DELETE
```

##Data Control Language (DCL)

The DCL commands deals with user's rights, permissions and other controls in database management system.

```
1. GRANT
2. REVOKE
```

##Transaction Control Language (TCL)

These commands deals with the control over transaction within the database. Transaction combines a set of tasks into single execution.

```
1. SET TRANSACTION
2. SAVEPOINT
3. ROLLBACK
4. COMMIT
```

Appendix E

Additional technical details for Windows users (95)

E.1   Docker for Windows settings

E.1.1   Shared drives

If you're going to mount host files into container file systems (as we do in the following chapters), you need to set up shared drives. Open the Docker settings dialog and select `Shared Drives`. Check the drives you want to share. In this screenshot, the `D:` drive is my 1 terabyte hard drive.
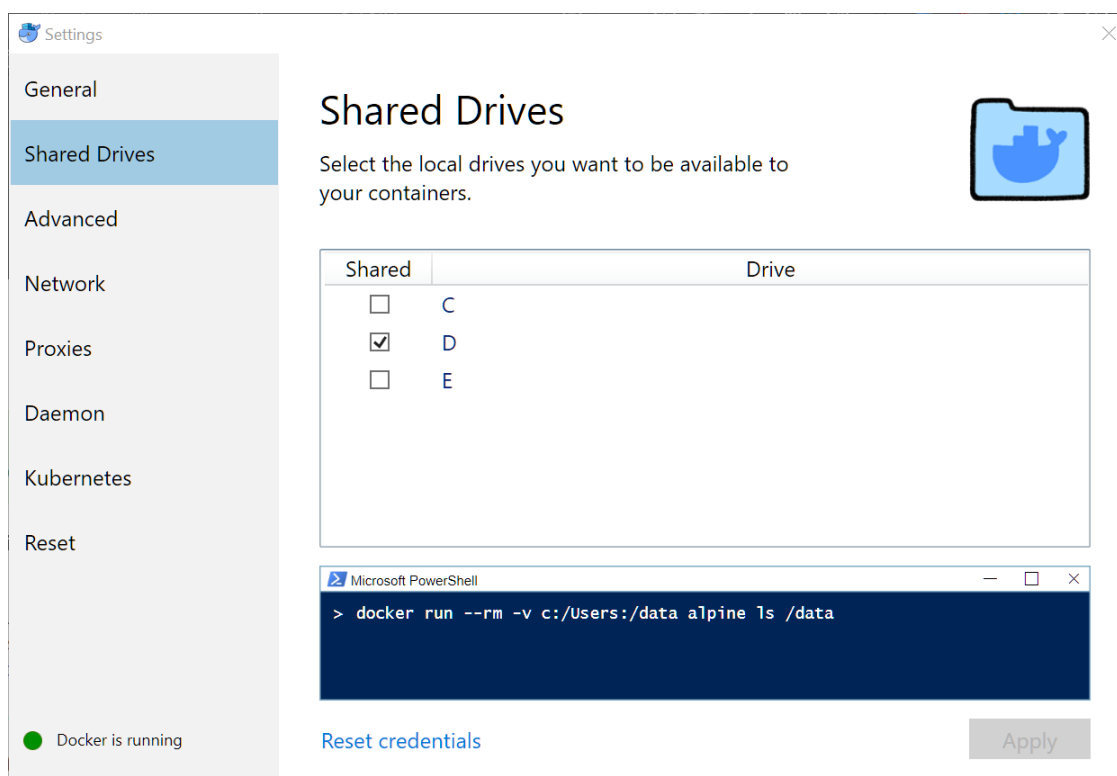
E.1.2   Kubernetes

Kubernetes is a container orchestration / cloud management package that's a major DevOps tool. It's heavily supported by Red Hat and Google, and as a result is becoming a required skill for DevOps.

However, it's overkill for this project at the moment. So you should make sure it's not enabled.

Go to the `Kubernetes` dialog and make sure the `Enable Kubernetes` checkbox is cleared.
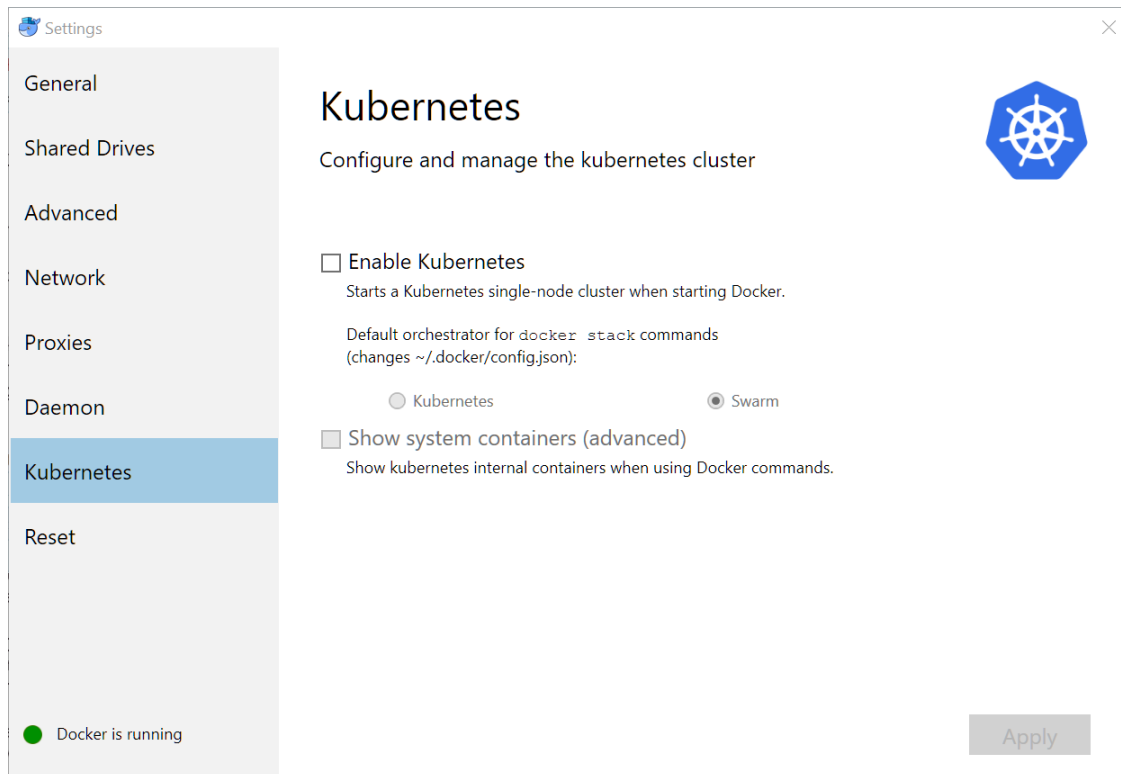
E.2   Git, GitHub and line endings

Git was originally developed for Linux - in fact, it was created by Linus Torvalds to manage hundreds of different versions of the Linux kernel on different machines all around the world. As usage has grown, Git has achieved a huge following and is the version control system used by most large open source projects, including this one.

If you're on Windows, there are some things about Git and GitHub you need to watch. First of all, there are quite a few tools for running Git on Windows, but the RStudio default and recommended one is Git for Windows (https://git-scm.com/download/win).

By default, text files on Linux end with a single linefeed (`\n`) character. But on Windows, text files end with a carriage return and a line feed (`\r\n`). See https://en.wikipedia.org/wiki/Newline for the gory details.

Git defaults to checking files out in the native mode. So if you're on Linux, a text file will show up with the Linux convention, and if you're on Windows, it will show up with the Windows convention.

Most of the time this doesn't cause any problems. But Docker containers usually run Linux, and if you have files from a repository on Windows that you've sent to the container, the container may malfunction or give weird results. *This kind of situation has caused a lot of grief for contributors to this project, so beware.*

In particular, executable `sh` or `bash` scripts will fail in a Docker container if they have Windows line endings. You may see an error message with `\r` in it, which means the shell saw the carriage return (`\r`) and gave up.

But often you'll see no hint at all what the problem was.

So you need a way to tell Git that some files need to be checked out with Linux line endings. See https: //help.github.com/articles/dealing-with-line-endings/ for the details. Summary:

1. You'll need a `.gitattributes` file in the root of the repository.
2. In that file, all text files (scripts, program source, data, etc.) that are destined for a Docker container will need to have the designator `<spec> text eol=lf`, where `<spec>` is the file name specifier, for example, `*.sh`.

This repo includes a sample: .gitattributes

Appendix F

Dplyr functions and SQL cross-walk (96)

Where are these covered and should they be included?

Dplyr Function	description	SQL Clause	When	Category
all_equal() all.equal()	Flexible equality comparison for data frames			Two-table verbs
all_vars() any_vars()	Apply predicate to all variables			scoped-Operate on a selection of variables
arrange()	Arrange rows by variables	ORDER BY	13.1.4 (21)	Basic single-table verbs
arrange_all() arrange_at() arrange_if()	Arrange rows by a selection of variables	ORDER BY		scoped-Operate on a selection of variables
auto_copy()	Copy tables to same source, if necessary			Remote tables
between()	Do values in a numeric vector fall in specified range?			Vector functions
bind_rows() bind_cols() combine()	Efficiently bind multiple data frames by row and column			Two-table verbs
case_when()	A general vectorised if			Vector functions
coalesce()	Find first non-missing element			Vector functions
compute() collect() collapse()	Force computation of a database query			Remote tables
copy_to()	Copy a local data frame to a remote src			Remote tables
cumall() cumany() cummean()	Cumulativate versions of any, all, and mean			Vector functions
desc()	Descending order			Vector functions

Dplyr Function	description	SQL Clause	When	Category
distinct()	Return rows with matching conditions	SELECT distinct *		Basic single-table verbs
distinct()	Select distinct/unique rows	SELECT distinct {colname1,...colnamen}		Basic single-table verbs
do()	Do anything	NA		Basic single-table verbs
explain() show_query()	Explain details of a tbl			Remote tables
filter_all() filter_if() filter_at()	Filter within a selection of variables			scoped-Operate on a selection of variables
funs()	Create a list of functions calls.			scoped-Operate on a selection of variables
group_by() ungroup()	Objects exported from other packages	GROUP BY no ungroup		Basic single-table verbs
group_by_all() group_by_at() group_by_if()	Group by a selection of variables			scoped-Operate on a selection of variables
groups() group_vars()	Return grouping variables			Metadata
ident()	Flag a character vector as SQL identifiers			Remote tables
if_else()	Vectorised if			Vector functions
inner_join() left_join() right_join() full_join() semi_join() anti_join()	Join two tbls together			Two-table verbs
inner_join()left_join() right_join() full_join() semi_join() anti_join()	Join data frame tbls			Two-table verbs
intersect() union() union_all() setdiff() setequal()	Set operations			Two-table verbs
lead() lag()	Lead and lag.			Vector functions
mutate() transmute()	Add new variables	SELECT computed_value computed_name	11.5.2 (13)	Basic single-table verbs
n()	The number of observations in the current group.			Vector functions
n_distinct()	Efficiently count the number of unique values in a set of vector			Vector functions
na_if()	Convert values to NA			Vector functions
near()	Compare two numeric vectors			Vector functions
nth() first() last()	Extract the first, last or nth value from a vector			Vector functions

Dplyr Function	description	SQL Clause	When	Category
order_by()	A helper function for ordering window function output			Vector functions
pull()	Pull out a single variable	SELECT column_name;		Basic single-table verbs
recode() recode_factor()	Recode values			Vector functions
row_number() ntile() min_rank() dense_rank() percent_rank() cume_dist()	Windowed rank functions.			Vector functions
rowwise()	Group input by rows			Other backends
sample_n() sample_frac()	Sample n rows from a table	ORDER BY RANDOM() LIMIT 10		Basic single-table verbs
select() rename()	Select/rename variables by name	SELECT column_name alias_name	9.1.8 (11)	Basic single-table verbs
select_all() rename_all() select_if() rename_if() select_at() rename_at()	Select and rename a selection of variables			scoped-Operate on a selection of variables
slice()	Select rows by position	SELECT row_number() over (partition by expression(s) order_by exp)		Basic single-table verbs
sql()	SQL escaping.			Remote tables
src_mysql() src_postgres() src_sqlite()	Source for database backends			Remote tables
summarise_all() summarise_if() summarise_at() summarize_all() summarize_if() summarize_at() mutate_all() mutate_if() mutate_at() transmute_all() transmute_if() transmute_at()	Summarise and mutate multiple columns.			scoped-Operate on a selection of variables
summarize()	Reduces multiple values down to a single value	SELECT aggregate_functions GROUP BY	11.5. (13)	Basic single-table verbs
tally() count()add_tally() add_count()	Count/tally observations by group	GROUP BY	9.1.6 (11)	Single-table helpers
tbl() is.tbl() as.tbl()	Create a table from a data source			Remote tables

Dplyr Function	description	SQL Clause	Where	Category
top_n()	Select top (or bottom) n rows (by value)	ORDER BY VALUE {DESC} LIMIT 10		Single-table helpers
vars()	Select variables			scoped-Operate on a selection of variables

Appendix G

DBI package functions - coverage (96b)

Where are these covered and should the by included?

DBI	1st time	Call Example/Notes
DBIConnct dbAppendTable dbCreateTable	6.3.2 (04)	in sp_get_postgres_connection
dbDisconnect	6.4n (04)	dbDisconnect(con)
dbExecute	10.4.2 (13)	Executes a statement and returns the number of rows affected. dbExecute() comes with a default implementation (which should work with most backends) that calls dbSendStatement(), then dbGetRowsAffected(), ensuring that the result is always free-d by dbClearResult().
dbExistsTable		dbExistsTable(con,'actor')
dbFetch dbGetException	17.1 (72)	dbFetch(rs)
dbGetInfo		dbGetInfo(con)
dbGetQuery	10.4.1 (13)	dbGetQuery(con,'select * from store;')
dbIsReadOnly		dbIsReadOnly(con)
dbIsValid		dbIsValid(con)
dbListFields	6.3.3 (04)	DBI::dbListFields(con, "mtcars")
dbListObjects		dbListObjects(con)
dbListTables	6.3.2 (04)	DBI::dbListTables(con, con)
dbReadTable dbRemoveTable	8.1.2	DBI::dbReadTable(con, "rental")
dbSendQuery	17.1 (72)	rs <- dbSendQuery(con, "SELECT * FROM mtcars WHERE cyl = 4")
dbSendStatement		The dbSendStatement() method only submits and synchronously executes the SQL data manipulation statement (e.g., UPDATE, DELETE, INSERT INTO, DROP TABLE, …) to the database engine.
dbWriteTable	6.3.3 (04)	dbWriteTable(con, "mtcars", mtcars, overwrite = TRUE)

Bibliography

Baumer, B. S. (2018). A grammar for reproducible and painless extract-transform-load operations on medium data. *Journal of Computational and Graphical Statistics*, pages 1–9.

Wickham, H. (2018). A grammar of data manipulation • dplyr. .